

```
In [1]: import kagglehub
```

```
# Download latest version
path = kagglehub.dataset_download("sohansakib75/pmrabrain")
```

```
print("Path to dataset files:", path)
```

```
Path to dataset files: /kaggle/input/pmrabrain
```

```
In [2]: import os
```

```
for root, dirs, files in os.walk("/kaggle/input/pmrabrain"):
    for d in dirs:
        print(os.path.join(root, d))
```

```
/kaggle/input/pmrabrain/Raw
```

```
/kaggle/input/pmrabrain/Raw/512Glioma
```

```
/kaggle/input/pmrabrain/Raw/512Meningioma
```

```
/kaggle/input/pmrabrain/Raw/512Pituitary
```

```
/kaggle/input/pmrabrain/Raw/512Normal
```

```
In [3]: import os
```

```
import cv2
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from skimage.feature import peak_local_max
```

```
# Paths and params
```

```
base_path = "/kaggle/input/pmrabrain/Raw"
```

```
output_base = "/kaggle/working/preprocessed"
```

```
class_dirs = ["512Glioma", "512Meningioma", "512Pituitary", "512Normal"]
```

```
img_size = (224, 224)
```

```
def ensure_dir(path):
```

```
    if not os.path.exists(path):
        os.makedirs(path)
```

```
def save_image(path, img):
    cv2.imwrite(path, img)
```

```
def show_images(images, titles):
```

```
    plt.figure(figsize=(18, 6))
```

```
    for i, (img, title) in enumerate(zip(images, titles)):
        plt.subplot(1, len(images), i+1)
```

```
        if len(img.shape) == 2:
```

```
            plt.imshow(img, cmap='gray')
```

```
        else:
```

```
            plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
```

```
        plt.title(title)
```

```
        plt.axis('off')
```

```
    plt.tight_layout()
```

```
    plt.show()
```

```
def preprocess_image(img):
```

```
    results = {}
```

```
# Step 1: Resize
```

```
    img_resized = cv2.resize(img, img_size)
```

```
    results['resized'] = img_resized
```

```
# Step 2: Z-score normalization (grayscale)
```

```
    img_gray = cv2.cvtColor(img_resized, cv2.COLOR_BGR2GRAY)
```

```
    img_norm = (img_gray - np.mean(img_gray)) / (np.std(img_gray) + 1e-8)
```

```
    img_norm = np.clip(img_norm, -3, 3)
```

```
    img_norm = ((img_norm - img_norm.min()) / (img_norm.max() - img_norm.min()) * 255).astype(np.uint8)
```

```
    results['normalized'] = img_norm
```

```
# Step 3: CLAHE
```

```
    clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8,8))
```

```
    img_clahe = clahe.apply(img_norm)
```

```
    results['clahe'] = img_clahe
```

```
# Step 4: Denoising
```

```
    img_denoised = cv2.fastNlMeansDenoising(img_clahe, h=10, templateWindowSize=7, searchWindowSize=21)
```

```
    results['denoised'] = img_denoised
```

```
return results
```

```
# Visualize and save images for one example per class
```

```
for class_name in class_dirs:
```

```
    class_path = os.path.join(base_path, class_name)
```

```
    save_class_base = os.path.join(output_base, class_name)
```

```
    sample_image_name = next((f for f in os.listdir(class_path) if f.lower().endswith('.png', '.jpg', '.jpeg')))
```

```

if sample_image_name:
    img_path = os.path.join(class_path, sample_image_name)
    img = cv2.imread(img_path) # Read as color BGR
    print(f"\nClass: {class_name}, Image: {sample_image_name}")

    processed_imgs = preprocess_image(img)

    # Show preprocessing steps (excluding watershed color overlay and final grayscale)
    show_images(
        [img, processed_imgs['resized'], processed_imgs['normalized'], processed_imgs['clahe'],
        processed_imgs['denoised']],
        ['Original', 'Resized', 'Normalized', 'CLAHE', 'Denoised']
    )

    # Save all preprocessed images for this class (one sample)
    for step_name, proc_img in processed_imgs.items():
        save_folder = os.path.join(save_class_base, step_name)
        ensure_dir(save_folder)
        save_path = os.path.join(save_folder, os.path.splitext(sample_image_name)[0] + '.png')
        save_image(save_path, proc_img)

# Batch preprocess and save all images (comment out to speed up)
print("\nStarting batch preprocessing & saving all images...")

for class_name in class_dirs:
    class_path = os.path.join(base_path, class_name)
    save_class_base = os.path.join(output_base, class_name)


    for filename in os.listdir(class_path):
        if not filename.lower().endswith(('.png', '.jpg', '.jpeg')):
            continue
        img_path = os.path.join(class_path, filename)
        img = cv2.imread(img_path) # Read as color BGR

        processed_imgs = preprocess_image(img)
        for step_name, proc_img in processed_imgs.items():
            save_folder = os.path.join(save_class_base, step_name)
            ensure_dir(save_folder)
            save_path = os.path.join(save_folder, os.path.splitext(filename)[0] + '.png')
            save_image(save_path, proc_img)


print("\nAll images preprocessed and saved in /kaggle/working/preprocessed/")

```


Class: 512Glioma, Image: glioma (123).jpg




Class: 512Meningioma, Image: M\_187.jpg



Class: 512Pituitary, Image: pituitary (205).jpg



Class: 512Normal, Image: normal (279).jpg



Starting batch preprocessing & saving all images...

All images preprocessed and saved in /kaggle/working/preprocessed/

```
In [4]: import os
import cv2
import random

# Paths and params
preprocessed_base = "/kaggle/working/preprocessed"
augmented_base = "/kaggle/working/augmented"
class_dirs = ["512Glioma", "512Meningioma", "512Pituitary", "512Normal"]
target_total = 5500 # approximate total images across all classes

def ensure_dir(path):
    if not os.path.exists(path):
        os.makedirs(path)

def save_image(path, img):
    cv2.imwrite(path, img)

def augment_image_flip_rotate(img):
    """Return a dictionary of possible flip and rotation augmentations."""
    aug_dict = {}
    # Horizontal flip
    aug_dict['hflip'] = cv2.flip(img, 1)
    # Vertical flip
    aug_dict['vflip'] = cv2.flip(img, 0)
    # Rotations ±15 degrees
    for angle in [-15, -10, -5, 5, 10, 15]:
        h, w = img.shape[:2]
        M = cv2.getRotationMatrix2D((w//2, h//2), angle, 1)
        aug_dict[f'rot{angle}'] = cv2.warpAffine(img, M, (w, h), borderMode=cv2.BORDER_REFLECT)
    return aug_dict

# Dictionary to hold summary
summary = {}

# Augment only the final preprocessed image (denoised)
for class_name in class_dirs:
    denoised_path = os.path.join(preprocessed_base, class_name, "denoised")
    save_aug_class_path = os.path.join(augmented_base, class_name, "denoised")
    ensure_dir(save_aug_class_path)
    summary[class_name] = 0

    for filename in os.listdir(denoised_path):
        if not filename.lower().endswith('.png'):
            continue
        img_path = os.path.join(denoised_path, filename)
        img = cv2.imread(img_path, cv2.IMREAD_UNCHANGED)

        # Save original
        save_image(os.path.join(save_aug_class_path, filename), img)
        summary[class_name] += 1

        # Decide how many augmentations per image
        aug_dict = augment_image_flip_rotate(img)
        aug_per_image = 2 # choose 2 random augmentations per image
        selected_keys = random.sample(list(aug_dict.keys()), min(aug_per_image, len(aug_dict)))

        for key in selected_keys:
            save_path = os.path.join(save_aug_class_path, f"{os.path.splitext(filename)[0]}_{key}.png")
            save_image(save_path, aug_dict[key])
            summary[class_name] += 1

# Print summary
print("\nAugmentation Summary (including original):")
total_aug = 0
for class_name, count in summary.items():
    print(f"Class: {class_name} -> Total images: {count}")
```

```

total_aug += count

print(f"\nTotal images across all classes: {total_aug}")
print(f"\nAll augmented images saved in {augmented_base}")

```

Augmentation Summary (including original):  
 Class: 512Glioma -> Total images: 1119  
 Class: 512Meningioma -> Total images: 1089  
 Class: 512Pituitary -> Total images: 1119  
 Class: 512Normal -> Total images: 1188

Total images across all classes: 4515

All augmented images saved in /kaggle/working/augmented

```
In [5]: import os
import cv2
import random
import shutil

# Paths
augmented_base = "/kaggle/working/augmented"
split_base = "/kaggle/working/augmented_split"
class_dirs = ["512Glioma", "512Meningioma", "512Pituitary", "512Normal"]

# Split ratios
train_ratio = 0.75
val_ratio = 0.10
test_ratio = 0.15

def ensure_dir(path):
    if not os.path.exists(path):
        os.makedirs(path)

# Create split directories
for split in ['train', 'val', 'test']:
    for class_name in class_dirs:
        ensure_dir(os.path.join(split_base, split, class_name))

# Function to split images
split_summary = {}

for class_name in class_dirs:
    class_aug_path = os.path.join(augmented_base, class_name)

    # Gather all images from all preprocessing steps
    all_imgs = []
    for step_name in os.listdir(class_aug_path):
        step_path = os.path.join(class_aug_path, step_name)
        for fname in os.listdir(step_path):
            if fname.lower().endswith('.png'):
                all_imgs.append(os.path.join(step_path, fname))

    random.shuffle(all_imgs)
    total = len(all_imgs)
    train_end = int(total * train_ratio)
    val_end = train_end + int(total * val_ratio)

    train_imgs = all_imgs[:train_end]
    val_imgs = all_imgs[train_end:val_end]
    test_imgs = all_imgs[val_end:]

    # Copy images to respective folders
    for img_path in train_imgs:
        shutil.copy(img_path, os.path.join(split_base, 'train', class_name))
    for img_path in val_imgs:
        shutil.copy(img_path, os.path.join(split_base, 'val', class_name))
    for img_path in test_imgs:
        shutil.copy(img_path, os.path.join(split_base, 'test', class_name))

    # Save summary for this class
    split_summary[class_name] = {
        'total': total,
        'train': len(train_imgs),
        'val': len(val_imgs),
        'test': len(test_imgs)
    }

# Print summary
print("Augmented Dataset Split Summary:")
for class_name, counts in split_summary.items():
    print(f"{class_name}: Total={counts['total']}, Train={counts['train']}, Val={counts['val']}, Test={counts['test']}")

print("\nAll augmented images split and saved in {split_base}")
```

```
Augmented Dataset Split Summary:  
512Glioma: Total=1119, Train=839, Val=111, Test=169  
512Meningioma: Total=1089, Train=816, Val=108, Test=165  
512Pituitary: Total=1119, Train=839, Val=111, Test=169  
512Normal: Total=1188, Train=891, Val=118, Test=179
```

```
All augmented images split and saved in /kaggle/working/augmented_split
```

```
In [6]: # =====  
# Swin-Tiny Training + Evaluation Pipeline  
# =====  
  
import torch  
import torch.nn as nn  
import torch.optim as optim  
from torch.utils.data import DataLoader  
from torchvision import datasets, transforms  
from timm import create_model  
from torch.optim.lr_scheduler import ReduceLROnPlateau  
  
from sklearn.metrics import (  
    classification_report, confusion_matrix, matthews_corrcoef,  
    roc_auc_score, average_precision_score, roc_curve,  
    precision_recall_curve, cohen_kappa_score  
)  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
import time  
  
# =====  
# Config  
# =====  
data_dir = "/kaggle/working/augmented_split"  
batch_size = 32  
num_epochs = 35  
patience = 5  
num_classes = 4 # Glioma, Meningioma, Pituitary, Normal  
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
class_names = ["Glioma", "Meningioma", "Pituitary", "Normal"]  
  
# =====  
# Data (no extra augmentation)  
# =====  
common_tfms = transforms.Compose([  
    transforms.Resize((224, 224)),  
    transforms.ToTensor(),  
    transforms.Normalize([0.5]*3, [0.5]*3)  
)  
  
train_ds = datasets.ImageFolder(root=f"{data_dir}/train", transform=common_tfms)  
val_ds = datasets.ImageFolder(root=f"{data_dir}/val", transform=common_tfms)  
test_ds = datasets.ImageFolder(root=f"{data_dir}/test", transform=common_tfms)  
  
train_loader = DataLoader(train_ds, batch_size=batch_size, shuffle=True, num_workers=2)  
val_loader = DataLoader(val_ds, batch_size=batch_size, shuffle=False, num_workers=2)  
test_loader = DataLoader(test_ds, batch_size=batch_size, shuffle=False, num_workers=2)  
  
# =====  
# Model  
# =====  
model = create_model("swin_tiny_patch4_window7_224", pretrained=True, num_classes=num_classes)  
model = model.to(device)  
  
criterion = nn.CrossEntropyLoss()  
optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)  
scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=True)  
  
# =====  
# Training Loop with Early Stopping  
# =====  
best_val_loss = float("inf")  
patience_counter = 0  
  
for epoch in range(num_epochs):  
    start_time = time.time()  
  
    # ---- Train ----  
    model.train()  
    train_loss, train_correct = 0, 0  
    for imgs, labels in train_loader:  
        imgs, labels = imgs.to(device), labels.to(device)
```

```

optimizer.zero_grad()
outputs = model(imgs)
loss = criterion(outputs, labels)
loss.backward()
optimizer.step()

train_loss += loss.item() * imgs.size(0)
train_correct += (outputs.argmax(1) == labels).sum().item()

train_loss /= len(train_loader.dataset)
train_acc = train_correct / len(train_loader.dataset)

# ---- Validation ----
model.eval()
val_loss, val_correct = 0, 0
with torch.no_grad():
    for imgs, labels in val_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        outputs = model(imgs)
        loss = criterion(outputs, labels)

        val_loss += loss.item() * imgs.size(0)
        val_correct += (outputs.argmax(1) == labels).sum().item()

val_loss /= len(val_loader.dataset)
val_acc = val_correct / len(val_loader.dataset)

scheduler.step(val_loss)

elapsed = time.time() - start_time
print(f"Epoch [{epoch+1}/{num_epochs}] "
      f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f} "
      f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f} "
      f"Time: {elapsed:.2f}s")

# Early Stopping
if val_loss < best_val_loss:
    best_val_loss = val_loss
    patience_counter = 0
    torch.save(model.state_dict(), "swin_tiny_best.pth")
else:
    patience_counter += 1
    if patience_counter >= patience:
        print("Early stopping triggered!")
        break

print("Training finished Best model saved as swin_tiny_best.pth")

# =====
# Evaluation Functions
# =====
def plot_confusion_matrix(cm, classes, title="Confusion Matrix"):
    plt.figure(figsize=(6,5))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=classes, yticklabels=classes)
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title(title)
    plt.show()

def print_report(y_true, y_pred, y_prob, title="", plot_curves=False):
    print(f"\n--- {title} ---")
    print(classification_report(y_true, y_pred, digits=4))

    # Confusion Matrix
    cm = confusion_matrix(y_true, y_pred)
    plot_confusion_matrix(cm, class_names, title=f"{title} Confusion Matrix")

    # MCC
    mcc = matthews_corrcoef(y_true, y_pred)
    print(f"MCC: {mcc:.4f}")

    # Cohen's Kappa
    kappa = cohen_kappa_score(y_true, y_pred)
    print(f"Cohen's Kappa: {kappa:.4f}")

    # Class-wise NPV + PPV
    npv_list, ppv_list = [], []
    for i in range(len(cm)):
        TN = np.sum(np.delete(np.delete(cm, i, axis=0), i, axis=1))
        FN = np.sum(cm[i, :]) - cm[i, i]
        FP = np.sum(cm[:, i]) - cm[i, i]
        TP = cm[i, i]

```

```

NPV = TN / (TN + FN) if (TN + FN) > 0 else 0
PPV = TP / (TP + FP) if (TP + FP) > 0 else 0
npv_list.append(NPV)
ppv_list.append(PPV)

print(f"Mean NPV: {np.mean(npv_list):.4f}")
print(f"Mean PPV (Precision): {np.mean(ppv_list):.4f}")

if plot_curves:
    # ROC AUC + PR AUC
    y_onehot = np.eye(num_classes)[y_true]
    roc_auc = roc_auc_score(y_onehot, y_prob, average='macro', multi_class='ovr')
    pr_auc = average_precision_score(y_onehot, y_prob, average='macro')
    print(f"ROC AUC: {roc_auc:.4f}, PR AUC: {pr_auc:.4f}")

    # --- ROC Curve ---
    plt.figure(figsize=(6,5))
    for i in range(num_classes):
        fpr, tpr, _ = roc_curve(y_onehot[:,i], y_prob[:,i])
        plt.plot(fpr, tpr, label=f"{class_names[i]} (AUC={roc_auc_score(y_onehot[:,i], y_prob[:,i]):.4f})")
    plt.plot([0,1],[0,1], 'k--')
    plt.title(f"{title} ROC Curve")
    plt.xlabel("FPR")
    plt.ylabel("TPR")
    plt.legend()
    plt.show()

    # --- PR Curve ---
    plt.figure(figsize=(6,5))
    for i in range(num_classes):
        precision, recall, _ = precision_recall_curve(y_onehot[:,i], y_prob[:,i])
        plt.plot(recall, precision, label=f"{class_names[i]} (AP={average_precision_score(y_onehot[:,i], y_prob[:,i]):.4f})")
    plt.title(f"{title} Precision-Recall Curve")
    plt.xlabel("Recall")
    plt.ylabel("Precision")
    plt.legend()
    plt.show()

def evaluate_model(model, loader, title="", plot_curves=False):
    model.eval()
    y_true, y_pred, y_prob = [], [], []
    start_time = time.time()

    with torch.no_grad():
        for imgs, labels in loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            probs = torch.softmax(outputs, dim=1)
            preds = outputs.argmax(1)

            y_true.extend(labels.cpu().numpy())
            y_pred.extend(preds.cpu().numpy())
            y_prob.extend(probs.cpu().numpy())

    infer_time = time.time() - start_time
    y_true, y_pred, y_prob = np.array(y_true), np.array(y_pred), np.array(y_prob)

    print_report(y_true, y_pred, y_prob, title, plot_curves)
    print(f"{title} inference time: {infer_time:.2f} sec")
    return y_true, y_pred, y_prob, infer_time

# =====
# Final Evaluation
# =====
model.load_state_dict(torch.load("swin_tiny_best.pth"))

# Only plot curves for Test Set
y_true_train, y_pred_train, y_prob_train, train_infer_time = evaluate_model(model, train_loader, "Train Set", plot_curves)
y_true_val, y_pred_val, y_prob_val, val_infer_time = evaluate_model(model, val_loader, "Validation Set", plot_curves)
y_true_test, y_pred_test, y_prob_test, test_infer_time = evaluate_model(model, test_loader, "Test Set", plot_curves)

print("\n===== Summary =====")
print(f"Training inference time: {train_infer_time:.2f} sec")
print(f"Validation inference time: {val_infer_time:.2f} sec")
print(f"Test inference time: {test_infer_time:.2f} sec")

model.safetensors:  0%|          | 0.00/114M [00:00<?, ?B/s]
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
  warnings.warn(

```

Epoch [1/35] Train Loss: 0.3674, Train Acc: 0.8656 Val Loss: 0.1095, Val Acc: 0.9643 Time: 29.94s


Epoch [2/35] Train Loss: 0.0638, Train Acc: 0.9784 Val Loss: 0.2573, Val Acc: 0.9196 Time: 28.66s  
Epoch [3/35] Train Loss: 0.0336, Train Acc: 0.9900 Val Loss: 0.0347, Val Acc: 0.9866 Time: 28.73s  
Epoch [4/35] Train Loss: 0.0080, Train Acc: 0.9979 Val Loss: 0.0225, Val Acc: 0.9933 Time: 28.71s  
Epoch [5/35] Train Loss: 0.0078, Train Acc: 0.9973 Val Loss: 0.0603, Val Acc: 0.9821 Time: 28.80s  
Epoch [6/35] Train Loss: 0.0291, Train Acc: 0.9920 Val Loss: 0.0529, Val Acc: 0.9866 Time: 28.74s  
Epoch [7/35] Train Loss: 0.0573, Train Acc: 0.9823 Val Loss: 0.0674, Val Acc: 0.9844 Time: 28.73s  
Epoch [8/35] Train Loss: 0.0159, Train Acc: 0.9956 Val Loss: 0.0308, Val Acc: 0.9955 Time: 28.71s  
Epoch [9/35] Train Loss: 0.0018, Train Acc: 0.9997 Val Loss: 0.0246, Val Acc: 0.9955 Time: 28.71s  
Early stopping triggered!

Training finished Best model saved as swin\_tiny\_best.pth

--- Train Set ---

	precision	recall	f1-score	support
0	1.0000	0.9988	0.9994	839
1	0.9988	1.0000	0.9994	816
2	1.0000	1.0000	1.0000	891
3	1.0000	1.0000	1.0000	839
accuracy			0.9997	3385
macro avg	0.9997	0.9997	0.9997	3385
weighted avg	0.9997	0.9997	0.9997	3385

Train Set Confusion Matrix



MCC: 0.9996

Cohen's Kappa: 0.9996


Mean NPV: 0.9999

Mean PPV (Precision): 0.9997

Train Set inference time: 9.87 sec

--- Validation Set ---

	precision	recall	f1-score	support
0	0.9910	0.9910	0.9910	111
1	0.9907	0.9815	0.9860	108
2	0.9916	1.0000	0.9958	118
3	1.0000	1.0000	1.0000	111
accuracy			0.9933	448
macro avg	0.9933	0.9931	0.9932	448
weighted avg	0.9933	0.9933	0.9933	448



MCC: 0.9911


Cohen's Kappa: 0.9911

Mean NPV: 0.9978



Mean PPV (Precision): 0.9933

Validation Set inference time: 1.49 sec

		precision	recall	f1-score	support
	0	0.9882	0.9941	0.9912	169
	1	0.9879	0.9879	0.9879	165
	2	1.0000	0.9944	0.9972	179
	3	1.0000	1.0000	1.0000	169
		accuracy		0.9941	682
		macro avg	0.9940	0.9941	682
		weighted avg	0.9942	0.9941	682



MCC: 0.9922  
Cohen's Kappa: 0.9922  
Mean NPV: 0.9980  
Mean PPV (Precision): 0.9940  
ROC AUC: 0.9999, PR AUC: 0.9998



Test Set inference time: 2.19 sec

===== Summary =====  
Training inference time: 9.87 sec  
Validation inference time: 1.49 sec  
Test inference time: 2.19 sec

In [7]: # =====

```
# =====
# MobileViT Training + Evaluation Pipeline
# =====

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
from timm import create_model
from torch.optim.lr_scheduler import ReduceLROnPlateau

from sklearn.metrics import (
    classification_report, confusion_matrix, matthews_corrcoef,
    roc_auc_score, average_precision_score, roc_curve,
    precision_recall_curve, cohen_kappa_score
)
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import time

# =====
# Config
# =====
data_dir = "/kaggle/working/augmented_split"
batch_size = 32
num_epochs = 35
patience = 5
num_classes = 4 # Glioma, Meningioma, Pituitary, Normal
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
class_names = ["Glioma", "Meningioma", "Pituitary", "Normal"]

# =====
# Data (no extra augmentation)
# =====
common_tfms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])

train_ds = datasets.ImageFolder(root=f"{data_dir}/train", transform=common_tfms)
val_ds = datasets.ImageFolder(root=f"{data_dir}/val", transform=common_tfms)
test_ds = datasets.ImageFolder(root=f"{data_dir}/test", transform=common_tfms)

train_loader = DataLoader(train_ds, batch_size=batch_size, shuffle=True, num_workers=2)
val_loader = DataLoader(val_ds, batch_size=batch_size, shuffle=False, num_workers=2)
test_loader = DataLoader(test_ds, batch_size=batch_size, shuffle=False, num_workers=2)

# =====
# Model (MobileViT Small)
# Options: "mobilevit_xxs", "mobilevit_xs", "mobilevit_s"
# =====
model = create_model("mobilevit_s", pretrained=True, num_classes=num_classes)
model = model.to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)
scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=True)

# =====
# Training Loop with Early Stopping
# =====
best_val_loss = float("inf")
patience_counter = 0

for epoch in range(num_epochs):
    start_time = time.time()

    # ---- Train ----
    model.train()
    train_loss, train_correct = 0, 0
    for imgs, labels in train_loader:
        imgs, labels = imgs.to(device), labels.to(device)
```

```

optimizer.zero_grad()
outputs = model(imgs)
loss = criterion(outputs, labels)
loss.backward()
optimizer.step()

train_loss += loss.item() * imgs.size(0)
train_correct += (outputs.argmax(1) == labels).sum().item()

train_loss /= len(train_loader.dataset)
train_acc = train_correct / len(train_loader.dataset)

# ---- Validation ----
model.eval()
val_loss, val_correct = 0, 0
with torch.no_grad():
    for imgs, labels in val_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        outputs = model(imgs)
        loss = criterion(outputs, labels)

        val_loss += loss.item() * imgs.size(0)
        val_correct += (outputs.argmax(1) == labels).sum().item()

val_loss /= len(val_loader.dataset)
val_acc = val_correct / len(val_loader.dataset)

scheduler.step(val_loss)

elapsed = time.time() - start_time
print(f"Epoch [{epoch+1}/{num_epochs}] "
      f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f} "
      f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f} "
      f"Time: {elapsed:.2f}s")

# Early Stopping
if val_loss < best_val_loss:
    best_val_loss = val_loss
    patience_counter = 0
    torch.save(model.state_dict(), "mobilevit_best.pth")
else:
    patience_counter += 1
    if patience_counter >= patience:
        print("Early stopping triggered!")
        break

print("Training finished Best model saved as mobilevit_best.pth")

# =====
# Evaluation Functions
# =====
def plot_confusion_matrix(cm, classes, title="Confusion Matrix"):
    plt.figure(figsize=(6,5))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=classes, yticklabels=classes)
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title(title)
    plt.show()

def print_report(y_true, y_pred, y_prob, title="", plot_curves=False):
    print(f"\n--- {title} ---")
    print(classification_report(y_true, y_pred, digits=4))

    # Confusion Matrix
    cm = confusion_matrix(y_true, y_pred)
    plot_confusion_matrix(cm, class_names, title=f"{title} Confusion Matrix")

    # MCC
    mcc = matthews_corrcoef(y_true, y_pred)
    print(f"MCC: {mcc:.4f}")

    # Cohen's Kappa
    kappa = cohen_kappa_score(y_true, y_pred)
    print(f"Cohen's Kappa: {kappa:.4f}")

    # Class-wise NPV + PPV
    npv_list, ppv_list = [], []
    for i in range(len(cm)):
        TN = np.sum(np.delete(np.delete(cm, i, axis=0), i, axis=1))
        FN = np.sum(cm[i, :]) - cm[i, i]
        FP = np.sum(cm[:, i]) - cm[i, i]
        TP = cm[i, i]

```

```

NPV = TN / (TN + FN) if (TN + FN) > 0 else 0
PPV = TP / (TP + FP) if (TP + FP) > 0 else 0
npv_list.append(NPV)
ppv_list.append(PPV)

print(f"Mean NPV: {np.mean(npv_list):.4f}")
print(f"Mean PPV (Precision): {np.mean(ppv_list):.4f}")

if plot_curves:
    # ROC AUC + PR AUC
    y_onehot = np.eye(num_classes)[y_true]
    roc_auc = roc_auc_score(y_onehot, y_prob, average='macro', multi_class='ovr')
    pr_auc = average_precision_score(y_onehot, y_prob, average='macro')
    print(f"ROC AUC: {roc_auc:.4f}, PR AUC: {pr_auc:.4f}")

    # --- ROC Curve ---
    plt.figure(figsize=(6,5))
    for i in range(num_classes):
        fpr, tpr, _ = roc_curve(y_onehot[:,i], y_prob[:,i])
        plt.plot(fpr, tpr, label=f"{class_names[i]} (AUC={roc_auc_score(y_onehot[:,i], y_prob[:,i]):.4f})")
    plt.plot([0,1],[0,1], 'k--')
    plt.title(f"{title} ROC Curve")
    plt.xlabel("FPR")
    plt.ylabel("TPR")
    plt.legend()
    plt.show()

    # --- PR Curve ---
    plt.figure(figsize=(6,5))
    for i in range(num_classes):
        precision, recall, _ = precision_recall_curve(y_onehot[:,i], y_prob[:,i])
        plt.plot(recall, precision, label=f"{class_names[i]} (AP={average_precision_score(y_onehot[:,i], y_prob[:,i]):.4f})")
    plt.title(f"{title} Precision-Recall Curve")
    plt.xlabel("Recall")
    plt.ylabel("Precision")
    plt.legend()
    plt.show()

def evaluate_model(model, loader, title="", plot_curves=False):
    model.eval()
    y_true, y_pred, y_prob = [], [], []
    start_time = time.time()

    with torch.no_grad():
        for imgs, labels in loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            probs = torch.softmax(outputs, dim=1)
            preds = outputs.argmax(1)

            y_true.extend(labels.cpu().numpy())
            y_pred.extend(preds.cpu().numpy())
            y_prob.extend(probs.cpu().numpy())

    infer_time = time.time() - start_time
    y_true, y_pred, y_prob = np.array(y_true), np.array(y_pred), np.array(y_prob)

    print_report(y_true, y_pred, y_prob, title, plot_curves)
    print(f"{title} inference time: {infer_time:.2f} sec")
    return y_true, y_pred, y_prob, infer_time

# =====
# Final Evaluation
# =====
model.load_state_dict(torch.load("mobilevit_best.pth"))

# Only plot curves for Test Set
y_true_train, y_pred_train, y_prob_train, train_infer_time = evaluate_model(model, train_loader, "Train Set", plot_curves)
y_true_val, y_pred_val, y_prob_val, val_infer_time = evaluate_model(model, val_loader, "Validation Set", plot_curves)
y_true_test, y_pred_test, y_prob_test, test_infer_time = evaluate_model(model, test_loader, "Test Set", plot_curves)

print("\n===== Summary =====")
print(f"Training inference time: {train_infer_time:.2f} sec")
print(f"Validation inference time: {val_infer_time:.2f} sec")
print(f"Test inference time: {test_infer_time:.2f} sec")

model.safetensors: 0%|          | 0.00/22.4M [00:00<?, ?B/s]
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
  warnings.warn(

```

Epoch [1/35] Train Loss: 0.8362, Train Acc: 0.8399 Val Loss: 0.3077, Val Acc: 0.9353 Time: 21.01s


Epoch [2/35] Train Loss: 0.2154, Train Acc: 0.9533 Val Loss: 0.0899, Val Acc: 0.9844 Time: 20.64s  
Epoch [3/35] Train Loss: 0.1087, Train Acc: 0.9758 Val Loss: 0.0689, Val Acc: 0.9844 Time: 20.76s  
Epoch [4/35] Train Loss: 0.0678, Train Acc: 0.9861 Val Loss: 0.0349, Val Acc: 0.9933 Time: 20.66s  
Epoch [5/35] Train Loss: 0.0437, Train Acc: 0.9911 Val Loss: 0.0398, Val Acc: 0.9888 Time: 20.64s  
Epoch [6/35] Train Loss: 0.0348, Train Acc: 0.9911 Val Loss: 0.0734, Val Acc: 0.9754 Time: 20.65s  
Epoch [7/35] Train Loss: 0.0249, Train Acc: 0.9938 Val Loss: 0.1708, Val Acc: 0.9487 Time: 20.59s  
Epoch [8/35] Train Loss: 0.0225, Train Acc: 0.9950 Val Loss: 0.0462, Val Acc: 0.9888 Time: 20.74s  
Epoch [9/35] Train Loss: 0.0179, Train Acc: 0.9953 Val Loss: 0.0310, Val Acc: 0.9888 Time: 20.68s  
Epoch [10/35] Train Loss: 0.0172, Train Acc: 0.9968 Val Loss: 0.0217, Val Acc: 0.9933 Time: 20.82s  
Epoch [11/35] Train Loss: 0.0124, Train Acc: 0.9976 Val Loss: 0.0239, Val Acc: 0.9933 Time: 20.63s  
Epoch [12/35] Train Loss: 0.0108, Train Acc: 0.9970 Val Loss: 0.0211, Val Acc: 0.9955 Time: 20.76s  
Epoch [13/35] Train Loss: 0.0108, Train Acc: 0.9970 Val Loss: 0.0306, Val Acc: 0.9911 Time: 20.83s  
Epoch [14/35] Train Loss: 0.0115, Train Acc: 0.9982 Val Loss: 0.0380, Val Acc: 0.9911 Time: 20.73s  
Epoch [15/35] Train Loss: 0.0063, Train Acc: 0.9988 Val Loss: 0.0372, Val Acc: 0.9933 Time: 20.74s  
Epoch [16/35] Train Loss: 0.0068, Train Acc: 0.9985 Val Loss: 0.0384, Val Acc: 0.9911 Time: 20.74s  
Epoch [17/35] Train Loss: 0.0161, Train Acc: 0.9970 Val Loss: 0.0283, Val Acc: 0.9933 Time: 20.73s  
Early stopping triggered!

Training finished Best model saved as mobilevit\_best.pth

--- Train Set ---

	precision	recall	f1-score	support
0	1.0000	0.9988	0.9994	839
1	0.9988	1.0000	0.9994	816
2	1.0000	1.0000	1.0000	891
3	1.0000	1.0000	1.0000	839
accuracy			0.9997	3385
macro avg	0.9997	0.9997	0.9997	3385
weighted avg	0.9997	0.9997	0.9997	3385

Train Set Confusion Matrix



MCC: 0.9996

Cohen's Kappa: 0.9996


Mean NPV: 0.9999

Mean PPV (Precision): 0.9997

Train Set inference time: 6.48 sec

--- Validation Set ---

	precision	recall	f1-score	support
0	1.0000	0.9820	0.9909	111
1	0.9818	1.0000	0.9908	108
2	1.0000	1.0000	1.0000	118
3	1.0000	1.0000	1.0000	111
accuracy			0.9955	448
macro avg	0.9955	0.9955	0.9954	448
weighted avg	0.9956	0.9955	0.9955	448



MCC: 0.9941

Cohen's Kappa: 0.9940


Mean NPV: 0.9985

Mean PPV (Precision): 0.9955

Validation Set inference time: 1.03 sec

--- Test Set ---

	precision	recall	f1-score	support
0	1.0000	0.9882	0.9940	169
1	0.9880	1.0000	0.9940	165
2	1.0000	1.0000	1.0000	179
3	1.0000	1.0000	1.0000	169
accuracy			0.9971	682
macro avg	0.9970	0.9970	0.9970	682
weighted avg	0.9971	0.9971	0.9971	682





MCC: 0.9961

Cohen's Kappa: 0.9961

Mean NPV: 0.9990

Mean PPV (Precision): 0.9970

ROC AUC: 1.0000, PR AUC: 1.0000



Test Set inference time: 1.56 sec

```
===== Summary =====
Training inference time: 6.48 sec
Validation inference time: 1.03 sec
Test inference time: 1.56 sec
```

In [8]: # =====

```
# TinyViT Training + Evaluation Pipeline
# =====

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
from timm import create_model
from torch.optim.lr_scheduler import ReduceLROnPlateau

from sklearn.metrics import (
    classification_report, confusion_matrix, matthews_corrcoef,
    roc_auc_score, average_precision_score, roc_curve,
    precision_recall_curve, cohen_kappa_score
)
import numpy as np
import matplotlib.pyplot as plt
import time
```

```

# =====
# Config
# =====
data_dir = "/kaggle/working/augmented_split"
batch_size = 32
num_epochs = 35
patience = 5
num_classes = 4 # Glioma, Meningioma, Pituitary, Normal
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# =====
# Data (same as before, no new augmentation)
# =====
common_tfms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])

train_ds = datasets.ImageFolder(root=f"{data_dir}/train", transform=common_tfms)
val_ds = datasets.ImageFolder(root=f"{data_dir}/val", transform=common_tfms)
test_ds = datasets.ImageFolder(root=f"{data_dir}/test", transform=common_tfms)

train_loader = DataLoader(train_ds, batch_size=batch_size, shuffle=True, num_workers=2)
val_loader = DataLoader(val_ds, batch_size=batch_size, shuffle=False, num_workers=2)
test_loader = DataLoader(test_ds, batch_size=batch_size, shuffle=False, num_workers=2)

# =====
# Model (TinyViT)
# Options: "tiny_vit_5m_224", "tiny_vit_11m_224", "tiny_vit_21m_224"
# =====
model = create_model("tiny_vit_5m_224", pretrained=True, num_classes=num_classes)
model = model.to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)
scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=True)

# =====
# Training Loop with Early Stopping
# =====
best_val_loss = float("inf")
patience_counter = 0

for epoch in range(num_epochs):
    start_time = time.time()

    # ---- Train ----
    model.train()
    train_loss, train_correct = 0, 0
    for imgs, labels in train_loader:
        imgs, labels = imgs.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(imgs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        train_loss += loss.item() * imgs.size(0)
        train_correct += (outputs.argmax(1) == labels).sum().item()

    train_loss /= len(train_loader.dataset)
    train_acc = train_correct / len(train_loader.dataset)

    # ---- Validation ----
    model.eval()
    val_loss, val_correct = 0, 0
    with torch.no_grad():
        for imgs, labels in val_loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            loss = criterion(outputs, labels)

            val_loss += loss.item() * imgs.size(0)
            val_correct += (outputs.argmax(1) == labels).sum().item()

    val_loss /= len(val_loader.dataset)
    val_acc = val_correct / len(val_loader.dataset)

    scheduler.step(val_loss)

    elapsed = time.time() - start_time

```

```

print(f"Epoch [{epoch+1}/{num_epochs}] "
      f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f} "
      f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f} "
      f"Time: {elapsed:.2f}s")

# Early Stopping
if val_loss < best_val_loss:
    best_val_loss = val_loss
    patience_counter = 0
    torch.save(model.state_dict(), "tinyvit_best.pth")
else:
    patience_counter += 1
    if patience_counter >= patience:
        print("Early stopping triggered!")
        break

print("Training finished Best model saved as tinyvit_best.pth")

# =====
# Evaluation Functions
# =====
def plot_confusion_matrix(cm, classes, title="Confusion Matrix"):
    plt.figure(figsize=(6,5))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=classes, yticklabels=classes)
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title(title)
    plt.show()

def print_report(y_true, y_pred, y_prob, title="", plot_curves=False):
    print(f"\n--- {title} ---")
    print(classification_report(y_true, y_pred, digits=4))

    # Confusion Matrix
    cm = confusion_matrix(y_true, y_pred)
    plot_confusion_matrix(cm, class_names, title=f"{title} Confusion Matrix")

    # MCC
    mcc = matthews_corrcoef(y_true, y_pred)
    print(f"MCC: {mcc:.4f}")

    # Cohen's Kappa
    kappa = cohen_kappa_score(y_true, y_pred)
    print(f"Cohen's Kappa: {kappa:.4f}")

    # Class-wise NPV + PPV
    npv_list, ppv_list = [], []
    for i in range(len(cm)):
        TN = np.sum(np.delete(np.delete(cm, i, axis=0), i, axis=1))
        FN = np.sum(cm[i, :]) - cm[i, i]
        FP = np.sum(cm[:, i]) - cm[i, i]
        TP = cm[i, i]

        NPV = TN / (TN + FN) if (TN + FN) > 0 else 0
        PPV = TP / (TP + FP) if (TP + FP) > 0 else 0
        npv_list.append(NPV)
        ppv_list.append(PPV)

    print(f"Mean NPV: {np.mean(npv_list):.4f}")
    print(f"Mean PPV (Precision): {np.mean(ppv_list):.4f}")

    if plot_curves:
        # ROC AUC + PR AUC
        y_onehot = np.eye(num_classes)[y_true]
        roc_auc = roc_auc_score(y_onehot, y_prob, average='macro', multi_class='ovr')
        pr_auc = average_precision_score(y_onehot, y_prob, average='macro')
        print(f"ROC AUC: {roc_auc:.4f}, PR AUC: {pr_auc:.4f}")

        # --- ROC Curve ---
        plt.figure(figsize=(6,5))
        for i in range(num_classes):
            fpr, tpr, _ = roc_curve(y_onehot[:,i], y_prob[:,i])
            plt.plot(fpr, tpr, label=f"{class_names[i]} (AUC={roc_auc_score(y_onehot[:,i], y_prob[:,i]):.4f})")
        plt.plot([0,1],[0,1], 'k--')
        plt.title(f"{title} ROC Curve")
        plt.xlabel("FPR")
        plt.ylabel("TPR")
        plt.legend()
        plt.show()

        # --- PR Curve ---
        plt.figure(figsize=(6,5))

```

```

    for i in range(num_classes):
        precision, recall, _ = precision_recall_curve(y_onehot[:,i], y_prob[:,i])
        plt.plot(recall, precision, label=f"{{class_names[i]}} (AP={average_precision_score(y_onehot[:,i], y_p
    plt.title(f"{{title}} Precision-Recall Curve")
    plt.xlabel("Recall")
    plt.ylabel("Precision")
    plt.legend()
    plt.show()

def evaluate_model(model, loader, title="", plot_curves=False):
    model.eval()
    y_true, y_pred, y_prob = [], [], []
    start_time = time.time()

    with torch.no_grad():
        for imgs, labels in loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            probs = torch.softmax(outputs, dim=1)
            preds = outputs.argmax(1)

            y_true.extend(labels.cpu().numpy())
            y_pred.extend(preds.cpu().numpy())
            y_prob.extend(probs.cpu().numpy())

    infer_time = time.time() - start_time
    y_true, y_pred, y_prob = np.array(y_true), np.array(y_pred), np.array(y_prob)

    print_report(y_true, y_pred, y_prob, title, plot_curves)
    print(f"{{title}} inference time: {infer_time:.2f} sec")
    return y_true, y_pred, y_prob, infer_time

# =====
# Final Evaluation
# =====
model.load_state_dict(torch.load("tinyvit_best.pth"))

# Only plot curves for Test Set
y_true_train, y_pred_train, y_prob_train, train_infer_time = evaluate_model(model, train_loader, "Train Set", p
y_true_val, y_pred_val, y_prob_val, val_infer_time = evaluate_model(model, val_loader, "Validation Set",
y_true_test, y_pred_test, y_prob_test, test_infer_time = evaluate_model(model, test_loader, "Test Set", plot

print("\n===== Summary =====")
print(f"Training inference time: {train_infer_time:.2f} sec")
print(f"Validation inference time: {val_infer_time:.2f} sec")
print(f"Test inference time: {test_infer_time:.2f} sec")

model.safetensors: 0% | 0.00/48.4M [00:00<?, ?B/s]
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
precated. Please use get_last_lr() to access the learning rate.
warnings.warn(

```


Epoch [1/35] Train Loss: 0.5656, Train Acc: 0.8331 Val Loss: 0.1971, Val Acc: 0.9621 Time: 16.38s  
 Epoch [2/35] Train Loss: 0.1211, Train Acc: 0.9761 Val Loss: 0.0870, Val Acc: 0.9821 Time: 16.20s  
 Epoch [3/35] Train Loss: 0.0416, Train Acc: 0.9953 Val Loss: 0.0353, Val Acc: 0.9978 Time: 16.26s  
 Epoch [4/35] Train Loss: 0.0221, Train Acc: 0.9973 Val Loss: 0.0331, Val Acc: 0.9955 Time: 16.16s  
 Epoch [5/35] Train Loss: 0.0128, Train Acc: 0.9988 Val Loss: 0.0361, Val Acc: 0.9933 Time: 16.33s  
 Epoch [6/35] Train Loss: 0.0105, Train Acc: 0.9991 Val Loss: 0.0283, Val Acc: 0.9955 Time: 16.17s  
 Epoch [7/35] Train Loss: 0.0132, Train Acc: 0.9973 Val Loss: 0.0208, Val Acc: 0.9978 Time: 16.18s  
 Epoch [8/35] Train Loss: 0.0134, Train Acc: 0.9973 Val Loss: 0.0325, Val Acc: 0.9933 Time: 16.21s  
 Epoch [9/35] Train Loss: 0.0245, Train Acc: 0.9905 Val Loss: 0.0961, Val Acc: 0.9732 Time: 16.30s  
 Epoch [10/35] Train Loss: 0.0120, Train Acc: 0.9976 Val Loss: 0.0222, Val Acc: 0.9978 Time: 16.04s  
 Epoch [11/35] Train Loss: 0.0071, Train Acc: 0.9988 Val Loss: 0.0218, Val Acc: 0.9955 Time: 16.13s  
 Epoch [12/35] Train Loss: 0.0040, Train Acc: 0.9997 Val Loss: 0.0190, Val Acc: 0.9978 Time: 16.24s  
 Epoch [13/35] Train Loss: 0.0046, Train Acc: 0.9991 Val Loss: 0.0215, Val Acc: 0.9955 Time: 16.17s  
 Epoch [14/35] Train Loss: 0.0024, Train Acc: 1.0000 Val Loss: 0.0210, Val Acc: 0.9955 Time: 16.10s  
 Epoch [15/35] Train Loss: 0.0021, Train Acc: 1.0000 Val Loss: 0.0202, Val Acc: 0.9978 Time: 16.11s  
 Epoch [16/35] Train Loss: 0.0019, Train Acc: 1.0000 Val Loss: 0.0212, Val Acc: 0.9955 Time: 16.18s  
 Epoch [17/35] Train Loss: 0.0018, Train Acc: 1.0000 Val Loss: 0.0196, Val Acc: 0.9978 Time: 16.08s  
 Early stopping triggered!

Training finished Best model saved as tinyvit\_best.pth

--- Train Set ---

	precision	recall	f1-score	support
0	1.0000	1.0000	1.0000	839
1	1.0000	1.0000	1.0000	816
2	1.0000	1.0000	1.0000	891
3	1.0000	1.0000	1.0000	839
accuracy			1.0000	3385
macro avg	1.0000	1.0000	1.0000	3385
weighted avg	1.0000	1.0000	1.0000	3385

Train Set Confusion Matrix



MCC: 1.0000

Cohen's Kappa: 1.0000


Mean NPV: 1.0000

Mean PPV (Precision): 1.0000

Train Set inference time: 5.46 sec

--- Validation Set ---

	precision	recall	f1-score	support
0	1.0000	0.9910	0.9955	111
1	0.9908	1.0000	0.9954	108
2	1.0000	1.0000	1.0000	118
3	1.0000	1.0000	1.0000	111
accuracy			0.9978	448
macro avg	0.9977	0.9977	0.9977	448
weighted avg	0.9978	0.9978	0.9978	448



MCC: 0.9970

Cohen's Kappa: 0.9970


Mean NPV: 0.9993

Mean PPV (Precision): 0.9977

Validation Set inference time: 0.90 sec

--- Test Set ---

	precision	recall	f1-score	support
0	1.0000	0.9882	0.9940	169
1	0.9880	1.0000	0.9940	165
2	1.0000	1.0000	1.0000	179
3	1.0000	1.0000	1.0000	169
accuracy			0.9971	682
macro avg	0.9970	0.9970	0.9970	682
weighted avg	0.9971	0.9971	0.9971	682





MCC: 0.9961

Cohen's Kappa: 0.9961

Mean NPV: 0.9990

Mean PPV (Precision): 0.9970

ROC AUC: 1.0000, PR AUC: 1.0000



Test Set inference time: 1.27 sec

```
===== Summary =====
Training inference time: 5.46 sec
Validation inference time: 0.90 sec
Test inference time: 1.27 sec
```

In [11]: # ======  
# LeViT Training + Evaluation Pipeline  
# ======

```
import os
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
from timm import create_model
from torch.optim.lr_scheduler import ReduceLROnPlateau

from sklearn.metrics import (
    classification_report, confusion_matrix, matthews_corrcoef,
    roc_auc_score, average_precision_score, roc_curve,
    precision_recall_curve, cohen_kappa_score
)
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```

import time

# =====
# Config
# =====
data_dir = "/kaggle/working/augmented_split"
batch_size = 32
num_epochs = 35
patience = 5
num_classes = 4 # Glioma, Meningioma, Pituitary, Normal
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
class_names = ["Glioma", "Meningioma", "Pituitary", "Normal"]

# =====
# Data (no extra augmentation)
# =====
common_tfms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])

train_ds = datasets.ImageFolder(root=f"{data_dir}/train", transform=common_tfms)
val_ds = datasets.ImageFolder(root=f"{data_dir}/val", transform=common_tfms)
test_ds = datasets.ImageFolder(root=f"{data_dir}/test", transform=common_tfms)

train_loader = DataLoader(train_ds, batch_size=batch_size, shuffle=True, num_workers=2)
val_loader = DataLoader(val_ds, batch_size=batch_size, shuffle=False, num_workers=2)
test_loader = DataLoader(test_ds, batch_size=batch_size, shuffle=False, num_workers=2)

# =====
# Model (LeViT)
# Options: "levit_128s", "levit_128", "levit_192", "levit_256"
# =====
model = create_model("levit_128s", pretrained=True, num_classes=num_classes)
model = model.to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)
scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=True)

# =====
# Training Loop with Early Stopping
# =====
best_val_loss = float("inf")
patience_counter = 0

for epoch in range(num_epochs):
    start_time = time.time()

    # ---- Train ----
    model.train()
    train_loss, train_correct = 0, 0
    for imgs, labels in train_loader:
        imgs, labels = imgs.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(imgs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        train_loss += loss.item() * imgs.size(0)
        train_correct += (outputs.argmax(1) == labels).sum().item()

    train_loss /= len(train_loader.dataset)
    train_acc = train_correct / len(train_loader.dataset)

    # ---- Validation ----
    model.eval()
    val_loss, val_correct = 0, 0
    with torch.no_grad():
        for imgs, labels in val_loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            loss = criterion(outputs, labels)

            val_loss += loss.item() * imgs.size(0)
            val_correct += (outputs.argmax(1) == labels).sum().item()

    val_loss /= len(val_loader.dataset)
    val_acc = val_correct / len(val_loader.dataset)

    if val_loss < best_val_loss:
        best_val_loss = val_loss
        patience_counter = 0
    else:
        patience_counter += 1

    if patience_counter == patience:
        print(f"Early stopping at epoch {epoch} due to no improvement in validation loss for {patience} epochs")
        break

```

```

scheduler.step(val_loss)

elapsed = time.time() - start_time
print(f"Epoch [{epoch+1}/{num_epochs}] "
      f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f} "
      f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f} "
      f"Time: {elapsed:.2f}s")

# Early Stopping
if val_loss < best_val_loss:
    best_val_loss = val_loss
    patience_counter = 0
    torch.save(model.state_dict(), "levit_best.pth")
else:
    patience_counter += 1
    if patience_counter >= patience:
        print("Early stopping triggered!")
        break

print("Training finished Best model saved as levit_best.pth")

# =====
# Evaluation Functions
# =====

def plot_confusion_matrix(cm, classes, title="Confusion Matrix"):
    plt.figure(figsize=(6,5))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=classes, yticklabels=classes)
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title(title)
    plt.show()

def print_report(y_true, y_pred, y_prob, title="", plot_curves=False):
    print(f"\n--- {title} ---")
    print(classification_report(y_true, y_pred, target_names=class_names, digits=4))

    # Confusion Matrix
    cm = confusion_matrix(y_true, y_pred)
    plot_confusion_matrix(cm, class_names, title=f"{title} Confusion Matrix")

    # MCC
    mcc = matthews_corrcoef(y_true, y_pred)
    print(f"MCC: {mcc:.4f}")

    # Cohen's Kappa
    kappa = cohen_kappa_score(y_true, y_pred)
    print(f"Cohen's Kappa: {kappa:.4f}")

    # Class-wise NPV + PPV
    npv_list, ppv_list = [], []
    for i in range(len(cm)):
        TN = np.sum(np.delete(np.delete(cm, i, axis=0), i, axis=1))
        FN = np.sum(cm[i, :]) - cm[i, i]
        FP = np.sum(cm[:, i]) - cm[i, i]
        TP = cm[i, i]

        NPV = TN / (TN + FN) if (TN + FN) > 0 else 0
        PPV = TP / (TP + FP) if (TP + FP) > 0 else 0
        npv_list.append(NPV)
        ppv_list.append(PPV)

    print(f"Mean NPV: {np.mean(npv_list):.4f}")
    print(f"Mean PPV (Precision): {np.mean(ppv_list):.4f}")

    if plot_curves:
        # ROC AUC + PR AUC
        y_onehot = np.eye(num_classes)[y_true]
        roc_auc = roc_auc_score(y_onehot, y_prob, average='macro', multi_class='ovr')
        pr_auc = average_precision_score(y_onehot, y_prob, average='macro')
        print(f"ROC AUC: {roc_auc:.4f}, PR AUC: {pr_auc:.4f}")

        # --- ROC Curve ---
        plt.figure(figsize=(6,5))
        for i in range(num_classes):
            fpr, tpr, _ = roc_curve(y_onehot[:,i], y_prob[:,i])
            plt.plot(fpr, tpr, label=f"{class_names[i]} (AUC={roc_auc_score(y_onehot[:,i], y_prob[:,i]):.4f})")
        plt.plot([0,1],[0,1],'k--')
        plt.title(f"{title} ROC Curve")
        plt.xlabel("FPR")
        plt.ylabel("TPR")
        plt.legend()
        plt.show()

```

```

# --- PR Curve ---
plt.figure(figsize=(6,5))
for i in range(num_classes):
    precision, recall, _ = precision_recall_curve(y_onehot[:,i], y_prob[:,i])
    plt.plot(recall, precision, label=f'{class_names[i]} (AP={average_precision_score(y_onehot[:,i], y_prob[:,i])})')
plt.title(f'{title} Precision-Recall Curve')
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.legend()
plt.show()

def evaluate_model(model, loader, title="", plot_curves=False):
    model.eval()
    y_true, y_pred, y_prob = [], [], []
    start_time = time.time()

    with torch.no_grad():
        for imgs, labels in loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            probs = torch.softmax(outputs, dim=1)
            preds = outputs.argmax(1)

            y_true.extend(labels.cpu().numpy())
            y_pred.extend(preds.cpu().numpy())
            y_prob.extend(probs.cpu().numpy())

    infer_time = time.time() - start_time
    y_true, y_pred, y_prob = np.array(y_true), np.array(y_pred), np.array(y_prob)

    print_report(y_true, y_pred, y_prob, title, plot_curves)
    print(f'{title} inference time: {infer_time:.2f} sec")
    return y_true, y_pred, y_prob, infer_time

# =====
# Final Evaluation
# =====
model.load_state_dict(torch.load("levit_best.pth"))

# Only plot curves for Test Set
y_true_train, y_pred_train, y_prob_train, train_infer_time = evaluate_model(model, train_loader, "Train Set", plot_curves=False)
y_true_val, y_pred_val, y_prob_val, val_infer_time = evaluate_model(model, val_loader, "Validation Set", plot_curves=False)
y_true_test, y_pred_test, y_prob_test, test_infer_time = evaluate_model(model, test_loader, "Test Set", plot_curves=False)

print("\n===== Summary =====")
print(f"Training inference time: {train_infer_time:.2f} sec")
print(f"Validation inference time: {val_infer_time:.2f} sec")
print(f"Test inference time: {test_infer_time:.2f} sec")

```

model.safetensors: 0% | 0.00/31.3M [00:00<?, ?B/s]

/usr/local/lib/python3.11/dist-packages/torch/optim/lr\_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get\_last\_lr() to access the learning rate.  
 warnings.warn(


Epoch [1/35] Train Loss: 0.9599, Train Acc: 0.7288 Val Loss: 0.5414, Val Acc: 0.8661 Time: 6.76s  
 Epoch [2/35] Train Loss: 0.4157, Train Acc: 0.9019 Val Loss: 1.0715, Val Acc: 0.9286 Time: 6.76s  
 Epoch [3/35] Train Loss: 0.2169, Train Acc: 0.9518 Val Loss: 0.1574, Val Acc: 0.9665 Time: 7.51s  
 Epoch [4/35] Train Loss: 0.1163, Train Acc: 0.9808 Val Loss: 0.1213, Val Acc: 0.9732 Time: 7.53s  
 Epoch [5/35] Train Loss: 0.0676, Train Acc: 0.9935 Val Loss: 0.0763, Val Acc: 0.9866 Time: 7.55s  
 Epoch [6/35] Train Loss: 0.0447, Train Acc: 0.9950 Val Loss: 0.0612, Val Acc: 0.9911 Time: 7.40s  
 Epoch [7/35] Train Loss: 0.0341, Train Acc: 0.9962 Val Loss: 0.0730, Val Acc: 0.9777 Time: 7.75s  
 Epoch [8/35] Train Loss: 0.0295, Train Acc: 0.9970 Val Loss: 0.0783, Val Acc: 0.9844 Time: 7.35s  
 Epoch [9/35] Train Loss: 0.0197, Train Acc: 0.9979 Val Loss: 0.0656, Val Acc: 0.9866 Time: 7.46s  
 Epoch [10/35] Train Loss: 0.0209, Train Acc: 0.9965 Val Loss: 0.0349, Val Acc: 0.9866 Time: 7.40s  
 Epoch [11/35] Train Loss: 0.0144, Train Acc: 0.9988 Val Loss: 0.0428, Val Acc: 0.9866 Time: 6.76s  
 Epoch [12/35] Train Loss: 0.0132, Train Acc: 0.9985 Val Loss: 0.0442, Val Acc: 0.9844 Time: 7.07s  
 Epoch [13/35] Train Loss: 0.0132, Train Acc: 0.9985 Val Loss: 0.0426, Val Acc: 0.9888 Time: 7.29s  
 Epoch [14/35] Train Loss: 0.0118, Train Acc: 0.9988 Val Loss: 0.0305, Val Acc: 0.9955 Time: 7.09s  
 Epoch [15/35] Train Loss: 0.0082, Train Acc: 0.9994 Val Loss: 0.0372, Val Acc: 0.9888 Time: 7.23s  
 Epoch [16/35] Train Loss: 0.0086, Train Acc: 0.9997 Val Loss: 0.0428, Val Acc: 0.9866 Time: 7.46s  
 Epoch [17/35] Train Loss: 0.0104, Train Acc: 0.9988 Val Loss: 0.0333, Val Acc: 0.9888 Time: 7.01s  
 Epoch [18/35] Train Loss: 0.0078, Train Acc: 1.0000 Val Loss: 0.0389, Val Acc: 0.9866 Time: 6.78s  
 Epoch [19/35] Train Loss: 0.0089, Train Acc: 0.9994 Val Loss: 0.0455, Val Acc: 0.9866 Time: 7.12s  
 Early stopping triggered!

Training finished Best model saved as levit\_best.pth

--- Train Set ---

	precision	recall	f1-score	support
Glioma	1.0000	1.0000	1.0000	839
Meningioma	1.0000	1.0000	1.0000	816
Pituitary	1.0000	0.9978	0.9989	891
Normal	0.9976	1.0000	0.9988	839
accuracy			0.9994	3385
macro avg	0.9994	0.9994	0.9994	3385
weighted avg	0.9994	0.9994	0.9994	3385

Train Set Confusion Matrix



MCC: 0.9992

Cohen's Kappa: 0.9992


Mean NPV: 0.9998

Mean PPV (Precision): 0.9994

Train Set inference time: 4.81 sec


--- Validation Set ---

	precision	recall	f1-score	support
Glioma	1.0000	0.9820	0.9909	111
Meningioma	0.9908	1.0000	0.9954	108
Pituitary	1.0000	1.0000	1.0000	118
Normal	0.9911	1.0000	0.9955	111
accuracy			0.9955	448
macro avg	0.9955	0.9955	0.9955	448
weighted avg	0.9956	0.9955	0.9955	448





MCC: 0.9941  
 Cohen's Kappa: 0.9940  
 Mean NPV: 0.9985  
 Mean PPV (Precision): 0.9955  
 Validation Set inference time: 0.78 sec

--- Test Set ---				
	precision	recall	f1-score	support
Glioma	1.0000	0.9822	0.9910	169
Meningioma	0.9821	1.0000	0.9910	165
Pituitary	1.0000	1.0000	1.0000	179
Normal	1.0000	1.0000	1.0000	169
accuracy			0.9956	682
macro avg	0.9955	0.9956	0.9955	682
weighted avg	0.9957	0.9956	0.9956	682



MCC: 0.9942  
 Cohen's Kappa: 0.9941  
 Mean NPV: 0.9985  
 Mean PPV (Precision): 0.9955  
 ROC AUC: 0.9999, PR AUC: 0.9998



Test Set inference time: 1.09 sec

```
===== Summary =====
Training inference time: 4.81 sec
Validation inference time: 0.78 sec
Test inference time: 1.09 sec
```

In [39]:

```
# =====
# Full TinyViT Training + Grad-CAM Visualization
# =====

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
from timm import create_model
from torch.optim.lr_scheduler import ReduceLROnPlateau

import numpy as np
import matplotlib.pyplot as plt
import cv2
import time

# =====
# Grad-CAM Class
```

```

# =====
class GradCAM:
    def __init__(self, model, target_layer):
        self.model = model
        self.target_layer = target_layer
        self.gradients = None
        self.activations = None
        self._register_hooks()

    def _register_hooks(self):
        def forward_hook(module, input, output):
            self.activations = output.detach()
        def backward_hook(module, grad_input, grad_output):
            self.gradients = grad_output[0].detach()

        self.target_layer.register_forward_hook(forward_hook)
        self.target_layer.register_full_backward_hook(backward_hook)

    def __call__(self, x, class_idx=None):
        self.model.eval()
        x = x.to(next(self.model.parameters()).device)
        output = self.model(x)

        if class_idx is None:
            class_idx = output.argmax(dim=1)[0]

        self.model.zero_grad()
        loss = output[0, class_idx]
        loss.backward(retain_graph=True)

        weights = self.gradients.mean(dim=(2, 3), keepdim=True)
        cam = (weights * self.activations).sum(dim=1, keepdim=True)
        cam = torch.nn.functional.relu(cam)
        cam = torch.nn.functional.interpolate(cam, size=(x.size(2), x.size(3)), mode='bilinear', align_corners=False)
        cam = cam.squeeze().cpu().numpy()
        cam = (cam - cam.min()) / (cam.max() - cam.min() + 1e-8)
        return cam

# =====
# Config
# =====
data_dir = "/kaggle/working/augmented_split"
batch_size = 32
num_epochs = 35
patience = 5
num_classes = 4 # Glioma, Meningioma, Pituitary, Normal
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# =====
# Data
# =====
common_tfms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])

train_ds = datasets.ImageFolder(root=f"{data_dir}/train", transform=common_tfms)
val_ds = datasets.ImageFolder(root=f"{data_dir}/val", transform=common_tfms)
test_ds = datasets.ImageFolder(root=f"{data_dir}/test", transform=common_tfms)

train_loader = DataLoader(train_ds, batch_size=batch_size, shuffle=True, num_workers=2)
val_loader = DataLoader(val_ds, batch_size=batch_size, shuffle=False, num_workers=2)
test_loader = DataLoader(test_ds, batch_size=batch_size, shuffle=False, num_workers=2)

# =====
# Model
# =====
model = create_model("tiny_vit_5m_224", pretrained=True, num_classes=num_classes)
model = model.to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)
scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=True)

# =====
# Training Loop with Early Stopping
# =====
best_val_loss = float("inf")
patience_counter = 0

for epoch in range(num_epochs):
    start_time = time.time()

```

```

# ---- Train ----
model.train()
train_loss = 0
for imgs, labels in train_loader:
    imgs, labels = imgs.to(device), labels.to(device)

    optimizer.zero_grad()
    outputs = model(imgs)
    loss = criterion(outputs, labels)
    loss.backward()
    optimizer.step()

    train_loss += loss.item() * imgs.size(0)

train_loss /= len(train_loader.dataset)

# ---- Validation ----
model.eval()
val_loss = 0
with torch.no_grad():
    for imgs, labels in val_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        outputs = model(imgs)
        loss = criterion(outputs, labels)
        val_loss += loss.item() * imgs.size(0)

val_loss /= len(val_loader.dataset)
scheduler.step(val_loss)

elapsed = time.time() - start_time
print(f"Epoch [{epoch+1}/{num_epochs}] Train Loss: {train_loss:.4f} Val Loss: {val_loss:.4f} Time: {elapsed:.2f}s")

# Early Stopping
if val_loss < best_val_loss:
    best_val_loss = val_loss
    patience_counter = 0
    torch.save(model.state_dict(), "tinyvit_best.pth")
else:
    patience_counter += 1
    if patience_counter >= patience:
        print("Early stopping triggered!")
        break

print("Training finished Best model saved as tinyvit_best.pth")

# =====
# Grad-CAM Visualization
# =====
target_layer = model.stages[-1].blocks[-1].local_conv
grad_cam = GradCAM(model, target_layer)

# Collect first 20 test images
imgs_collected = []
labels_collected = []
for imgs, labels in test_loader:
    for i in range(len(imgs)):
        imgs_collected.append(imgs[i])
        labels_collected.append(labels[i].item())
        if len(imgs_collected) >= 20:
            break
    if len(imgs_collected) >= 20:
        break

# Plot Grad-CAM
plt.figure(figsize=(15, 10))
for idx in range(20):
    img = imgs_collected[idx].unsqueeze(0)
    label = labels_collected[idx]

    mask = grad_cam(img, class_idx=label)

    plt.subplot(4, 5, idx+1)
    img_np = img[0].permute(1, 2, 0).cpu().numpy()
    img_np = (img_np - img_np.min()) / (img_np.max() - img_np.min())
    heatmap = cv2.applyColorMap(np.uint8(255*mask), cv2.COLORMAP_JET)
    heatmap = cv2.cvtColor(heatmap, cv2.COLOR_BGR2RGB) / 255.0
    cam_img = 0.5*heatmap + 0.5*img_np

    plt.imshow(cam_img)
    plt.axis('off')

plt.tight_layout()


```

```
plt.show()
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.  
warnings.warn(
```

```
Epoch [1/35] Train Loss: 0.5339 Val Loss: 0.2788 Time: 16.08s  
Epoch [2/35] Train Loss: 0.1072 Val Loss: 0.0554 Time: 16.03s  
Epoch [3/35] Train Loss: 0.0434 Val Loss: 0.1607 Time: 16.04s  
Epoch [4/35] Train Loss: 0.0299 Val Loss: 0.0509 Time: 16.04s  
Epoch [5/35] Train Loss: 0.0134 Val Loss: 0.0233 Time: 16.09s  
Epoch [6/35] Train Loss: 0.0295 Val Loss: 0.0263 Time: 16.15s  
Epoch [7/35] Train Loss: 0.0109 Val Loss: 0.0301 Time: 15.98s  
Epoch [8/35] Train Loss: 0.0101 Val Loss: 0.0245 Time: 16.12s  
Epoch [9/35] Train Loss: 0.0070 Val Loss: 0.0184 Time: 16.13s  
Epoch [10/35] Train Loss: 0.0040 Val Loss: 0.0179 Time: 16.02s  
Epoch [11/35] Train Loss: 0.0036 Val Loss: 0.0190 Time: 16.11s  
Epoch [12/35] Train Loss: 0.0030 Val Loss: 0.0211 Time: 16.01s  
Epoch [13/35] Train Loss: 0.0028 Val Loss: 0.0207 Time: 16.13s  
Epoch [14/35] Train Loss: 0.0022 Val Loss: 0.0188 Time: 16.13s  
Epoch [15/35] Train Loss: 0.0023 Val Loss: 0.0187 Time: 16.07s  
Early stopping triggered!
```

```
Training finished Best model saved as tinyvit_best.pth
```



```
In [7]: # =====  
# TinyViT + 5-Fold Cross Validation  
# =====  
  
# -----  
# Imports  
# -----  
import torch  
import torch.nn as nn  
import torch.optim as optim  
from torch.utils.data import DataLoader, Subset  
from torchvision import datasets, transforms  
from timm import create_model  
from torch.optim.lr_scheduler import ReduceLROnPlateau  
  
import numpy as np  
from sklearn.model_selection import StratifiedKFold  
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score  
  
# -----  
# Config  
# -----  
data_dir = "/kaggle/working/augmented_split"  
batch_size = 32
```

```

num_epochs = 35
patience = 5
num_classes = 4 # Glioma, Meningioma, Pituitary, Normal
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# -----
# Data transforms
# -----
common_tfms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])

# Test dataset (kept separate for final evaluation)
test_ds = datasets.ImageFolder(root=f"{data_dir}/test", transform=common_tfms)
test_loader = DataLoader(test_ds, batch_size=batch_size, shuffle=False, num_workers=2)

# For CV (use only train split here)
full_ds = datasets.ImageFolder(root=f"{data_dir}/train", transform=common_tfms)
X = np.arange(len(full_ds))
y = [label for _, label in full_ds.samples]

# -----
# Stratified 5-Fold CV
# -----
k_folds = 5
skf = StratifiedKFold(n_splits=k_folds, shuffle=True, random_state=42)

fold_metrics = {"accuracy": [], "precision": [], "recall": [], "f1": []}

for fold, (train_idx, val_idx) in enumerate(skf.split(X, y), 1):
    print(f"\n===== Fold {fold} =====")

    # Subset datasets
    train_subset = Subset(full_ds, train_idx)
    val_subset = Subset(full_ds, val_idx)

    train_loader = DataLoader(train_subset, batch_size=batch_size, shuffle=True, num_workers=2)
    val_loader = DataLoader(val_subset, batch_size=batch_size, shuffle=False, num_workers=2)

    # Reinitialize model for each fold
    model = create_model("tiny_vit_5m_224", pretrained=True, num_classes=num_classes).to(device)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)
    scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=False)

    best_val_loss = float("inf")
    patience_counter = 0

    # -----
    # Training Loop
    # -----
    for epoch in range(num_epochs):
        # --- Train ---
        model.train()
        for imgs, labels in train_loader:
            imgs, labels = imgs.to(device), labels.to(device)
            optimizer.zero_grad()
            outputs = model(imgs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

        # --- Validation ---
        model.eval()
        val_loss, y_true, y_pred = 0, [], []
        with torch.no_grad():
            for imgs, labels in val_loader:
                imgs, labels = imgs.to(device), labels.to(device)
                outputs = model(imgs)
                loss = criterion(outputs, labels)
                val_loss += loss.item() * imgs.size(0)
                preds = outputs.argmax(1)
                y_true.extend(labels.cpu().numpy())
                y_pred.extend(preds.cpu().numpy())

        val_loss /= len(val_loader.dataset)
        scheduler.step(val_loss)

        # Early stopping
        if val_loss < best_val_loss:
            best_val_loss = val_loss

```

```

        patience_counter = 0
        best_y_true, best_y_pred = y_true[:, :], y_pred[:, :]
    else:
        patience_counter += 1
        if patience_counter >= patience:
            break

# -----
# Metrics for this fold
# -----
acc = accuracy_score(best_y_true, best_y_pred)
prec = precision_score(best_y_true, best_y_pred, average="macro", zero_division=0)
rec = recall_score(best_y_true, best_y_pred, average="macro", zero_division=0)
f1 = f1_score(best_y_true, best_y_pred, average="macro", zero_division=0)

fold_metrics["accuracy"].append(acc)
fold_metrics["precision"].append(prec)
fold_metrics["recall"].append(rec)
fold_metrics["f1"].append(f1)

print(f"Fold {fold} - Acc: {acc:.4f}, Prec: {prec:.4f}, Recall: {rec:.4f}, F1: {f1:.4f}")

# -----
# Summary (mean ± std)
# -----
print("\n===== Cross-Validation Results =====")
for metric, values in fold_metrics.items():
    mean, std = np.mean(values), np.std(values)
    print(f"{metric.capitalize()}: {mean:.4f} ± {std:.4f}")

```

===== Fold 1 =====

model.safetensors: 0% | 0.00/48.4M [00:00<?, ?B/s]

```

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
precated. Please use get_last_lr() to access the learning rate.
warnings.warn(

```

Fold 1 - Acc: 0.9970, Prec: 0.9970, Recall: 0.9971, F1: 0.9970

===== Fold 2 =====

```

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
precated. Please use get_last_lr() to access the learning rate.
warnings.warn(

```

Fold 2 - Acc: 0.9970, Prec: 0.9970, Recall: 0.9970, F1: 0.9970

===== Fold 3 =====

```

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
precated. Please use get_last_lr() to access the learning rate.
warnings.warn(

```

Fold 3 - Acc: 0.9970, Prec: 0.9970, Recall: 0.9970, F1: 0.9970

===== Fold 4 =====

```

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
precated. Please use get_last_lr() to access the learning rate.
warnings.warn(

```

Fold 4 - Acc: 1.0000, Prec: 1.0000, Recall: 1.0000, F1: 1.0000

===== Fold 5 =====

```

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
precated. Please use get_last_lr() to access the learning rate.
warnings.warn(

```

Fold 5 - Acc: 0.9985, Prec: 0.9985, Recall: 0.9985, F1: 0.9985

==== Cross-Validation Results ====

Accuracy: 0.9979 ± 0.0012

Precision: 0.9979 ± 0.0012

Recall: 0.9979 ± 0.0012

F1: 0.9979 ± 0.0012