

```
In [1]: import kagglehub

# Download latest version
path = kagglehub.dataset_download("sohansakib75/cotton-4-class")

print("Path to dataset files:", path)
```

Path to dataset files: /kaggle/input/cotton-4-class

```
In [2]: import os
import kagglehub

# Download dataset
path = kagglehub.dataset_download("sohansakib75/cotton-4-class")

print("Dataset root path:", path)

# List folders and files inside
print("Contents inside dataset folder:")
for item in os.listdir(path):
    item_path = os.path.join(path, item)
    if os.path.isdir(item_path):
        print(f" {item}/")
    else:
        print(f" {item}")
```

```
Dataset root path: /kaggle/input/cotton-4-class
Contents inside dataset folder:
Cotton leaf/
```

```
In [3]: import os
import glob

dataset_path = os.path.join(path, "Cotton leaf")

print("Path to Dataset folder:", dataset_path)

for subdir in sorted(os.listdir(dataset_path)):
    subpath = os.path.join(dataset_path, subdir)
    if os.path.isdir(subpath):
        # Count images by common formats
        image_files = glob.glob(os.path.join(subpath, "*.jpg")) + \
                      glob.glob(os.path.join(subpath, "*.jpeg")) + \
                      glob.glob(os.path.join(subpath, "*.png"))
        print(f"{subdir}: {len(image_files)} images")
```

```
Path to Dataset folder: /kaggle/input/cotton-4-class/Cotton leaf
diseased cotton leaf: 346 images
diseased cotton plant: 921 images
fresh cotton leaf: 519 images
fresh cotton plant: 514 images
```

```
In [4]: import os
import glob
import matplotlib.pyplot as plt
import random

dataset_path = "/kaggle/input/cotton-4-class/Cotton leaf"

# Classes
classes = sorted(os.listdir(dataset_path))

# Plot 5 images per class
fig, axes = plt.subplots(len(classes), 5, figsize=(15, 12))

for i, cls in enumerate(classes):
    cls_path = os.path.join(dataset_path, cls)
    image_files = glob.glob(os.path.join(cls_path, "*.jpg")) + \
                  glob.glob(os.path.join(cls_path, "*.jpeg")) + \
                  glob.glob(os.path.join(cls_path, "*.png"))

    # Randomly pick 5 images
    sample_files = random.sample(image_files, 5)

    for j, img_path in enumerate(sample_files):
        img = plt.imread(img_path)
        axes[i, j].imshow(img)
        axes[i, j].axis("off")
        if j == 2: # center column
            axes[i, j].set_title(cls, fontsize=10)

plt.tight_layout()
```

```
plt.show()
```



```
In [5]: import os
import cv2
import numpy as np
import matplotlib.pyplot as plt
import random

# Paths
input_dir = "/kaggle/input/cotton-4-class/Cotton leaf"
output_dir = "/kaggle/working/preprocessed_dataset"
os.makedirs(output_dir, exist_ok=True)

img_size = (224, 224)
num_samples = 5

classes = sorted(os.listdir(input_dir))
fig, axes = plt.subplots(len(classes), num_samples, figsize=(15, 3*len(classes)))

clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8,8))

for i, cls in enumerate(classes):
    cls_path = os.path.join(input_dir, cls)
    if os.path.isdir(cls_path):
        save_cls_path = os.path.join(output_dir, cls)
        os.makedirs(save_cls_path, exist_ok=True)

    files = [f for f in os.listdir(cls_path) if f.lower().endswith(("jpg", "jpeg", "png"))]
    sample_files = random.sample(files, min(num_samples, len(files)))

    for j, file in enumerate(files):
        img_path = os.path.join(cls_path, file)
        img = cv2.imread(img_path)

        # 1) Resize
        img_resized = cv2.resize(img, img_size)

        # 2) CLAHE on Y channel
        yuv = cv2.cvtColor(img_resized, cv2.COLOR_BGR2YUV)
        yuv[:, :, 0] = clahe.apply(yuv[:, :, 0])
```

```

    img_clahe = cv2.cvtColor(yuv, cv2.COLOR_YUV2BGR)

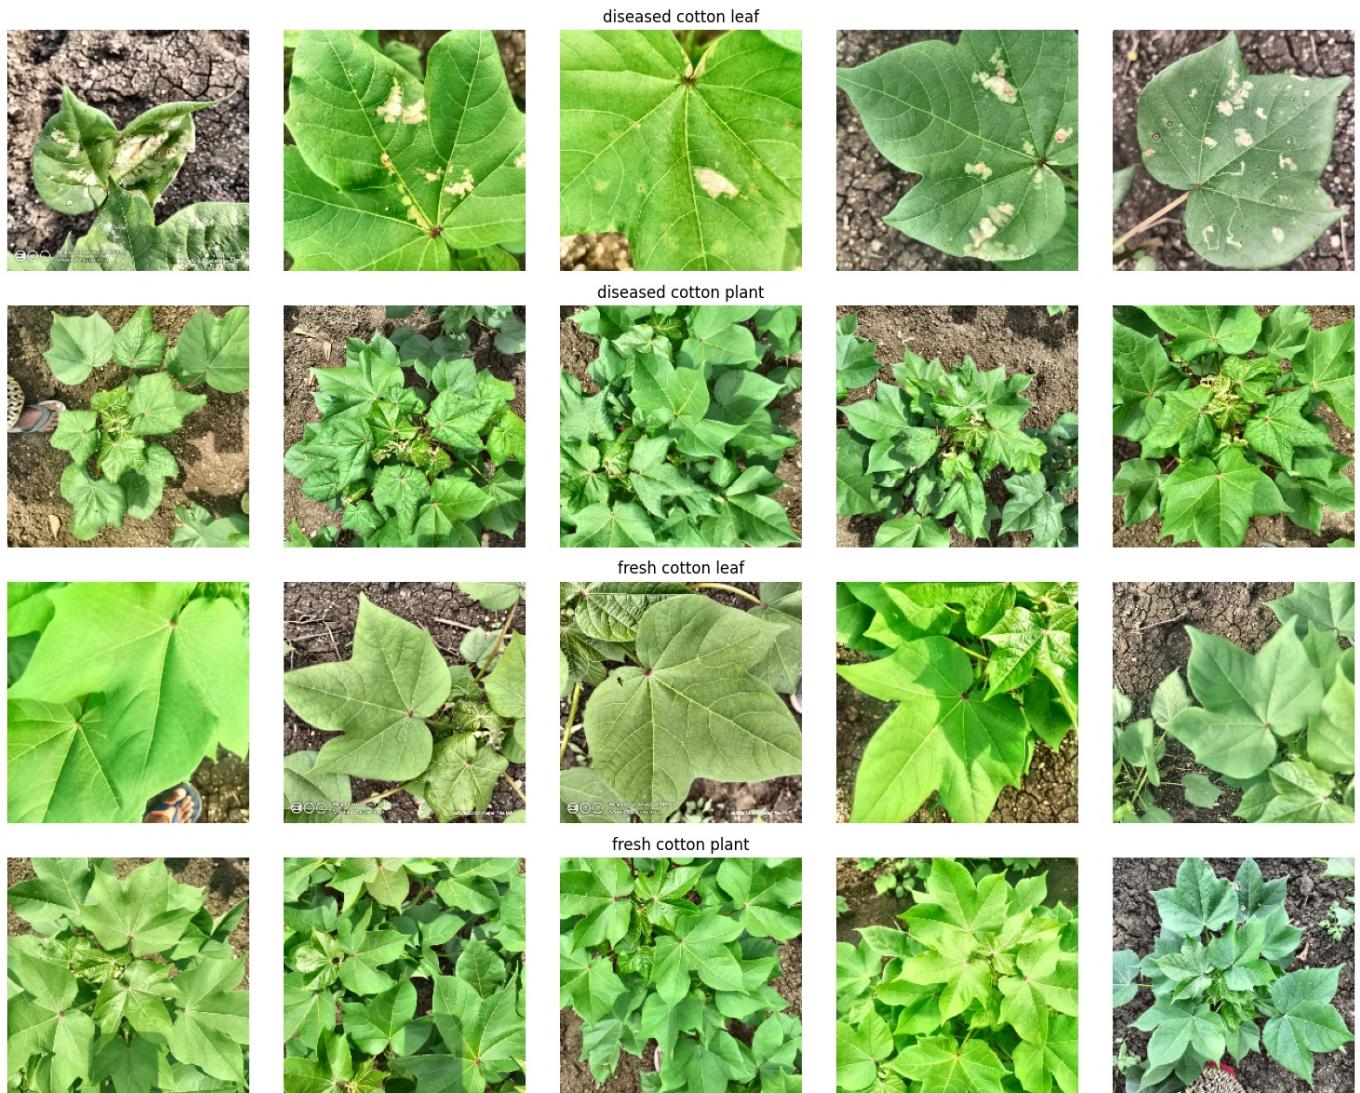
    # 3. Normalize to [0,1]
    img_final = img_clahe.astype(np.float32) / 255.0

    # Save preprocessed image
    save_img = (img_final * 255).astype(np.uint8)
    cv2.imwrite(os.path.join(save_cls_path, file), save_img)

    # Plot sample images
    if file in sample_files:
        axes[i, sample_files.index(file)].imshow(cv2.cvtColor(save_img, cv2.COLOR_BGR2RGB))
        axes[i, sample_files.index(file)].axis("off")
        if sample_files.index(file) == num_samples // 2:
            axes[i, sample_files.index(file)].set_title(cls, fontsize=12)

plt.tight_layout()
plt.show()
print("\n Preprocessed images saved in:", output_dir)

```



Preprocessed images saved in: /kaggle/working/preprocessed_dataset

In [6]:

```

import os
import shutil
import random

# Preprocessed dataset path
preprocessed_dir = "/kaggle/working/preprocessed_dataset"

# Split paths
split_base = "/kaggle/working/cotton_split"
train_dir = os.path.join(split_base, "train")
val_dir = os.path.join(split_base, "val")
test_dir = os.path.join(split_base, "test")

# Create split folders
for d in [train_dir, val_dir, test_dir]:
    os.makedirs(d, exist_ok=True)

# Split ratios
train_ratio = 0.75

```

```

val_ratio = 0.1
test_ratio = 0.15

classes = sorted(os.listdir(preprocessed_dir))

for cls in classes:
    cls_path = os.path.join(preprocessed_dir, cls)
    files = [f for f in os.listdir(cls_path) if f.lower().endswith((".jpg", ".jpeg", ".png"))]
    random.shuffle(files)

    n_total = len(files)
    n_train = int(train_ratio * n_total)
    n_val = int(val_ratio * n_total)
    n_test = n_total - n_train - n_val

    splits = {
        train_dir: files[:n_train],
        val_dir: files[n_train:n_train+n_val],
        test_dir: files[n_train+n_val:]
    }

    for split_folder, split_files in splits.items():
        cls_split_path = os.path.join(split_folder, cls)
        os.makedirs(cls_split_path, exist_ok=True)
        for f in split_files:
            shutil.copy(os.path.join(cls_path, f), os.path.join(cls_split_path, f))

print(" Dataset split into 75-10-15 and saved in:", split_base)

```

Dataset split into 75-10-15 and saved in: /kaggle/working/cotton_split

```

In [7]: import os

split_base = "/kaggle/working/cotton_split"
splits = ["train", "val", "test"]

for split in splits:
    split_path = os.path.join(split_base, split)
    print(f"\n {split.capitalize()} Split:")
    for cls in sorted(os.listdir(split_path)):
        cls_path = os.path.join(split_path, cls)
        num_images = len([f for f in os.listdir(cls_path) if f.lower().endswith((".jpg", ".jpeg", ".png"))])
        print(f"{cls}: {num_images} images")

Train Split:
diseased cotton leaf: 259 images
diseased cotton plant: 690 images
fresh cotton leaf: 389 images
fresh cotton plant: 385 images

Val Split:
diseased cotton leaf: 34 images
diseased cotton plant: 92 images
fresh cotton leaf: 51 images
fresh cotton plant: 51 images

Test Split:
diseased cotton leaf: 53 images
diseased cotton plant: 139 images
fresh cotton leaf: 79 images
fresh cotton plant: 78 images

```

```

In [8]: import os
import cv2
import numpy as np
import random

train_dir = "/kaggle/working/cotton_split/train"
aug_train_dir = "/kaggle/working/cotton_train_aug"
os.makedirs(aug_train_dir, exist_ok=True)

# Augmentation functions
def random_flip(img):
    flip_code = random.choice([-1, 0, 1])
    return cv2.flip(img, flip_code)

def random_rotate(img):
    angle = random.uniform(-25, 25)
    h, w = img.shape[:2]
    M = cv2.getRotationMatrix2D((w//2, h//2), angle, 1)
    return cv2.warpAffine(img, M, (w, h), borderMode=cv2.BORDER_REFLECT)

def random_zoom(img):
    zoom_factor = random.uniform(0.8, 1.2)

```

```

h, w = img.shape[:2]
new_h, new_w = int(h*zoom_factor), int(w*zoom_factor)
img_resized = cv2.resize(img, (new_w, new_h))
if zoom_factor < 1:
    pad_h = (h - new_h) // 2
    pad_w = (w - new_w) // 2
    img_padded = cv2.copyMakeBorder(img_resized, pad_h, h-new_h-pad_h,
                                    pad_w, w-new_w-pad_w, cv2.BORDER_REFLECT)
    return img_padded
else:
    start_h = (new_h - h)//2
    start_w = (new_w - w)//2
    return img_resized[start_h:start_h+h, start_w:start_w+w]

def random_brightness(img):
    factor = random.uniform(0.7, 1.3)
    img = img.astype(np.float32) * factor
    img = np.clip(img, 0, 255).astype(np.uint8)
    return img

augmentations = [random_flip, random_rotate, random_zoom, random_brightness]

# Apply augmentations
classes = sorted(os.listdir(train_dir))
for cls in classes:
    cls_path = os.path.join(train_dir, cls)
    save_cls_path = os.path.join(aug_train_dir, cls)
    os.makedirs(save_cls_path, exist_ok=True)

    for file in os.listdir(cls_path):
        if not file.lower().endswith((".jpg", ".jpeg", ".png")):
            continue
        img_path = os.path.join(cls_path, file)
        img = cv2.imread(img_path)

        # Save original
        cv2.imwrite(os.path.join(save_cls_path, file), img)

        # 3 random augmentations
        for k in range(3):
            aug_img = img.copy()
            aug_funcs = random.sample(augmentations, 2)
            for func in aug_funcs:
                aug_img = func(aug_img)
            filename, ext = os.path.splitext(file)
            aug_name = f"{filename}_aug{k+1}{ext}"
            cv2.imwrite(os.path.join(save_cls_path, aug_name), aug_img)

# Print image count per class
print("\n Image count per class after augmentation:")
for cls in classes:
    cls_path = os.path.join(aug_train_dir, cls)
    count = len([f for f in os.listdir(cls_path) if f.lower().endswith((".jpg", ".jpeg", ".png"))])
    print(f"[cls]: {count} images")

print("\n All train images and augmented images saved in: {aug_train_dir}")

```

Image count per class after augmentation:
diseased cotton leaf: 1036 images
diseased cotton plant: 2760 images
fresh cotton leaf: 1556 images
fresh cotton plant: 1540 images

All train images and augmented images saved in: /kaggle/working/cotton_train_aug

In [9]: !pip install timm

```

Requirement already satisfied: timm in /usr/local/lib/python3.11/dist-packages (1.0.19)
Requirement already satisfied: torch in /usr/local/lib/python3.11/dist-packages (from timm) (2.6.0+cu124)
Requirement already satisfied: torchvision in /usr/local/lib/python3.11/dist-packages (from timm) (0.21.0+cu124)
Requirement already satisfied: pyyaml in /usr/local/lib/python3.11/dist-packages (from timm) (6.0.3)
Requirement already satisfied: huggingface_hub in /usr/local/lib/python3.11/dist-packages (from timm) (1.0.0rc2)
Requirement already satisfied: safetensors in /usr/local/lib/python3.11/dist-packages (from timm) (0.5.3)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm) (3.19.1)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm) (2025.9.0)
Requirement already satisfied: packaging>=20.9 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm) (25.0)
Requirement already satisfied: httpx<1,>=0.23.0 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm) (0.28.1)
Requirement already satisfied: tqdm>=4.42.1 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm) (4.67.1)
Requirement already satisfied: typer-slim in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm)

```

```
) (0.19.2)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm) (4.15.0)
Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm) (1.1.10)
Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages (from torch->timm) (3.5)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (3.1.6)
Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch->timm)
    Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch->timm)
    Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch->timm)
    Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch->timm)
    Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cublas-cu12==12.4.5.8 (from torch->timm)
    Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cufft-cu12==11.2.1.3 (from torch->timm)
    Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-curand-cu12==10.3.5.147 (from torch->timm)
    Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch->timm)
    Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cusparse-cu12==12.3.1.170 (from torch->timm)
    Downloading nvidia_cusparse_cu12-12.3.1.170-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Requirement already satisfied: nvidia-cusparseelt-cu12==0.6.2 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (0.6.2)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (2.21.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (12.4.127)
Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch->timm)
    Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Requirement already satisfied: triton==3.2.0 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (3.2.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch->timm) (1.3.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from torchvision->timm) (1.26.4)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.11/dist-packages (from torchvision->timm) (11.3.0)
Requirement already satisfied: anyio in /usr/local/lib/python3.11/dist-packages (from httpx<1,>=0.23.0->huggingface_hub->timm) (4.11.0)
Requirement already satisfied: certifi in /usr/local/lib/python3.11/dist-packages (from httpx<1,>=0.23.0->huggingface_hub->timm) (2025.8.3)
Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.11/dist-packages (from httpx<1,>=0.23.0->huggingface_hub->timm) (1.0.9)
Requirement already satisfied: idna in /usr/local/lib/python3.11/dist-packages (from httpx<1,>=0.23.0->huggingface_hub->timm) (3.10)
Requirement already satisfied: h11=>0.16 in /usr/local/lib/python3.11/dist-packages (from httpcore==1.*->httpx<1,>=0.23.0->huggingface_hub->timm) (0.16.0)
Requirement already satisfied: MarkupSafe=>2.0 in /usr/local/lib/python3.11/dist-packages (from jinja2->torch->timm) (3.0.2)
Requirement already satisfied: mkl_fft in /usr/local/lib/python3.11/dist-packages (from numpy->torchvision->timm) (1.3.8)
Requirement already satisfied: mkl_random in /usr/local/lib/python3.11/dist-packages (from numpy->torchvision->timm) (1.2.4)
Requirement already satisfied: mkl_umath in /usr/local/lib/python3.11/dist-packages (from numpy->torchvision->timm) (0.1.1)
Requirement already satisfied: mkl in /usr/local/lib/python3.11/dist-packages (from numpy->torchvision->timm) (2.025.2.0)
Requirement already satisfied: tbb4py in /usr/local/lib/python3.11/dist-packages (from numpy->torchvision->timm) (2022.2.0)
Requirement already satisfied: mkl-service in /usr/local/lib/python3.11/dist-packages (from numpy->torchvision->timm) (2.4.1)
Requirement already satisfied: click=>8.0.0 in /usr/local/lib/python3.11/dist-packages (from typer-slim->huggingface_hub->timm) (8.3.0)
Requirement already satisfied: sniffio=>1.1 in /usr/local/lib/python3.11/dist-packages (from anyio->httpx<1,>=0.23.0->huggingface_hub->timm) (1.3.1)
Requirement already satisfied: intel-openmp<2026,>=2024 in /usr/local/lib/python3.11/dist-packages (from mkl->numpy->torchvision->timm) (2024.2.0)
Requirement already satisfied: tbb==2022.* in /usr/local/lib/python3.11/dist-packages (from mkl->numpy->torchvision->timm) (2022.2.0)
Requirement already satisfied: tcmlib==1.* in /usr/local/lib/python3.11/dist-packages (from tbb==2022.*->mkl->numpy->torchvision->timm) (1.4.0)
Requirement already satisfied: intel-cmplr-lib-rt in /usr/local/lib/python3.11/dist-packages (from mkl_umath->numpy->torchvision->timm) (2024.2.0)
Requirement already satisfied: intel-cmplr-lib-ur==2024.2.0 in /usr/local/lib/python3.11/dist-packages (from intel-openmp<2026,>=2024->mkl->numpy->torchvision->timm) (2024.2.0)
Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl (363.4 MB)
```

363.4/363.4 MB 5.0 MB/s eta 0:00:00:00:0100:01

```

Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (13.8 MB)
  13.8/13.8 MB 109.3 MB/s eta 0:00:000:010:01
Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (24.6 MB)
  24.6/24.6 MB 61.4 MB/s eta 0:00:00:00:0100:01
Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (883 kB)
  883.7/883.7 KB 46.9 MB/s eta 0:00:00
Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl (664.8 MB)
  664.8/664.8 MB 2.1 MB/s eta 0:00:00:00:0100:01
Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl (211.5 MB)
  211.5/211.5 MB 3.4 MB/s eta 0:00:00:00:0100:01
Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl (56.3 MB)
  56.3/56.3 MB 24.6 MB/s eta 0:00:000:00:0100:01
Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl (127.9 MB)
  127.9/127.9 MB 12.4 MB/s eta 0:00:000:00:0100:01
Downloading nvidia_cusparse_cu12-12.3.1.170-py3-none-manylinux2014_x86_64.whl (207.5 MB)
  207.5/207.5 MB 8.0 MB/s eta 0:00:00:00:0100:01
Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (21.1 MB)
  21.1/21.1 MB 63.7 MB/s eta 0:00:00:00:0100:01
Installing collected packages: nvidia-nvjitlink-cu12, nvidia-curand-cu12, nvidia-cufft-cu12, nvidia-cuda-runtime-cu12, nvidia-cuda-nvrtc-cu12, nvidia-cuda-cupti-cu12, nvidia-cublas-cu12, nvidia-cusparse-cu12, nvidia-cudnn-cu12, nvidia-cusolver-cu12
  Attempting uninstall: nvidia-nvjitlink-cu12
    Found existing installation: nvidia-nvjitlink-cu12 12.5.82
    Uninstalling nvidia-nvjitlink-cu12-12.5.82:
      Successfully uninstalled nvidia-nvjitlink-cu12-12.5.82
  Attempting uninstall: nvidia-curand-cu12
    Found existing installation: nvidia-curand-cu12 10.3.6.82
    Uninstalling nvidia-curand-cu12-10.3.6.82:
      Successfully uninstalled nvidia-curand-cu12-10.3.6.82
  Attempting uninstall: nvidia-cufft-cu12
    Found existing installation: nvidia-cufft-cu12 11.2.3.61
    Uninstalling nvidia-cufft-cu12-11.2.3.61:
      Successfully uninstalled nvidia-cufft-cu12-11.2.3.61
  Attempting uninstall: nvidia-cuda-runtime-cu12
    Found existing installation: nvidia-cuda-runtime-cu12 12.5.82
    Uninstalling nvidia-cuda-runtime-cu12-12.5.82:
      Successfully uninstalled nvidia-cuda-runtime-cu12-12.5.82
  Attempting uninstall: nvidia-cuda-nvrtc-cu12
    Found existing installation: nvidia-cuda-nvrtc-cu12 12.5.82
    Uninstalling nvidia-cuda-nvrtc-cu12-12.5.82:
      Successfully uninstalled nvidia-cuda-nvrtc-cu12-12.5.82
  Attempting uninstall: nvidia-cuda-cupti-cu12
    Found existing installation: nvidia-cuda-cupti-cu12 12.5.82
    Uninstalling nvidia-cuda-cupti-cu12-12.5.82:
      Successfully uninstalled nvidia-cuda-cupti-cu12-12.5.82
  Attempting uninstall: nvidia-cublas-cu12
    Found existing installation: nvidia-cublas-cu12 12.5.3.2
    Uninstalling nvidia-cublas-cu12-12.5.3.2:
      Successfully uninstalled nvidia-cublas-cu12-12.5.3.2
  Attempting uninstall: nvidia-cusparse-cu12
    Found existing installation: nvidia-cusparse-cu12 12.5.1.3
    Uninstalling nvidia-cusparse-cu12-12.5.1.3:
      Successfully uninstalled nvidia-cusparse-cu12-12.5.1.3
  Attempting uninstall: nvidia-cudnn-cu12
    Found existing installation: nvidia-cudnn-cu12 9.3.0.75
    Uninstalling nvidia-cudnn-cu12-9.3.0.75:
      Successfully uninstalled nvidia-cudnn-cu12-9.3.0.75
  Attempting uninstall: nvidia-cusolver-cu12
    Found existing installation: nvidia-cusolver-cu12 11.6.3.83
    Uninstalling nvidia-cusolver-cu12-11.6.3.83:
      Successfully uninstalled nvidia-cusolver-cu12-11.6.3.83
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.
libcugraph-cu12 25.6.0 requires libraft-cu12==25.6.*, but you have libraft-cu12 25.2.0 which is incompatible.
pylibcugraph-cu12 25.6.0 requires pylibraft-cu12==25.6.*, but you have pylibraft-cu12 25.2.0 which is incompatible.
pylibcugraph-cu12 25.6.0 requires rmm-cu12==25.6.*, but you have rmm-cu12 25.2.0 which is incompatible.
Successfully installed nvidia-cublas-cu12-12.4.5.8 nvidia-cuda-cupti-cu12-12.4.127 nvidia-cuda-nvrtc-cu12-12.4.127 nvidia-cuda-runtime-cu12-12.4.127 nvidia-cudnn-cu12-9.1.0.70 nvidia-cufft-cu12-11.2.1.3 nvidia-curand-cu12-10.3.5.147 nvidia-cusolver-cu12-11.6.1.9 nvidia-cusparse-cu12-12.3.1.170 nvidia-nvjitlink-cu12-12.4.127
```

In [10]: `import timm`

```
/usr/local/lib/python3.11/dist-packages/pydantic/_internal/_generate_schema.py:2225: UnsupportedFieldAttributeWarning: The 'repr' attribute with value False was provided to the `Field()` function, which has no effect in the context it was used. 'repr' is field-specific metadata, and can only be attached to a model field using `Annotated` metadata or by assignment. This may have happened because an `Annotated` type alias using the `type` statement was used, or if the `Field()` function was attached to a single member of a union type.  
    warnings.warn()  
/usr/local/lib/python3.11/dist-packages/pydantic/_internal/_generate_schema.py:2225: UnsupportedFieldAttributeWarning: The 'frozen' attribute with value True was provided to the `Field()` function, which has no effect in the context it was used. 'frozen' is field-specific metadata, and can only be attached to a model field using `Annotated` metadata or by assignment. This may have happened because an `Annotated` type alias using the `type` statement was used, or if the `Field()` function was attached to a single member of a union type.  
    warnings.warn()
```

```
In [11]:  
import torch  
import torch.nn as nn  
import torch.optim as optim  
from torch.utils.data import DataLoader, Subset  
from torchvision import datasets, transforms  
from timm import create_model  
from torch.optim.lr_scheduler import ReduceLROnPlateau  
from sklearn.model_selection import KFold  
from sklearn.metrics import accuracy_score  
import numpy as np  
import time  
import gc  
  
# =====  
# CONFIGURATION  
# =====  
train_dir = "/kaggle/working/cotton_train_aug"  
val_dir = "/kaggle/working/cotton_split/val"  
test_dir = "/kaggle/working/cotton_split/test"  
  
num_classes = 4  
batch_size = 32  
num_epochs = 35  
patience = 5  
num_folds = 5  
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
class_names = ["diseased cotton leaf", "diseased cotton plant", "fresh cotton leaf", "fresh cotton plant"]  
  
# =====  
# DATASET & TRANSFORMS  
# =====  
transform = transforms.Compose([  
    transforms.Resize((224, 224)),  
    transforms.ToTensor(),  
    transforms.Normalize([0.5]*3, [0.5]*3)  
])  
  
dataset = datasets.ImageFolder(root=train_dir, transform=transform)  
  
# =====  
# 5-FOLD CROSS VALIDATION  
# =====  
kf = KFold(n_splits=num_folds, shuffle=True, random_state=42)  
fold_accuracies = []  
  
for fold, (train_idx, val_idx) in enumerate(kf.split(dataset)):  
    print(f"\n===== Fold {fold+1}/{num_folds} =====")  
  
    train_subset = Subset(dataset, train_idx)  
    val_subset = Subset(dataset, val_idx)  
  
    train_loader = DataLoader(train_subset, batch_size=batch_size, shuffle=True, num_workers=2)  
    val_loader = DataLoader(val_subset, batch_size=batch_size, shuffle=False, num_workers=2)  
  
    # =====  
    # Model Setup  
    # =====  
    model = create_model("mobilevit_s", pretrained=True, num_classes=num_classes, drop_rate=0.2)  
    model = model.to(device)  
  
    criterion = nn.CrossEntropyLoss()  
    optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)  
    scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=False)  
  
    best_val_loss = float("inf")  
    patience_counter = 0  
  
    # =====  
    # Training  
    # =====
```

```

for epoch in range(num_epochs):
    start_time = time.time()
    model.train()
    running_loss, correct = 0, 0

    for imgs, labels in train_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(imgs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item() * imgs.size(0)
        correct += (outputs.argmax(1) == labels).sum().item()

    train_loss = running_loss / len(train_loader.dataset)
    train_acc = correct / len(train_loader.dataset)

# ----- Validation -----
model.eval()
val_loss, val_correct = 0, 0
with torch.no_grad():
    for imgs, labels in val_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        outputs = model(imgs)
        loss = criterion(outputs, labels)

        val_loss += loss.item() * imgs.size(0)
        val_correct += (outputs.argmax(1) == labels).sum().item()

    val_loss /= len(val_loader.dataset)
    val_acc = val_correct / len(val_loader.dataset)
    scheduler.step(val_loss)

elapsed = time.time() - start_time
print(f"Epoch [{epoch+1}/{num_epochs}] | Train Acc: {train_acc:.4f} | Val Acc: {val_acc:.4f} | Time: {elapsed:.2f}s")

# Early Stopping
if val_loss < best_val_loss:
    best_val_loss = val_loss
    patience_counter = 0
    torch.save(model.state_dict(), f"fold_{fold+1}_best.pth")
else:
    patience_counter += 1
    if patience_counter >= patience:
        print("Early stopping triggered!")
        break

# =====
# Evaluation on Validation Set
# =====
model.load_state_dict(torch.load(f"fold_{fold+1}_best.pth"))
model.eval()
y_true, y_pred = [], []

with torch.no_grad():
    for imgs, labels in val_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        outputs = model(imgs)
        preds = outputs.argmax(1)
        y_true.extend(labels.cpu().numpy())
        y_pred.extend(preds.cpu().numpy())

acc = accuracy_score(y_true, y_pred)
fold_accuracies.append(acc)
print(f"Fold {fold+1} Accuracy: {acc:.4f}")

# =====
# MEMORY CLEANUP
# =====
del model, optimizer, scheduler, train_loader, val_loader, train_subset, val_subset
torch.cuda.empty_cache()
gc.collect()

# =====
# RESULTS SUMMARY
# =====
mean_acc = np.mean(fold_accuracies)
std_acc = np.std(fold_accuracies)
print("\n===== 5-FOLD CROSS VALIDATION RESULTS =====")
for i, acc in enumerate(fold_accuracies):
    print(f"Fold {i+1}: {acc:.4f}")

```

```
print(f"\nMean Accuracy: {mean_acc:.4f}")
print(f"Standard Deviation: {std_acc:.4f}")
```

===== Fold 1/5 =====

model.safetensors: 0% | 0.00/22.4M [00:00<?, ?B/s]

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.

```
warnings.warn(
```

```
Epoch [1/35] | Train Acc: 0.9048 | Val Acc: 0.9920 | Time: 35.12s
Epoch [2/35] | Train Acc: 0.9829 | Val Acc: 0.9935 | Time: 34.29s
Epoch [3/35] | Train Acc: 0.9904 | Val Acc: 0.9964 | Time: 34.14s
Epoch [4/35] | Train Acc: 0.9949 | Val Acc: 0.9956 | Time: 34.27s
Epoch [5/35] | Train Acc: 0.9955 | Val Acc: 0.9993 | Time: 34.24s
Epoch [6/35] | Train Acc: 0.9949 | Val Acc: 0.9971 | Time: 34.21s
Epoch [7/35] | Train Acc: 0.9953 | Val Acc: 0.9949 | Time: 34.42s
Epoch [8/35] | Train Acc: 0.9980 | Val Acc: 0.9971 | Time: 34.26s
Epoch [9/35] | Train Acc: 0.9985 | Val Acc: 0.9985 | Time: 34.25s
Epoch [10/35] | Train Acc: 0.9984 | Val Acc: 0.9942 | Time: 34.40s
Epoch [11/35] | Train Acc: 0.9985 | Val Acc: 0.9978 | Time: 34.13s
Epoch [12/35] | Train Acc: 0.9998 | Val Acc: 0.9978 | Time: 34.19s
Epoch [13/35] | Train Acc: 0.9996 | Val Acc: 0.9985 | Time: 34.32s
Epoch [14/35] | Train Acc: 0.9991 | Val Acc: 0.9964 | Time: 34.32s
Epoch [15/35] | Train Acc: 0.9993 | Val Acc: 0.9978 | Time: 34.11s
Epoch [16/35] | Train Acc: 0.9998 | Val Acc: 0.9978 | Time: 34.23s
Epoch [17/35] | Train Acc: 1.0000 | Val Acc: 0.9985 | Time: 34.43s
Epoch [18/35] | Train Acc: 0.9998 | Val Acc: 0.9993 | Time: 34.30s
Epoch [19/35] | Train Acc: 0.9998 | Val Acc: 0.9993 | Time: 34.18s
Epoch [20/35] | Train Acc: 0.9998 | Val Acc: 0.9993 | Time: 34.13s
Epoch [21/35] | Train Acc: 0.9996 | Val Acc: 0.9978 | Time: 34.36s
Epoch [22/35] | Train Acc: 0.9996 | Val Acc: 0.9978 | Time: 34.32s
Epoch [23/35] | Train Acc: 0.9995 | Val Acc: 0.9978 | Time: 34.38s
Epoch [24/35] | Train Acc: 0.9998 | Val Acc: 0.9978 | Time: 34.30s
Epoch [25/35] | Train Acc: 0.9998 | Val Acc: 0.9978 | Time: 34.20s
```

Early stopping triggered!

Fold 1 Accuracy: 0.9993

===== Fold 2/5 =====

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.

```
warnings.warn(
```

```
Epoch [1/35] | Train Acc: 0.8776 | Val Acc: 0.9826 | Time: 34.30s
Epoch [2/35] | Train Acc: 0.9793 | Val Acc: 0.9869 | Time: 34.20s
Epoch [3/35] | Train Acc: 0.9880 | Val Acc: 0.9920 | Time: 34.21s
Epoch [4/35] | Train Acc: 0.9906 | Val Acc: 0.9993 | Time: 34.25s
Epoch [5/35] | Train Acc: 0.9933 | Val Acc: 0.9978 | Time: 34.21s
Epoch [6/35] | Train Acc: 0.9978 | Val Acc: 0.9964 | Time: 34.26s
Epoch [7/35] | Train Acc: 0.9964 | Val Acc: 0.9993 | Time: 34.36s
Epoch [8/35] | Train Acc: 0.9940 | Val Acc: 0.9985 | Time: 34.25s
Epoch [9/35] | Train Acc: 0.9978 | Val Acc: 0.9993 | Time: 34.28s
Epoch [10/35] | Train Acc: 0.9973 | Val Acc: 0.9985 | Time: 34.38s
Epoch [11/35] | Train Acc: 0.9989 | Val Acc: 0.9993 | Time: 34.14s
Epoch [12/35] | Train Acc: 0.9995 | Val Acc: 0.9978 | Time: 34.19s
Epoch [13/35] | Train Acc: 0.9987 | Val Acc: 0.9949 | Time: 34.19s
Epoch [14/35] | Train Acc: 0.9973 | Val Acc: 0.9985 | Time: 34.11s
Epoch [15/35] | Train Acc: 0.9993 | Val Acc: 0.9956 | Time: 34.16s
Epoch [16/35] | Train Acc: 0.9996 | Val Acc: 0.9993 | Time: 34.23s
```

Early stopping triggered!

Fold 2 Accuracy: 0.9993

===== Fold 3/5 =====

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.

```
warnings.warn(
```

```
Epoch [1/35] | Train Acc: 0.8894 | Val Acc: 0.9877 | Time: 34.25s
Epoch [2/35] | Train Acc: 0.9804 | Val Acc: 0.9927 | Time: 34.25s
Epoch [3/35] | Train Acc: 0.9917 | Val Acc: 0.9964 | Time: 34.18s
Epoch [4/35] | Train Acc: 0.9944 | Val Acc: 0.9942 | Time: 34.15s
Epoch [5/35] | Train Acc: 0.9949 | Val Acc: 0.9964 | Time: 34.33s
Epoch [6/35] | Train Acc: 0.9955 | Val Acc: 0.9942 | Time: 34.19s
Epoch [7/35] | Train Acc: 0.9966 | Val Acc: 0.9877 | Time: 34.31s
Epoch [8/35] | Train Acc: 0.9975 | Val Acc: 0.9927 | Time: 34.30s
Epoch [9/35] | Train Acc: 0.9987 | Val Acc: 0.9971 | Time: 34.25s
Epoch [10/35] | Train Acc: 0.9985 | Val Acc: 0.9978 | Time: 34.14s
Epoch [11/35] | Train Acc: 0.9995 | Val Acc: 0.9978 | Time: 34.14s
Epoch [12/35] | Train Acc: 0.9980 | Val Acc: 0.9985 | Time: 34.29s
Epoch [13/35] | Train Acc: 0.9996 | Val Acc: 0.9971 | Time: 34.20s
Epoch [14/35] | Train Acc: 0.9998 | Val Acc: 0.9978 | Time: 34.18s
Epoch [15/35] | Train Acc: 0.9995 | Val Acc: 0.9978 | Time: 34.33s
```

Early stopping triggered!

Fold 3 Accuracy: 0.9978

===== Fold 4/5 =====

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de  
precated. Please use get_last_lr() to access the learning rate.  
    warnings.warn(  
Epoch [1/35] | Train Acc: 0.8932 | Val Acc: 0.9862 | Time: 34.30s  
Epoch [2/35] | Train Acc: 0.9839 | Val Acc: 0.9927 | Time: 34.32s  
Epoch [3/35] | Train Acc: 0.9900 | Val Acc: 0.9935 | Time: 34.16s  
Epoch [4/35] | Train Acc: 0.9895 | Val Acc: 0.9964 | Time: 34.19s  
Epoch [5/35] | Train Acc: 0.9962 | Val Acc: 0.9978 | Time: 34.23s  
Epoch [6/35] | Train Acc: 0.9956 | Val Acc: 0.9964 | Time: 34.22s  
Epoch [7/35] | Train Acc: 0.9975 | Val Acc: 0.9942 | Time: 34.26s  
Epoch [8/35] | Train Acc: 0.9982 | Val Acc: 0.9949 | Time: 34.17s  
Epoch [9/35] | Train Acc: 0.9967 | Val Acc: 0.9978 | Time: 34.37s  
Epoch [10/35] | Train Acc: 0.9985 | Val Acc: 0.9978 | Time: 34.27s  
Epoch [11/35] | Train Acc: 0.9995 | Val Acc: 0.9985 | Time: 34.23s  
Epoch [12/35] | Train Acc: 0.9987 | Val Acc: 0.9993 | Time: 34.30s  
Epoch [13/35] | Train Acc: 0.9998 | Val Acc: 0.9985 | Time: 34.12s  
Epoch [14/35] | Train Acc: 0.9998 | Val Acc: 0.9985 | Time: 34.25s  
Epoch [15/35] | Train Acc: 0.9991 | Val Acc: 0.9978 | Time: 34.33s  
Epoch [16/35] | Train Acc: 0.9987 | Val Acc: 0.9978 | Time: 34.24s  
Epoch [17/35] | Train Acc: 0.9996 | Val Acc: 0.9985 | Time: 34.17s  
Early stopping triggered!  
Fold 4 Accuracy: 0.9993
```

```
===== Fold 5/5 =====
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de  
precated. Please use get_last_lr() to access the learning rate.  
    warnings.warn(  
Epoch [1/35] | Train Acc: 0.8886 | Val Acc: 0.9797 | Time: 34.26s  
Epoch [2/35] | Train Acc: 0.9799 | Val Acc: 0.9913 | Time: 34.30s  
Epoch [3/35] | Train Acc: 0.9927 | Val Acc: 0.9927 | Time: 34.25s  
Epoch [4/35] | Train Acc: 0.9933 | Val Acc: 0.9891 | Time: 34.40s  
Epoch [5/35] | Train Acc: 0.9915 | Val Acc: 0.9985 | Time: 34.23s  
Epoch [6/35] | Train Acc: 0.9940 | Val Acc: 0.9978 | Time: 34.34s  
Epoch [7/35] | Train Acc: 0.9967 | Val Acc: 0.9964 | Time: 34.18s  
Epoch [8/35] | Train Acc: 0.9973 | Val Acc: 0.9964 | Time: 34.19s  
Epoch [9/35] | Train Acc: 0.9984 | Val Acc: 0.9985 | Time: 34.14s  
Epoch [10/35] | Train Acc: 0.9980 | Val Acc: 0.9985 | Time: 34.37s  
Epoch [11/35] | Train Acc: 0.9991 | Val Acc: 0.9985 | Time: 34.26s  
Epoch [12/35] | Train Acc: 0.9993 | Val Acc: 0.9993 | Time: 34.20s  
Epoch [13/35] | Train Acc: 0.9996 | Val Acc: 0.9985 | Time: 34.13s  
Epoch [14/35] | Train Acc: 0.9995 | Val Acc: 0.9985 | Time: 34.33s  
Epoch [15/35] | Train Acc: 0.9996 | Val Acc: 0.9985 | Time: 34.23s  
Epoch [16/35] | Train Acc: 0.9995 | Val Acc: 0.9985 | Time: 34.39s  
Epoch [17/35] | Train Acc: 0.9996 | Val Acc: 0.9985 | Time: 34.38s  
Early stopping triggered!  
Fold 5 Accuracy: 0.9993
```

```
===== 5-FOLD CROSS VALIDATION RESULTS =====
```

```
Fold 1: 0.9993  
Fold 2: 0.9993  
Fold 3: 0.9978  
Fold 4: 0.9993  
Fold 5: 0.9993
```

```
Mean Accuracy: 0.9990  
Standard Deviation: 0.0006
```

```
In [12]:
```

```
import torch  
import torch.nn as nn  
import torch.optim as optim  
from torch.utils.data import DataLoader, Subset  
from torchvision import datasets, transforms  
from timm import create_model  
from torch.optim.lr_scheduler import ReduceLROnPlateau  
from sklearn.model_selection import KFold  
from sklearn.metrics import accuracy_score  
import numpy as np  
import time  
import gc  
  
# =====  
# CONFIGURATION  
# =====  
train_dir = "/kaggle/working/cotton_train_aug"  
val_dir = "/kaggle/working/cotton_split/val"  
test_dir = "/kaggle/working/cotton_split/test"  
  
num_classes = 4  
batch_size = 32  
num_epochs = 35  
patience = 5  
num_folds = 5  
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```

class_names = ["diseased cotton leaf", "diseased cotton plant", "fresh cotton leaf", "fresh cotton plant"]

# =====
# DATASET & TRANSFORMS
# =====
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])

dataset = datasets.ImageFolder(root=train_dir, transform=transform)

# =====
# 5-FOLD CROSS VALIDATION
# =====
kf = KFold(n_splits=num_folds, shuffle=True, random_state=42)
fold_accuracies = []

for fold, (train_idx, val_idx) in enumerate(kf.split(dataset)):
    print(f"\n===== Fold {fold+1}/{num_folds} =====")

    train_subset = Subset(dataset, train_idx)
    val_subset = Subset(dataset, val_idx)

    train_loader = DataLoader(train_subset, batch_size=batch_size, shuffle=True, num_workers=2)
    val_loader = DataLoader(val_subset, batch_size=batch_size, shuffle=False, num_workers=2)

    # =====
    # Model Setup
    # =====
    model = create_model("tiny_vit_5m_224", pretrained=True, num_classes=num_classes)
    model = model.to(device)

    criterion = nn.CrossEntropyLoss()
    optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)
    scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=False)

    best_val_loss = float("inf")
    patience_counter = 0

    # =====
    # Training
    # =====
    for epoch in range(num_epochs):
        start_time = time.time()
        model.train()
        running_loss, correct = 0, 0

        for imgs, labels in train_loader:
            imgs, labels = imgs.to(device), labels.to(device)
            optimizer.zero_grad()
            outputs = model(imgs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            running_loss += loss.item() * imgs.size(0)
            correct += (outputs.argmax(1) == labels).sum().item()

        train_loss = running_loss / len(train_loader.dataset)
        train_acc = correct / len(train_loader.dataset)

        # ---- Validation ---
        model.eval()
        val_loss, val_correct = 0, 0
        with torch.no_grad():
            for imgs, labels in val_loader:
                imgs, labels = imgs.to(device), labels.to(device)
                outputs = model(imgs)
                loss = criterion(outputs, labels)

                val_loss += loss.item() * imgs.size(0)
                val_correct += (outputs.argmax(1) == labels).sum().item()

        val_loss /= len(val_loader.dataset)
        val_acc = val_correct / len(val_loader.dataset)
        scheduler.step(val_loss)

        elapsed = time.time() - start_time
        print(f"Epoch [{epoch+1}/{num_epochs}] | Train Acc: {train_acc:.4f} | Val Acc: {val_acc:.4f} | Time: {elapsed:.2f} s")

    # Early Stopping

```

```

    if val_loss < best_val_loss:
        best_val_loss = val_loss
        patience_counter = 0
        torch.save(model.state_dict(), f"fold_{fold+1}_best.pth")
    else:
        patience_counter += 1
        if patience_counter >= patience:
            print("Early stopping triggered!")
            break

# =====
# Evaluation on Validation Set
# =====
model.load_state_dict(torch.load(f"fold_{fold+1}_best.pth"))
model.eval()
y_true, y_pred = [], []

with torch.no_grad():
    for imgs, labels in val_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        outputs = model(imgs)
        preds = outputs.argmax(1)
        y_true.extend(labels.cpu().numpy())
        y_pred.extend(preds.cpu().numpy())

acc = accuracy_score(y_true, y_pred)
fold_accuracies.append(acc)
print(f"Fold {fold+1} Accuracy: {acc:.4f}")

# =====
# MEMORY CLEANUP
# =====
del model, optimizer, scheduler, train_loader, val_loader, train_subset, val_subset
torch.cuda.empty_cache()
gc.collect()

# =====
# RESULTS SUMMARY
# =====
mean_acc = np.mean(fold_accuracies)
std_acc = np.std(fold_accuracies)
print("\n===== 5-FOLD CROSS VALIDATION RESULTS =====")
for i, acc in enumerate(fold_accuracies):
    print(f"Fold {i+1}: {acc:.4f}")
print(f"\nMean Accuracy: {mean_acc:.4f}")
print(f"Standard Deviation: {std_acc:.4f}")

```

===== Fold 1/5 =====

model.safetensors: 0%| 0.00/48.4M [00:00<?, ?B/s]

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
 warnings.warn(

Epoch [1/35] | Train Acc: 0.9144 | Val Acc: 0.9898 | Time: 27.11s
 Epoch [2/35] | Train Acc: 0.9898 | Val Acc: 0.9978 | Time: 26.65s
 Epoch [3/35] | Train Acc: 0.9969 | Val Acc: 0.9978 | Time: 26.80s
 Epoch [4/35] | Train Acc: 0.9953 | Val Acc: 0.9978 | Time: 26.72s
 Epoch [5/35] | Train Acc: 0.9978 | Val Acc: 0.9978 | Time: 26.76s
 Epoch [6/35] | Train Acc: 0.9967 | Val Acc: 0.9985 | Time: 26.63s
 Epoch [7/35] | Train Acc: 0.9989 | Val Acc: 0.9971 | Time: 26.72s
 Epoch [8/35] | Train Acc: 0.9973 | Val Acc: 0.9993 | Time: 26.83s
 Epoch [9/35] | Train Acc: 0.9991 | Val Acc: 0.9978 | Time: 26.73s
 Epoch [10/35] | Train Acc: 0.9995 | Val Acc: 0.9978 | Time: 26.71s
 Epoch [11/35] | Train Acc: 0.9989 | Val Acc: 0.9978 | Time: 26.77s
 Epoch [12/35] | Train Acc: 0.9985 | Val Acc: 0.9978 | Time: 26.86s
 Epoch [13/35] | Train Acc: 0.9996 | Val Acc: 0.9964 | Time: 26.83s
 Epoch [14/35] | Train Acc: 0.9998 | Val Acc: 0.9949 | Time: 26.72s
 Epoch [15/35] | Train Acc: 0.9995 | Val Acc: 1.0000 | Time: 26.91s
 Epoch [16/35] | Train Acc: 1.0000 | Val Acc: 0.9985 | Time: 26.85s
 Epoch [17/35] | Train Acc: 0.9996 | Val Acc: 0.9993 | Time: 26.73s
 Epoch [18/35] | Train Acc: 1.0000 | Val Acc: 0.9978 | Time: 26.79s
 Epoch [19/35] | Train Acc: 0.9998 | Val Acc: 0.9993 | Time: 26.83s
 Epoch [20/35] | Train Acc: 0.9998 | Val Acc: 0.9993 | Time: 26.82s
 Early stopping triggered!

Fold 1 Accuracy: 1.0000

===== Fold 2/5 =====

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
 warnings.warn(

```
Epoch [1/35] | Train Acc: 0.9068 | Val Acc: 0.9927 | Time: 27.04s
Epoch [2/35] | Train Acc: 0.9891 | Val Acc: 0.9964 | Time: 26.85s
Epoch [3/35] | Train Acc: 0.9964 | Val Acc: 0.9927 | Time: 26.76s
Epoch [4/35] | Train Acc: 0.9978 | Val Acc: 0.9978 | Time: 26.91s

Epoch [5/35] | Train Acc: 0.9958 | Val Acc: 0.9964 | Time: 26.81s
Epoch [6/35] | Train Acc: 0.9964 | Val Acc: 0.9993 | Time: 26.81s
Epoch [7/35] | Train Acc: 0.9989 | Val Acc: 0.9985 | Time: 26.85s
Epoch [8/35] | Train Acc: 0.9998 | Val Acc: 0.9993 | Time: 26.84s
Epoch [9/35] | Train Acc: 1.0000 | Val Acc: 0.9978 | Time: 26.84s
Epoch [10/35] | Train Acc: 0.9996 | Val Acc: 0.9985 | Time: 26.84s
Epoch [11/35] | Train Acc: 0.9980 | Val Acc: 0.9949 | Time: 26.85s
Epoch [12/35] | Train Acc: 0.9984 | Val Acc: 0.9985 | Time: 26.79s
Epoch [13/35] | Train Acc: 0.9985 | Val Acc: 0.9985 | Time: 27.04s
Epoch [14/35] | Train Acc: 0.9995 | Val Acc: 0.9956 | Time: 26.97s
Epoch [15/35] | Train Acc: 0.9996 | Val Acc: 1.0000 | Time: 26.87s
Epoch [16/35] | Train Acc: 1.0000 | Val Acc: 0.9993 | Time: 26.85s
Epoch [17/35] | Train Acc: 0.9993 | Val Acc: 0.9985 | Time: 26.90s
Epoch [18/35] | Train Acc: 0.9998 | Val Acc: 0.9993 | Time: 26.85s
Epoch [19/35] | Train Acc: 1.0000 | Val Acc: 0.9993 | Time: 26.89s
Epoch [20/35] | Train Acc: 1.0000 | Val Acc: 0.9993 | Time: 26.83s
Early stopping triggered!
Fold 2 Accuracy: 1.0000
```

===== Fold 3/5 =====

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
precated. Please use get_last_lr() to access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/35] | Train Acc: 0.8899 | Val Acc: 0.9840 | Time: 26.75s
Epoch [2/35] | Train Acc: 0.9898 | Val Acc: 0.9891 | Time: 26.75s
Epoch [3/35] | Train Acc: 0.9960 | Val Acc: 0.9913 | Time: 26.92s
Epoch [4/35] | Train Acc: 0.9951 | Val Acc: 0.9949 | Time: 26.93s
Epoch [5/35] | Train Acc: 0.9975 | Val Acc: 0.9942 | Time: 26.75s
Epoch [6/35] | Train Acc: 0.9975 | Val Acc: 0.9971 | Time: 26.94s
Epoch [7/35] | Train Acc: 0.9982 | Val Acc: 0.9971 | Time: 26.95s
Epoch [8/35] | Train Acc: 0.9995 | Val Acc: 0.9978 | Time: 26.76s
Epoch [9/35] | Train Acc: 0.9982 | Val Acc: 0.9964 | Time: 26.84s
Epoch [10/35] | Train Acc: 0.9995 | Val Acc: 0.9956 | Time: 26.77s
Epoch [11/35] | Train Acc: 0.9998 | Val Acc: 0.9985 | Time: 26.86s
Epoch [12/35] | Train Acc: 1.0000 | Val Acc: 0.9978 | Time: 26.84s
Epoch [13/35] | Train Acc: 0.9995 | Val Acc: 0.9971 | Time: 26.85s
Early stopping triggered!
Fold 3 Accuracy: 0.9978
```

===== Fold 4/5 =====

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
precated. Please use get_last_lr() to access the learning rate.
```

```
    warnings.warn(
```

```
Epoch [1/35] | Train Acc: 0.9120 | Val Acc: 0.9906 | Time: 26.86s
Epoch [2/35] | Train Acc: 0.9898 | Val Acc: 0.9978 | Time: 26.70s
Epoch [3/35] | Train Acc: 0.9962 | Val Acc: 0.9956 | Time: 26.82s
Epoch [4/35] | Train Acc: 0.9984 | Val Acc: 0.9956 | Time: 26.87s
Epoch [5/35] | Train Acc: 0.9980 | Val Acc: 0.9956 | Time: 26.83s
Epoch [6/35] | Train Acc: 0.9960 | Val Acc: 0.9971 | Time: 26.93s
Epoch [7/35] | Train Acc: 0.9980 | Val Acc: 0.9971 | Time: 26.90s
Epoch [8/35] | Train Acc: 0.9993 | Val Acc: 0.9964 | Time: 26.82s
Epoch [9/35] | Train Acc: 0.9980 | Val Acc: 0.9978 | Time: 26.84s
Epoch [10/35] | Train Acc: 0.9980 | Val Acc: 0.9964 | Time: 26.93s
Epoch [11/35] | Train Acc: 0.9993 | Val Acc: 0.9978 | Time: 26.86s
Epoch [12/35] | Train Acc: 0.9995 | Val Acc: 0.9971 | Time: 26.90s
Early stopping triggered!
Fold 4 Accuracy: 0.9971
```

===== Fold 5/5 =====

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
precated. Please use get_last_lr() to access the learning rate.
```

```
    warnings.warn(
```

```

Epoch [1/35] | Train Acc: 0.8954 | Val Acc: 0.9913 | Time: 26.75s
Epoch [2/35] | Train Acc: 0.9908 | Val Acc: 0.9964 | Time: 26.82s
Epoch [3/35] | Train Acc: 0.9955 | Val Acc: 0.9724 | Time: 26.84s
Epoch [4/35] | Train Acc: 0.9958 | Val Acc: 0.9782 | Time: 26.93s
Epoch [5/35] | Train Acc: 0.9975 | Val Acc: 0.9956 | Time: 27.07s
Epoch [6/35] | Train Acc: 0.9975 | Val Acc: 0.9978 | Time: 26.81s
Epoch [7/35] | Train Acc: 0.9998 | Val Acc: 0.9985 | Time: 26.86s
Epoch [8/35] | Train Acc: 0.9989 | Val Acc: 0.9935 | Time: 26.89s
Epoch [9/35] | Train Acc: 0.9993 | Val Acc: 0.9971 | Time: 26.87s
Epoch [10/35] | Train Acc: 0.9951 | Val Acc: 0.9906 | Time: 26.78s
Epoch [11/35] | Train Acc: 0.9989 | Val Acc: 0.9978 | Time: 26.98s
Epoch [12/35] | Train Acc: 0.9995 | Val Acc: 0.9978 | Time: 26.75s
Early stopping triggered!
Fold 5 Accuracy: 0.9985

```

===== 5-FOLD CROSS VALIDATION RESULTS =====

```

Fold 1: 1.0000
Fold 2: 1.0000
Fold 3: 0.9978
Fold 4: 0.9971
Fold 5: 0.9985

```

Mean Accuracy: 0.9987
Standard Deviation: 0.0012

In [13]:

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, Subset
from torchvision import datasets, transforms
from timm import create_model
from torch.optim.lr_scheduler import ReduceLROnPlateau
from sklearn.model_selection import KFold
from sklearn.metrics import accuracy_score
import numpy as np
import time
import gc

# =====
# CONFIGURATION
# =====
train_dir = "/kaggle/working/cotton_train_aug"
val_dir = "/kaggle/working/cotton_split/val"
test_dir = "/kaggle/working/cotton_split/test"

num_classes = 4
batch_size = 32
num_epochs = 35
patience = 5
num_folds = 5
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
class_names = ["diseased cotton leaf", "diseased cotton plant", "fresh cotton leaf", "fresh cotton plant"]

# =====
# DATASET & TRANSFORMS
# =====
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])

dataset = datasets.ImageFolder(root=train_dir, transform=transform)

# =====
# 5-FOLD CROSS VALIDATION
# =====
kf = KFold(n_splits=num_folds, shuffle=True, random_state=42)
fold_accuracies = []

for fold, (train_idx, val_idx) in enumerate(kf.split(dataset)):
    print(f"\n===== Fold {fold+1}/{num_folds} =====")

    train_subset = Subset(dataset, train_idx)
    val_subset = Subset(dataset, val_idx)

    train_loader = DataLoader(train_subset, batch_size=batch_size, shuffle=True, num_workers=2)
    val_loader = DataLoader(val_subset, batch_size=batch_size, shuffle=False, num_workers=2)

    # =====
    # Model Setup
    # =====
    model = create_model("levit_128", pretrained=True, num_classes=num_classes) # smaller LeViT

```

```

model = model.to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)
scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=False)

best_val_loss = float("inf")
patience_counter = 0

# =====
# Training
# =====
for epoch in range(num_epochs):
    start_time = time.time()
    model.train()
    running_loss, correct = 0, 0

    for imgs, labels in train_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(imgs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item() * imgs.size(0)
        correct += (outputs.argmax(1) == labels).sum().item()

    train_loss = running_loss / len(train_loader.dataset)
    train_acc = correct / len(train_loader.dataset)

    # ---- Validation ----
    model.eval()
    val_loss, val_correct = 0, 0
    with torch.no_grad():
        for imgs, labels in val_loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            loss = criterion(outputs, labels)

            val_loss += loss.item() * imgs.size(0)
            val_correct += (outputs.argmax(1) == labels).sum().item()

    val_loss /= len(val_loader.dataset)
    val_acc = val_correct / len(val_loader.dataset)
    scheduler.step(val_loss)

    elapsed = time.time() - start_time
    print(f"Epoch [{epoch+1}/{num_epochs}] | Train Acc: {train_acc:.4f} | Val Acc: {val_acc:.4f} | Time: {elapsed:.2f}s")

    # Early Stopping
    if val_loss < best_val_loss:
        best_val_loss = val_loss
        patience_counter = 0
        torch.save(model.state_dict(), f"fold_{fold+1}_best.pth")
    else:
        patience_counter += 1
        if patience_counter >= patience:
            print("Early stopping triggered!")
            break

# =====
# Evaluation on Validation Set
# =====
model.load_state_dict(torch.load(f"fold_{fold+1}_best.pth"))
model.eval()
y_true, y_pred = [], []

with torch.no_grad():
    for imgs, labels in val_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        outputs = model(imgs)
        preds = outputs.argmax(1)
        y_true.extend(labels.cpu().numpy())
        y_pred.extend(preds.cpu().numpy())

acc = accuracy_score(y_true, y_pred)
fold_accuracies.append(acc)
print(f"Fold {fold+1} Accuracy: {acc:.4f}")

# =====
# MEMORY CLEANUP
# =====

```

```
del model, optimizer, scheduler, train_loader, val_loader, train_subset, val_subset
torch.cuda.empty_cache()
gc.collect()

# =====
# RESULTS SUMMARY
# =====

mean_acc = np.mean(fold_accuracies)
std_acc = np.std(fold_accuracies)
print("\n===== 5-FOLD CROSS VALIDATION RESULTS =====")
for i, acc in enumerate(fold_accuracies):
    print(f"Fold {i+1}: {acc:.4f}")
print(f"\nMean Accuracy: {mean_acc:.4f}")
print(f"Standard Deviation: {std_acc:.4f}")

===== Fold 1/5 =====
```

```
model.safetensors: 0% | 0.00/37.1M [00:00<?, ?B/s]
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
warnings.warn()
```

```
Epoch [1/35] | Train Acc: 0.7684 | Val Acc: 0.8767 | Time: 14.25s
Epoch [2/35] | Train Acc: 0.9207 | Val Acc: 0.9384 | Time: 14.20s
Epoch [3/35] | Train Acc: 0.9478 | Val Acc: 0.9565 | Time: 14.38s
Epoch [4/35] | Train Acc: 0.9619 | Val Acc: 0.9434 | Time: 14.01s
Epoch [5/35] | Train Acc: 0.9746 | Val Acc: 0.9463 | Time: 13.89s
Epoch [6/35] | Train Acc: 0.9777 | Val Acc: 0.9594 | Time: 13.54s
Epoch [7/35] | Train Acc: 0.9797 | Val Acc: 0.9601 | Time: 14.03s
Epoch [8/35] | Train Acc: 0.9844 | Val Acc: 0.9608 | Time: 13.71s
Early stopping triggered!
Fold 1 Accuracy: 0.9565
```

```
===== Fold 2/5 =====
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
warnings.warn()
```

```
Epoch [1/35] | Train Acc: 0.7731 | Val Acc: 0.9333 | Time: 13.80s
Epoch [2/35] | Train Acc: 0.9469 | Val Acc: 0.9746 | Time: 14.15s
Epoch [3/35] | Train Acc: 0.9755 | Val Acc: 0.9840 | Time: 13.64s
Epoch [4/35] | Train Acc: 0.9842 | Val Acc: 0.9855 | Time: 13.82s
Epoch [5/35] | Train Acc: 0.9931 | Val Acc: 0.9891 | Time: 14.36s
Epoch [6/35] | Train Acc: 0.9960 | Val Acc: 0.9869 | Time: 14.00s
Epoch [7/35] | Train Acc: 0.9942 | Val Acc: 0.9869 | Time: 14.15s
Epoch [8/35] | Train Acc: 0.9964 | Val Acc: 0.9869 | Time: 13.95s
Epoch [9/35] | Train Acc: 0.9971 | Val Acc: 0.9877 | Time: 14.22s
Epoch [10/35] | Train Acc: 0.9973 | Val Acc: 0.9869 | Time: 14.11s
Early stopping triggered!
Fold 2 Accuracy: 0.9891
```

```
===== Fold 3/5 =====
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
warnings.warn()
```

```
Epoch [1/35] | Train Acc: 0.7711 | Val Acc: 0.9267 | Time: 13.74s
Epoch [2/35] | Train Acc: 0.9438 | Val Acc: 0.9557 | Time: 14.05s
Epoch [3/35] | Train Acc: 0.9672 | Val Acc: 0.9739 | Time: 14.04s
Epoch [4/35] | Train Acc: 0.9802 | Val Acc: 0.9855 | Time: 13.76s
Epoch [5/35] | Train Acc: 0.9884 | Val Acc: 0.9826 | Time: 13.96s
Epoch [6/35] | Train Acc: 0.9944 | Val Acc: 0.9855 | Time: 14.12s
Epoch [7/35] | Train Acc: 0.9947 | Val Acc: 0.9826 | Time: 14.01s
Epoch [8/35] | Train Acc: 0.9958 | Val Acc: 0.9884 | Time: 14.07s
Epoch [9/35] | Train Acc: 0.9982 | Val Acc: 0.9898 | Time: 13.63s
Early stopping triggered!
Fold 3 Accuracy: 0.9855
```

```
===== Fold 4/5 =====
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
warnings.warn()
```

```
Epoch [1/35] | Train Acc: 0.7807 | Val Acc: 0.9151 | Time: 13.95s
Epoch [2/35] | Train Acc: 0.9354 | Val Acc: 0.9724 | Time: 13.92s
Epoch [3/35] | Train Acc: 0.9625 | Val Acc: 0.9688 | Time: 14.08s
Epoch [4/35] | Train Acc: 0.9764 | Val Acc: 0.9790 | Time: 14.15s
Epoch [5/35] | Train Acc: 0.9833 | Val Acc: 0.9833 | Time: 14.25s
Epoch [6/35] | Train Acc: 0.9820 | Val Acc: 0.9804 | Time: 14.30s
Epoch [7/35] | Train Acc: 0.9873 | Val Acc: 0.9840 | Time: 14.01s
Epoch [8/35] | Train Acc: 0.9886 | Val Acc: 0.9731 | Time: 13.65s
Epoch [9/35] | Train Acc: 0.9908 | Val Acc: 0.9848 | Time: 13.61s
Epoch [10/35] | Train Acc: 0.9933 | Val Acc: 0.9855 | Time: 14.00s
Early stopping triggered!
Fold 4 Accuracy: 0.9833
```

```
===== Fold 5/5 =====
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de  
precated. Please use get_last_lr() to access the learning rate.  
    warnings.warn(
```

```
Epoch [1/35] | Train Acc: 0.7691 | Val Acc: 0.9180 | Time: 13.87s  
Epoch [2/35] | Train Acc: 0.9483 | Val Acc: 0.9485 | Time: 13.58s  
Epoch [3/35] | Train Acc: 0.9723 | Val Acc: 0.9565 | Time: 14.10s  
Epoch [4/35] | Train Acc: 0.9875 | Val Acc: 0.9739 | Time: 13.81s  
Epoch [5/35] | Train Acc: 0.9877 | Val Acc: 0.9768 | Time: 13.67s  
Epoch [6/35] | Train Acc: 0.9726 | Val Acc: 0.9485 | Time: 14.22s  
Epoch [7/35] | Train Acc: 0.9657 | Val Acc: 0.9731 | Time: 14.27s  
Epoch [8/35] | Train Acc: 0.9726 | Val Acc: 0.9731 | Time: 13.90s  
Epoch [9/35] | Train Acc: 0.9819 | Val Acc: 0.9710 | Time: 13.92s  
Epoch [10/35] | Train Acc: 0.9828 | Val Acc: 0.9695 | Time: 13.83s  
Epoch [11/35] | Train Acc: 0.9918 | Val Acc: 0.9652 | Time: 13.86s  
Epoch [12/35] | Train Acc: 0.9893 | Val Acc: 0.9695 | Time: 13.83s  
Early stopping triggered!  
Fold 5 Accuracy: 0.9731
```

```
===== 5-FOLD CROSS VALIDATION RESULTS =====
```

```
Fold 1: 0.9565  
Fold 2: 0.9891  
Fold 3: 0.9855  
Fold 4: 0.9833  
Fold 5: 0.9731
```

```
Mean Accuracy: 0.9775
```

```
Standard Deviation: 0.0118
```

```
In [14]:
```

```
import torch  
import torch.nn as nn  
import torch.optim as optim  
from torch.utils.data import DataLoader, Subset  
from torchvision import datasets, transforms  
from timm import create_model  
from torch.optim.lr_scheduler import ReduceLROnPlateau  
from sklearn.model_selection import KFold  
from sklearn.metrics import accuracy_score  
import numpy as np  
import time  
import gc  
  
# =====  
# CONFIGURATION  
# =====  
train_dir = "/kaggle/working/cotton_train_aug"  
val_dir = "/kaggle/working/cotton_split/val"  
test_dir = "/kaggle/working/cotton_split/test"  
  
num_classes = 4  
batch_size = 32  
num_epochs = 35  
patience = 5  
num_folds = 5  
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
class_names = ["diseased cotton leaf", "diseased cotton plant", "fresh cotton leaf", "fresh cotton plant"]  
  
# =====  
# DATASET & TRANSFORMS  
# =====  
transform = transforms.Compose([  
    transforms.Resize((224, 224)),  
    transforms.ToTensor(),  
    transforms.Normalize([0.5]*3, [0.5]*3)  
])  
  
dataset = datasets.ImageFolder(root=train_dir, transform=transform)  
  
# =====  
# 5-FOLD CROSS VALIDATION  
# =====  
kf = KFold(n_splits=num_folds, shuffle=True, random_state=42)  
fold_accuracies = []  
  
for fold, (train_idx, val_idx) in enumerate(kf.split(dataset)):  
    print(f"\n===== Fold {fold+1}/{num_folds} =====")  
  
    train_subset = Subset(dataset, train_idx)  
    val_subset = Subset(dataset, val_idx)  
  
    train_loader = DataLoader(train_subset, batch_size=batch_size, shuffle=True, num_workers=2)  
    val_loader = DataLoader(val_subset, batch_size=batch_size, shuffle=False, num_workers=2)
```

```

# =====
# Model Setup
# =====
model = create_model("deit_tiny_patch16_224", pretrained=True, num_classes=num_classes)
model = model.to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)
scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=False)

best_val_loss = float("inf")
patience_counter = 0

# =====
# Training
# =====
for epoch in range(num_epochs):
    start_time = time.time()
    model.train()
    running_loss, correct = 0, 0

    for imgs, labels in train_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(imgs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item() * imgs.size(0)
        correct += (outputs.argmax(1) == labels).sum().item()

    train_loss = running_loss / len(train_loader.dataset)
    train_acc = correct / len(train_loader.dataset)

    # ---- Validation ----
    model.eval()
    val_loss, val_correct = 0, 0
    with torch.no_grad():
        for imgs, labels in val_loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            loss = criterion(outputs, labels)

            val_loss += loss.item() * imgs.size(0)
            val_correct += (outputs.argmax(1) == labels).sum().item()

    val_loss /= len(val_loader.dataset)
    val_acc = val_correct / len(val_loader.dataset)
    scheduler.step(val_loss)

    elapsed = time.time() - start_time
    print(f"Epoch [{epoch+1}/{num_epochs}] | Train Acc: {train_acc:.4f} | Val Acc: {val_acc:.4f} | Time: {elapsed:.2f}s")

    # Early Stopping
    if val_loss < best_val_loss:
        best_val_loss = val_loss
        patience_counter = 0
        torch.save(model.state_dict(), f"fold_{fold+1}_best.pth")
    else:
        patience_counter += 1
        if patience_counter >= patience:
            print("Early stopping triggered!")
            break

# =====
# Evaluation on Validation Set
# =====
model.load_state_dict(torch.load(f"fold_{fold+1}_best.pth"))
model.eval()
y_true, y_pred = [], []

with torch.no_grad():
    for imgs, labels in val_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        outputs = model(imgs)
        preds = outputs.argmax(1)
        y_true.extend(labels.cpu().numpy())
        y_pred.extend(preds.cpu().numpy())

acc = accuracy_score(y_true, y_pred)
fold_accuracies.append(acc)
print(f"Fold {fold+1} Accuracy: {acc:.4f}")

```

```

# =====
# MEMORY CLEANUP
# =====
del model, optimizer, scheduler, train_loader, val_loader, train_subset, val_subset
torch.cuda.empty_cache()
gc.collect()

# =====
# RESULTS SUMMARY
# =====
mean_acc = np.mean(fold_accuracies)
std_acc = np.std(fold_accuracies)
print("\n===== 5-FOLD CROSS VALIDATION RESULTS =====")
for i, acc in enumerate(fold_accuracies):
    print(f"Fold {i+1}: {acc:.4f}")
print(f"\nMean Accuracy: {mean_acc:.4f}")
print(f"Standard Deviation: {std_acc:.4f}")

===== Fold 1/5 =====
model.safetensors:  0%|          | 0.00/22.9M [00:00<?, ?B/s]
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
precated. Please use get_last_lr() to access the learning rate.
    warnings.warn(
Epoch [1/35] | Train Acc: 0.9356 | Val Acc: 0.9761 | Time: 17.10s
Epoch [2/35] | Train Acc: 0.9906 | Val Acc: 0.9891 | Time: 17.08s
Epoch [3/35] | Train Acc: 0.9964 | Val Acc: 0.9434 | Time: 17.12s
Epoch [4/35] | Train Acc: 0.9868 | Val Acc: 0.9659 | Time: 17.15s
Epoch [5/35] | Train Acc: 0.9976 | Val Acc: 0.9913 | Time: 17.16s
Epoch [6/35] | Train Acc: 0.9995 | Val Acc: 0.9920 | Time: 17.06s
Epoch [7/35] | Train Acc: 1.0000 | Val Acc: 0.9906 | Time: 17.31s
Epoch [8/35] | Train Acc: 0.9893 | Val Acc: 0.9877 | Time: 17.36s
Epoch [9/35] | Train Acc: 0.9993 | Val Acc: 0.9927 | Time: 17.05s
Epoch [10/35] | Train Acc: 1.0000 | Val Acc: 0.9935 | Time: 17.27s
Epoch [11/35] | Train Acc: 1.0000 | Val Acc: 0.9935 | Time: 17.26s
Early stopping triggered!
Fold 1 Accuracy: 0.9920

===== Fold 2/5 =====
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
precated. Please use get_last_lr() to access the learning rate.
    warnings.warn(
Epoch [1/35] | Train Acc: 0.9391 | Val Acc: 0.9797 | Time: 17.28s
Epoch [2/35] | Train Acc: 0.9864 | Val Acc: 0.9906 | Time: 17.25s
Epoch [3/35] | Train Acc: 0.9929 | Val Acc: 0.9906 | Time: 17.12s
Epoch [4/35] | Train Acc: 0.9955 | Val Acc: 0.9630 | Time: 17.25s
Epoch [5/35] | Train Acc: 0.9964 | Val Acc: 0.9869 | Time: 17.31s
Epoch [6/35] | Train Acc: 0.9991 | Val Acc: 0.9942 | Time: 17.38s
Epoch [7/35] | Train Acc: 1.0000 | Val Acc: 0.9942 | Time: 17.31s
Epoch [8/35] | Train Acc: 1.0000 | Val Acc: 0.9942 | Time: 17.27s
Epoch [9/35] | Train Acc: 1.0000 | Val Acc: 0.9942 | Time: 17.11s
Epoch [10/35] | Train Acc: 1.0000 | Val Acc: 0.9942 | Time: 17.33s
Epoch [11/35] | Train Acc: 1.0000 | Val Acc: 0.9949 | Time: 17.14s
Epoch [12/35] | Train Acc: 1.0000 | Val Acc: 0.9949 | Time: 17.17s
Epoch [13/35] | Train Acc: 1.0000 | Val Acc: 0.9949 | Time: 17.09s
Early stopping triggered!
Fold 2 Accuracy: 0.9942

===== Fold 3/5 =====
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
precated. Please use get_last_lr() to access the learning rate.
    warnings.warn(
Epoch [1/35] | Train Acc: 0.9284 | Val Acc: 0.9717 | Time: 17.43s
Epoch [2/35] | Train Acc: 0.9857 | Val Acc: 0.9833 | Time: 17.18s
Epoch [3/35] | Train Acc: 0.9964 | Val Acc: 0.9790 | Time: 17.25s
Epoch [4/35] | Train Acc: 0.9924 | Val Acc: 0.9833 | Time: 17.19s
Epoch [5/35] | Train Acc: 0.9938 | Val Acc: 0.9884 | Time: 17.37s
Epoch [6/35] | Train Acc: 0.9949 | Val Acc: 0.9949 | Time: 17.18s
Epoch [7/35] | Train Acc: 0.9956 | Val Acc: 0.9920 | Time: 17.33s
Epoch [8/35] | Train Acc: 0.9995 | Val Acc: 0.9913 | Time: 17.39s
Epoch [9/35] | Train Acc: 1.0000 | Val Acc: 0.9891 | Time: 17.41s
Epoch [10/35] | Train Acc: 1.0000 | Val Acc: 0.9898 | Time: 17.36s
Epoch [11/35] | Train Acc: 1.0000 | Val Acc: 0.9906 | Time: 17.21s
Early stopping triggered!
Fold 3 Accuracy: 0.9949

===== Fold 4/5 =====
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
precated. Please use get_last_lr() to access the learning rate.
    warnings.warn(

```

```
Epoch [1/35] | Train Acc: 0.9280 | Val Acc: 0.9855 | Time: 17.13s
Epoch [2/35] | Train Acc: 0.9909 | Val Acc: 0.9623 | Time: 17.20s
Epoch [3/35] | Train Acc: 0.9873 | Val Acc: 0.9855 | Time: 17.32s
Epoch [4/35] | Train Acc: 0.9913 | Val Acc: 0.9862 | Time: 17.12s
Epoch [5/35] | Train Acc: 0.9976 | Val Acc: 0.9819 | Time: 17.26s
Epoch [6/35] | Train Acc: 0.9991 | Val Acc: 0.9906 | Time: 17.18s
Epoch [7/35] | Train Acc: 0.9975 | Val Acc: 0.9920 | Time: 17.33s
Epoch [8/35] | Train Acc: 0.9859 | Val Acc: 0.9848 | Time: 17.25s
Epoch [9/35] | Train Acc: 0.9971 | Val Acc: 0.9978 | Time: 17.40s
Epoch [10/35] | Train Acc: 0.9993 | Val Acc: 0.9906 | Time: 17.30s
Epoch [11/35] | Train Acc: 0.9958 | Val Acc: 0.9956 | Time: 17.17s
Epoch [12/35] | Train Acc: 0.9998 | Val Acc: 0.9956 | Time: 17.28s
Epoch [13/35] | Train Acc: 1.0000 | Val Acc: 0.9971 | Time: 17.29s
Epoch [14/35] | Train Acc: 1.0000 | Val Acc: 0.9971 | Time: 17.27s
Early stopping triggered!
Fold 4 Accuracy: 0.9978
```

```
===== Fold 5/5 =====
```

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
  warnings.warn(
```

```
Epoch [1/35] | Train Acc: 0.9343 | Val Acc: 0.9819 | Time: 17.20s
Epoch [2/35] | Train Acc: 0.9906 | Val Acc: 0.9840 | Time: 17.31s
Epoch [3/35] | Train Acc: 0.9942 | Val Acc: 0.9906 | Time: 17.24s
Epoch [4/35] | Train Acc: 0.9918 | Val Acc: 0.9848 | Time: 17.17s
Epoch [5/35] | Train Acc: 0.9955 | Val Acc: 0.9790 | Time: 17.31s
Epoch [6/35] | Train Acc: 0.9995 | Val Acc: 0.9942 | Time: 17.18s
Epoch [7/35] | Train Acc: 0.9946 | Val Acc: 0.9826 | Time: 17.18s
Epoch [8/35] | Train Acc: 0.9980 | Val Acc: 0.9949 | Time: 17.13s
Epoch [9/35] | Train Acc: 0.9998 | Val Acc: 0.9949 | Time: 17.31s
Epoch [10/35] | Train Acc: 1.0000 | Val Acc: 0.9949 | Time: 17.15s
Epoch [11/35] | Train Acc: 1.0000 | Val Acc: 0.9949 | Time: 17.22s
Early stopping triggered!
Fold 5 Accuracy: 0.9942
```

```
===== 5-FOLD CROSS VALIDATION RESULTS =====
```

```
Fold 1: 0.9920
Fold 2: 0.9942
Fold 3: 0.9949
Fold 4: 0.9978
Fold 5: 0.9942
```

```
Mean Accuracy: 0.9946
```

```
Standard Deviation: 0.0019
```

```
In [15]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, Subset
from torchvision import datasets, transforms
from timm import create_model
from torch.optim.lr_scheduler import ReduceLROnPlateau
from sklearn.model_selection import KFold
from sklearn.metrics import accuracy_score
import numpy as np
import time
import gc

# =====#
# CONFIGURATION
# =====#
train_dir = "/kaggle/working/cotton_train_aug"
val_dir = "/kaggle/working/cotton_split/val"
test_dir = "/kaggle/working/cotton_split/test"

num_classes = 4
batch_size = 32
num_epochs = 35
patience = 5
num_folds = 5
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
class_names = ["diseased cotton leaf", "diseased cotton plant", "fresh cotton leaf", "fresh cotton plant"]

# =====#
# DATASET & TRANSFORMS
# =====#
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])

dataset = datasets.ImageFolder(root=train_dir, transform=transform)
```

```

# =====
# 5-FOLD CROSS VALIDATION
# =====
kf = KFold(n_splits=num_folds, shuffle=True, random_state=42)
fold_accuracies = []

for fold, (train_idx, val_idx) in enumerate(kf.split(dataset)):
    print(f"\n===== Fold {fold+1}/{num_folds} =====")

    train_subset = Subset(dataset, train_idx)
    val_subset = Subset(dataset, val_idx)

    train_loader = DataLoader(train_subset, batch_size=batch_size, shuffle=True, num_workers=2)
    val_loader = DataLoader(val_subset, batch_size=batch_size, shuffle=False, num_workers=2)

    # =====
    # Model Setup
    # =====
    model = create_model("swin_tiny_patch4_window7_224", pretrained=True, num_classes=num_classes)
    model = model.to(device)

    criterion = nn.CrossEntropyLoss()
    optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)
    scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=False)

    best_val_loss = float("inf")
    patience_counter = 0

    # =====
    # Training
    # =====
    for epoch in range(num_epochs):
        start_time = time.time()
        model.train()
        running_loss, correct = 0, 0

        for imgs, labels in train_loader:
            imgs, labels = imgs.to(device), labels.to(device)
            optimizer.zero_grad()
            outputs = model(imgs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            running_loss += loss.item() * imgs.size(0)
            correct += (outputs.argmax(1) == labels).sum().item()

        train_loss = running_loss / len(train_loader.dataset)
        train_acc = correct / len(train_loader.dataset)

        # ---- Validation ----
        model.eval()
        val_loss, val_correct = 0, 0
        with torch.no_grad():
            for imgs, labels in val_loader:
                imgs, labels = imgs.to(device), labels.to(device)
                outputs = model(imgs)
                loss = criterion(outputs, labels)

                val_loss += loss.item() * imgs.size(0)
                val_correct += (outputs.argmax(1) == labels).sum().item()

        val_loss /= len(val_loader.dataset)
        val_acc = val_correct / len(val_loader.dataset)
        scheduler.step(val_loss)

        elapsed = time.time() - start_time
        print(f"Epoch [{epoch+1}/{num_epochs}] | Train Acc: {train_acc:.4f} | Val Acc: {val_acc:.4f} | Time: {elapsed:.2f}s")

        # Early Stopping
        if val_loss < best_val_loss:
            best_val_loss = val_loss
            patience_counter = 0
            torch.save(model.state_dict(), f"fold_{fold+1}_best.pth")
        else:
            patience_counter += 1
            if patience_counter >= patience:
                print("Early stopping triggered!")
                break

    # =====
    # Evaluation on Validation Set

```

```

# =====
model.load_state_dict(torch.load(f"fold_{fold+1}_best.pth"))
model.eval()
y_true, y_pred = [], []

with torch.no_grad():
    for imgs, labels in val_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        outputs = model(imgs)
        preds = outputs.argmax(1)
        y_true.extend(labels.cpu().numpy())
        y_pred.extend(preds.cpu().numpy())

acc = accuracy_score(y_true, y_pred)
fold_accuracies.append(acc)
print(f"Fold {fold+1} Accuracy: {acc:.4f}")

# =====
# MEMORY CLEANUP
# =====
del model, optimizer, scheduler, train_loader, val_loader, train_subset, val_subset
torch.cuda.empty_cache()
gc.collect()

# =====
# RESULTS SUMMARY
# =====
mean_acc = np.mean(fold_accuracies)
std_acc = np.std(fold_accuracies)
print("\n===== 5-FOLD CROSS VALIDATION RESULTS =====")
for i, acc in enumerate(fold_accuracies):
    print(f"Fold {i+1}: {acc:.4f}")
print(f"\nMean Accuracy: {mean_acc:.4f}")
print(f"Standard Deviation: {std_acc:.4f}")

```

===== Fold 1/5 =====

model.safetensors: 0%| 0.00/114M [00:00<?, ?B/s]

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
warnings.warn(

Epoch [1/35] | Train Acc: 0.9407 | Val Acc: 0.9840 | Time: 48.34s
Epoch [2/35] | Train Acc: 0.9891 | Val Acc: 0.9920 | Time: 48.38s
Epoch [3/35] | Train Acc: 0.9975 | Val Acc: 0.9971 | Time: 48.19s
Epoch [4/35] | Train Acc: 0.9929 | Val Acc: 0.9811 | Time: 48.35s
Epoch [5/35] | Train Acc: 0.9918 | Val Acc: 0.9949 | Time: 48.20s
Epoch [6/35] | Train Acc: 0.9996 | Val Acc: 0.9942 | Time: 48.15s
Epoch [7/35] | Train Acc: 1.0000 | Val Acc: 0.9964 | Time: 48.17s
Epoch [8/35] | Train Acc: 0.9995 | Val Acc: 0.9964 | Time: 48.21s
Early stopping triggered!

Fold 1 Accuracy: 0.9971

===== Fold 2/5 =====

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
warnings.warn(

Epoch [1/35] | Train Acc: 0.9231 | Val Acc: 0.9927 | Time: 48.31s
Epoch [2/35] | Train Acc: 0.9891 | Val Acc: 0.9964 | Time: 48.18s
Epoch [3/35] | Train Acc: 0.9931 | Val Acc: 0.9920 | Time: 48.16s
Epoch [4/35] | Train Acc: 0.9891 | Val Acc: 0.9884 | Time: 48.16s
Epoch [5/35] | Train Acc: 0.9967 | Val Acc: 0.9971 | Time: 48.10s
Epoch [6/35] | Train Acc: 0.9989 | Val Acc: 0.9993 | Time: 48.17s
Epoch [7/35] | Train Acc: 0.9996 | Val Acc: 0.9985 | Time: 48.26s
Epoch [8/35] | Train Acc: 0.9900 | Val Acc: 0.9848 | Time: 48.08s
Epoch [9/35] | Train Acc: 0.9980 | Val Acc: 0.9985 | Time: 48.14s
Epoch [10/35] | Train Acc: 0.9996 | Val Acc: 0.9985 | Time: 48.11s
Epoch [11/35] | Train Acc: 0.9971 | Val Acc: 0.9891 | Time: 48.14s
Epoch [12/35] | Train Acc: 0.9993 | Val Acc: 0.9964 | Time: 48.16s
Epoch [13/35] | Train Acc: 1.0000 | Val Acc: 0.9978 | Time: 48.09s
Epoch [14/35] | Train Acc: 1.0000 | Val Acc: 0.9985 | Time: 48.07s
Early stopping triggered!

Fold 2 Accuracy: 0.9985

===== Fold 3/5 =====

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
warnings.warn(

```
Epoch [1/35] | Train Acc: 0.9295 | Val Acc: 0.9819 | Time: 48.39s
Epoch [2/35] | Train Acc: 0.9897 | Val Acc: 0.9920 | Time: 48.14s
Epoch [3/35] | Train Acc: 0.9895 | Val Acc: 0.9913 | Time: 48.26s
Epoch [4/35] | Train Acc: 0.9964 | Val Acc: 0.9811 | Time: 48.18s
Epoch [5/35] | Train Acc: 0.9960 | Val Acc: 0.9906 | Time: 48.19s
Epoch [6/35] | Train Acc: 0.9991 | Val Acc: 0.9942 | Time: 48.24s
Epoch [7/35] | Train Acc: 1.0000 | Val Acc: 0.9942 | Time: 48.28s
Epoch [8/35] | Train Acc: 0.9998 | Val Acc: 0.9935 | Time: 48.28s
Epoch [9/35] | Train Acc: 0.9956 | Val Acc: 0.9884 | Time: 48.14s
Epoch [10/35] | Train Acc: 0.9996 | Val Acc: 0.9949 | Time: 48.31s
Epoch [11/35] | Train Acc: 0.9998 | Val Acc: 0.9956 | Time: 48.20s
Epoch [12/35] | Train Acc: 0.9998 | Val Acc: 0.9942 | Time: 48.23s
Epoch [13/35] | Train Acc: 1.0000 | Val Acc: 0.9949 | Time: 48.29s
Epoch [14/35] | Train Acc: 1.0000 | Val Acc: 0.9949 | Time: 48.21s
Epoch [15/35] | Train Acc: 1.0000 | Val Acc: 0.9956 | Time: 48.21s
Epoch [16/35] | Train Acc: 0.9998 | Val Acc: 0.9964 | Time: 48.27s
Early stopping triggered!
Fold 3 Accuracy: 0.9956
```

===== Fold 4/5 =====

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
preacted. Please use get_last_lr() to access the learning rate.
    warnings.warn(
Epoch [1/35] | Train Acc: 0.9207 | Val Acc: 0.9884 | Time: 48.26s
Epoch [2/35] | Train Acc: 0.9884 | Val Acc: 0.9949 | Time: 48.35s
Epoch [3/35] | Train Acc: 0.9917 | Val Acc: 0.9906 | Time: 48.34s
Epoch [4/35] | Train Acc: 0.9931 | Val Acc: 0.9913 | Time: 48.23s
Epoch [5/35] | Train Acc: 0.9973 | Val Acc: 0.9978 | Time: 48.32s
Epoch [6/35] | Train Acc: 0.9967 | Val Acc: 0.9978 | Time: 48.18s
Epoch [7/35] | Train Acc: 0.9975 | Val Acc: 0.9949 | Time: 48.21s
Epoch [8/35] | Train Acc: 0.9971 | Val Acc: 0.9985 | Time: 48.30s
Epoch [9/35] | Train Acc: 0.9996 | Val Acc: 0.9993 | Time: 48.29s
Epoch [10/35] | Train Acc: 1.0000 | Val Acc: 0.9985 | Time: 48.13s
Epoch [11/35] | Train Acc: 1.0000 | Val Acc: 0.9993 | Time: 48.30s
Epoch [12/35] | Train Acc: 1.0000 | Val Acc: 0.9993 | Time: 48.32s
Epoch [13/35] | Train Acc: 1.0000 | Val Acc: 0.9993 | Time: 48.23s
Epoch [14/35] | Train Acc: 1.0000 | Val Acc: 0.9993 | Time: 48.44s
Early stopping triggered!
Fold 4 Accuracy: 0.9993
```

===== Fold 5/5 =====

```
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is de
preacted. Please use get_last_lr() to access the learning rate.
    warnings.warn(
Epoch [1/35] | Train Acc: 0.9236 | Val Acc: 0.9913 | Time: 48.20s
Epoch [2/35] | Train Acc: 0.9859 | Val Acc: 0.9898 | Time: 48.17s
Epoch [3/35] | Train Acc: 0.9859 | Val Acc: 0.9927 | Time: 48.12s
Epoch [4/35] | Train Acc: 0.9966 | Val Acc: 0.9964 | Time: 48.17s
Epoch [5/35] | Train Acc: 0.9998 | Val Acc: 0.9964 | Time: 48.13s
Epoch [6/35] | Train Acc: 0.9969 | Val Acc: 0.9753 | Time: 48.33s
Epoch [7/35] | Train Acc: 0.9906 | Val Acc: 0.9949 | Time: 48.09s
Epoch [8/35] | Train Acc: 0.9940 | Val Acc: 0.9935 | Time: 48.17s
Epoch [9/35] | Train Acc: 0.9987 | Val Acc: 0.9971 | Time: 48.21s
Epoch [10/35] | Train Acc: 0.9980 | Val Acc: 0.9971 | Time: 48.22s
Epoch [11/35] | Train Acc: 0.9996 | Val Acc: 0.9978 | Time: 48.23s
Epoch [12/35] | Train Acc: 1.0000 | Val Acc: 0.9985 | Time: 48.27s
Epoch [13/35] | Train Acc: 1.0000 | Val Acc: 0.9985 | Time: 48.21s
Epoch [14/35] | Train Acc: 1.0000 | Val Acc: 0.9985 | Time: 48.20s
Epoch [15/35] | Train Acc: 1.0000 | Val Acc: 0.9978 | Time: 48.14s
Epoch [16/35] | Train Acc: 1.0000 | Val Acc: 0.9985 | Time: 48.18s
Epoch [17/35] | Train Acc: 1.0000 | Val Acc: 0.9985 | Time: 48.24s
Epoch [18/35] | Train Acc: 1.0000 | Val Acc: 0.9978 | Time: 48.21s
Epoch [19/35] | Train Acc: 1.0000 | Val Acc: 0.9985 | Time: 48.28s
Early stopping triggered!
Fold 5 Accuracy: 0.9985
```

===== 5-FOLD CROSS VALIDATION RESULTS =====

```
Fold 1: 0.9971
Fold 2: 0.9985
Fold 3: 0.9956
Fold 4: 0.9993
Fold 5: 0.9985
```

Mean Accuracy: 0.9978

Standard Deviation: 0.0013