

```
In [1]: import kagglehub

# Download latest version
path = kagglehub.dataset_download("sohansakib75/cotton-4-class")

print("Path to dataset files:", path)
```

Path to dataset files: /kaggle/input/cotton-4-class

```
In [2]: import os
import kagglehub

# Download dataset
path = kagglehub.dataset_download("sohansakib75/cotton-4-class")

print("Dataset root path:", path)

# List folders and files inside
print("Contents inside dataset folder:")
for item in os.listdir(path):
    item_path = os.path.join(path, item)
    if os.path.isdir(item_path):
        print(f" {item}/")
    else:
        print(f" {item}")
```

```
Dataset root path: /kaggle/input/cotton-4-class
Contents inside dataset folder:
Cotton leaf/
```

```
In [3]: import os
import glob

dataset_path = os.path.join(path, "Cotton leaf")

print("Path to Dataset folder:", dataset_path)

for subdir in sorted(os.listdir(dataset_path)):
    subpath = os.path.join(dataset_path, subdir)
    if os.path.isdir(subpath):
        # Count images by common formats
        image_files = glob.glob(os.path.join(subpath, "*.jpg")) + \
                      glob.glob(os.path.join(subpath, "*.jpeg")) + \
                      glob.glob(os.path.join(subpath, "*.png"))
        print(f"{subdir}: {len(image_files)} images")
```

```
Path to Dataset folder: /kaggle/input/cotton-4-class/Cotton leaf
diseased cotton leaf: 346 images
diseased cotton plant: 921 images
fresh cotton leaf: 519 images
fresh cotton plant: 514 images
```

```
In [4]: import os
import glob
import matplotlib.pyplot as plt
import random

dataset_path = "/kaggle/input/cotton-4-class/Cotton leaf"

# Classes
classes = sorted(os.listdir(dataset_path))

# Plot 5 images per class
fig, axes = plt.subplots(len(classes), 5, figsize=(15, 12))

for i, cls in enumerate(classes):
    cls_path = os.path.join(dataset_path, cls)
    image_files = glob.glob(os.path.join(cls_path, "*.jpg")) + \
                  glob.glob(os.path.join(cls_path, "*.jpeg")) + \
                  glob.glob(os.path.join(cls_path, "*.png"))

    # Randomly pick 5 images
    sample_files = random.sample(image_files, 5)

    for j, img_path in enumerate(sample_files):
        img = plt.imread(img_path)
        axes[i, j].imshow(img)
        axes[i, j].axis("off")
        if j == 2: # center column
            axes[i, j].set_title(cls, fontsize=10)

plt.tight_layout()
```

```
plt.show()
```



```
In [5]: import os
import cv2
import numpy as np
import matplotlib.pyplot as plt
import random

# Paths
input_dir = "/kaggle/input/cotton-4-class/Cotton leaf"
output_dir = "/kaggle/working/preprocessed_dataset"
os.makedirs(output_dir, exist_ok=True)

img_size = (224, 224)
num_samples = 5

classes = sorted(os.listdir(input_dir))
fig, axes = plt.subplots(len(classes), num_samples, figsize=(15, 3*len(classes)))

clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8,8))

for i, cls in enumerate(classes):
    cls_path = os.path.join(input_dir, cls)
    if os.path.isdir(cls_path):
        save_cls_path = os.path.join(output_dir, cls)
        os.makedirs(save_cls_path, exist_ok=True)

    files = [f for f in os.listdir(cls_path) if f.lower().endswith((".jpg", ".jpeg", ".png"))]
    sample_files = random.sample(files, min(num_samples, len(files)))

    for j, file in enumerate(files):
        img_path = os.path.join(cls_path, file)
        img = cv2.imread(img_path)

        # 1) Resize
        img_resized = cv2.resize(img, img_size)

        # 2) CLAHE on Y channel
        yuv = cv2.cvtColor(img_resized, cv2.COLOR_BGR2YUV)
        yuv[:, :, 0] = clahe.apply(yuv[:, :, 0])
```

```

        img_clahe = cv2.cvtColor(yuv, cv2.COLOR_YUV2BGR)


        # 3. Normalize to [0,1]
        img_final = img_clahe.astype(np.float32) / 255.0

        # Save preprocessed image
        save_img = (img_final * 255).astype(np.uint8)
        cv2.imwrite(os.path.join(save_cls_path, file), save_img)

        # Plot sample images
        if file in sample_files:
            axes[i, sample_files.index(file)].imshow(cv2.cvtColor(save_img, cv2.COLOR_BGR2RGB))
            axes[i, sample_files.index(file)].axis("off")
            if sample_files.index(file) == num_samples // 2:
                axes[i, sample_files.index(file)].set_title(cls, fontsize=12)

plt.tight_layout()
plt.show()
print("\n Preprocessed images saved in:", output_dir)

```



Preprocessed images saved in: /kaggle/working/preprocessed_dataset

In [6]:

```

import os
import shutil
import random

# Preprocessed dataset path
preprocessed_dir = "/kaggle/working/preprocessed_dataset"

# Split paths
split_base = "/kaggle/working/cotton_split"
train_dir = os.path.join(split_base, "train")
val_dir = os.path.join(split_base, "val")
test_dir = os.path.join(split_base, "test")

# Create split folders
for d in [train_dir, val_dir, test_dir]:
    os.makedirs(d, exist_ok=True)

# Split ratios
train_ratio = 0.75

```

```

val_ratio = 0.1
test_ratio = 0.15

classes = sorted(os.listdir(preprocessed_dir))

for cls in classes:
    cls_path = os.path.join(preprocessed_dir, cls)
    files = [f for f in os.listdir(cls_path) if f.lower().endswith((".jpg", ".jpeg", ".png"))]
    random.shuffle(files)

    n_total = len(files)
    n_train = int(train_ratio * n_total)
    n_val = int(val_ratio * n_total)
    n_test = n_total - n_train - n_val

    splits = {
        train_dir: files[:n_train],
        val_dir: files[n_train:n_train+n_val],
        test_dir: files[n_train+n_val:]
    }

    for split_folder, split_files in splits.items():
        cls_split_path = os.path.join(split_folder, cls)
        os.makedirs(cls_split_path, exist_ok=True)
        for f in split_files:
            shutil.copy(os.path.join(cls_path, f), os.path.join(cls_split_path, f))

print(" Dataset split into 75-10-15 and saved in:", split_base)

```

Dataset split into 75-10-15 and saved in: /kaggle/working/cotton_split

```

In [7]: import os

split_base = "/kaggle/working/cotton_split"
splits = ["train", "val", "test"]

for split in splits:
    split_path = os.path.join(split_base, split)
    print(f"\n {split.capitalize()} Split:")
    for cls in sorted(os.listdir(split_path)):
        cls_path = os.path.join(split_path, cls)
        num_images = len([f for f in os.listdir(cls_path) if f.lower().endswith((".jpg", ".jpeg", ".png"))])
        print(f"{cls}: {num_images} images")

Train Split:
diseased cotton leaf: 316 images
diseased cotton plant: 865 images
fresh cotton leaf: 487 images
fresh cotton plant: 481 images

Val Split:
diseased cotton leaf: 64 images
diseased cotton plant: 171 images
fresh cotton leaf: 98 images
fresh cotton plant: 100 images

Test Split:
diseased cotton leaf: 92 images
diseased cotton plant: 263 images
fresh cotton leaf: 145 images
fresh cotton plant: 149 images

```

```

In [8]: import os
import cv2
import numpy as np
import random

train_dir = "/kaggle/working/cotton_split/train"
aug_train_dir = "/kaggle/working/cotton_train_aug"
os.makedirs(aug_train_dir, exist_ok=True)

# Augmentation functions
def random_flip(img):
    flip_code = random.choice([-1, 0, 1])
    return cv2.flip(img, flip_code)

def random_rotate(img):
    angle = random.uniform(-25, 25)
    h, w = img.shape[:2]
    M = cv2.getRotationMatrix2D((w//2, h//2), angle, 1)
    return cv2.warpAffine(img, M, (w, h), borderMode=cv2.BORDER_REFLECT)

def random_zoom(img):
    zoom_factor = random.uniform(0.8, 1.2)

```

```

h, w = img.shape[:2]
new_h, new_w = int(h*zoom_factor), int(w*zoom_factor)
img_resized = cv2.resize(img, (new_w, new_h))
if zoom_factor < 1:
    pad_h = (h - new_h) // 2
    pad_w = (w - new_w) // 2
    img_padded = cv2.copyMakeBorder(img_resized, pad_h, h-new_h-pad_h,
                                    pad_w, w-new_w-pad_w, cv2.BORDER_REFLECT)
    return img_padded
else:
    start_h = (new_h - h)//2
    start_w = (new_w - w)//2
    return img_resized[start_h:start_h+h, start_w:start_w+w]

def random_brightness(img):
    factor = random.uniform(0.7, 1.3)
    img = img.astype(np.float32) * factor
    img = np.clip(img, 0, 255).astype(np.uint8)
    return img

augmentations = [random_flip, random_rotate, random_zoom, random_brightness]

# Apply augmentations
classes = sorted(os.listdir(train_dir))
for cls in classes:
    cls_path = os.path.join(train_dir, cls)
    save_cls_path = os.path.join(aug_train_dir, cls)
    os.makedirs(save_cls_path, exist_ok=True)

    for file in os.listdir(cls_path):
        if not file.lower().endswith((".jpg", ".jpeg", ".png")):
            continue
        img_path = os.path.join(cls_path, file)
        img = cv2.imread(img_path)

        # Save original
        cv2.imwrite(os.path.join(save_cls_path, file), img)

        # 3 random augmentations
        for k in range(3):
            aug_img = img.copy()
            aug_funcs = random.sample(augmentations, 2)
            for func in aug_funcs:
                aug_img = func(aug_img)
            filename, ext = os.path.splitext(file)
            aug_name = f"{filename}_aug{k+1}{ext}"
            cv2.imwrite(os.path.join(save_cls_path, aug_name), aug_img)

# Print image count per class
print("\n Image count per class after augmentation:")
for cls in classes:
    cls_path = os.path.join(aug_train_dir, cls)
    count = len([f for f in os.listdir(cls_path) if f.lower().endswith((".jpg", ".jpeg", ".png"))])
    print(f"{cls}: {count} images")

print(f"\n All train images and augmented images saved in: {aug_train_dir}")

```

Image count per class after augmentation:
diseased cotton leaf: 1264 images
diseased cotton plant: 3460 images
fresh cotton leaf: 1948 images
fresh cotton plant: 1924 images

All train images and augmented images saved in: /kaggle/working/cotton_train_aug

In [9]: !pip install timm

```

Requirement already satisfied: timm in /usr/local/lib/python3.11/dist-packages (1.0.19)
Requirement already satisfied: torch in /usr/local/lib/python3.11/dist-packages (from timm) (2.6.0+cu124)
Requirement already satisfied: torchvision in /usr/local/lib/python3.11/dist-packages (from timm) (0.21.0+cu124)
Requirement already satisfied: pyyaml in /usr/local/lib/python3.11/dist-packages (from timm) (6.0.3)
Requirement already satisfied: huggingface_hub in /usr/local/lib/python3.11/dist-packages (from timm) (1.0.0rc2)
Requirement already satisfied: safetensors in /usr/local/lib/python3.11/dist-packages (from timm) (0.5.3)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm) (3.19.1)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm) (2025.9.0)
Requirement already satisfied: packaging>=20.9 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm) (25.0)
Requirement already satisfied: httpx<1,>=0.23.0 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm) (0.28.1)
Requirement already satisfied: tqdm>=4.42.1 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm) (4.67.1)
Requirement already satisfied: typer-slim in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm)

```

```
) (0.19.2)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm) (4.15.0)
Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->timm) (1.1.10)
Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages (from torch->timm) (3.5)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (3.1.6)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (12.4.127)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (12.4.127)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (12.4.127)
Requirement already satisfied: nvidia-cudnn-cu12==9.1.0.70 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (9.1.0.70)
Requirement already satisfied: nvidia-cublas-cu12==12.4.5.8 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (12.4.5.8)
Requirement already satisfied: nvidia-cufft-cu12==11.2.1.3 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (11.2.1.3)
Requirement already satisfied: nvidia-curand-cu12==10.3.5.147 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (10.3.5.147)
Requirement already satisfied: nvidia-cusolver-cu12==11.6.1.9 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (11.6.1.9)
Requirement already satisfied: nvidia-cusparse-cu12==12.3.1.170 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (12.3.1.170)
Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (0.6.2)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (2.21.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (12.4.127)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (12.4.127)
Requirement already satisfied: triton==3.2.0 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (3.2.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-packages (from torch->timm) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch->timm) (1.3.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from torchvision->timm) (1.26.4)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.11/dist-packages (from torchvision->timm) (11.3.0)
Requirement already satisfied: anyio in /usr/local/lib/python3.11/dist-packages (from httpx<1,>=0.23.0->huggingface_hub->timm) (4.11.0)
Requirement already satisfied: certifi in /usr/local/lib/python3.11/dist-packages (from httpx<1,>=0.23.0->huggingface_hub->timm) (2025.8.3)
Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.11/dist-packages (from httpx<1,>=0.23.0->huggingface_hub->timm) (1.0.9)
Requirement already satisfied: idna in /usr/local/lib/python3.11/dist-packages (from httpx<1,>=0.23.0->huggingface_hub->timm) (3.10)
Requirement already satisfied: h11=>0.16 in /usr/local/lib/python3.11/dist-packages (from httpcore==1.*->httpx<1,>=0.23.0->huggingface_hub->timm) (0.16.0)
Requirement already satisfied: MarkupSafe=>2.0 in /usr/local/lib/python3.11/dist-packages (from jinja2->torch->timm) (3.0.2)
Requirement already satisfied: mkl_fft in /usr/local/lib/python3.11/dist-packages (from numpy->torchvision->timm) (1.3.8)
Requirement already satisfied: mkl_random in /usr/local/lib/python3.11/dist-packages (from numpy->torchvision->timm) (1.2.4)
Requirement already satisfied: mkl_umath in /usr/local/lib/python3.11/dist-packages (from numpy->torchvision->timm) (0.1.1)
Requirement already satisfied: mkl in /usr/local/lib/python3.11/dist-packages (from numpy->torchvision->timm) (2.025.2.0)
Requirement already satisfied: tbb4py in /usr/local/lib/python3.11/dist-packages (from numpy->torchvision->timm) (2022.2.0)
Requirement already satisfied: mkl-service in /usr/local/lib/python3.11/dist-packages (from numpy->torchvision->timm) (2.4.1)
Requirement already satisfied: click=>8.0.0 in /usr/local/lib/python3.11/dist-packages (from typer-slim->huggingface_hub->timm) (8.3.0)
Requirement already satisfied: sniffio=>1.1 in /usr/local/lib/python3.11/dist-packages (from anyio->httpx<1,>=0.23.0->huggingface_hub->timm) (1.3.1)
Requirement already satisfied: intel-openmp<2026,>=2024 in /usr/local/lib/python3.11/dist-packages (from mkl->numpy->torchvision->timm) (2024.2.0)
Requirement already satisfied: tbb==2022.* in /usr/local/lib/python3.11/dist-packages (from mkl->numpy->torchvision->timm) (2022.2.0)
Requirement already satisfied: tcmlib==1.* in /usr/local/lib/python3.11/dist-packages (from tbb==2022.*->mkl->numpy->torchvision->timm) (1.4.0)
Requirement already satisfied: intel-cmplr-lib-rt in /usr/local/lib/python3.11/dist-packages (from mkl_umath->numpy->torchvision->timm) (2024.2.0)
Requirement already satisfied: intel-cmplr-lib-ur==2024.2.0 in /usr/local/lib/python3.11/dist-packages (from intel-openmp<2026,>=2024->mkl->numpy->torchvision->timm) (2024.2.0)
```

In [10]: `import timm`

```
/usr/local/lib/python3.11/dist-packages/pydantic/_internal/_generate_schema.py:2225: UnsupportedFieldAttributeWarning: The 'repr' attribute with value False was provided to the `Field()` function, which has no effect in the context it was used. 'repr' is field-specific metadata, and can only be attached to a model field using `Annotated` metadata or by assignment. This may have happened because an `Annotated` type alias using the `type` statement was used, or if the `Field()` function was attached to a single member of a union type.  
    warnings.warn()  
/usr/local/lib/python3.11/dist-packages/pydantic/_internal/_generate_schema.py:2225: UnsupportedFieldAttributeWarning: The 'frozen' attribute with value True was provided to the `Field()` function, which has no effect in the context it was used. 'frozen' is field-specific metadata, and can only be attached to a model field using `Annotated` metadata or by assignment. This may have happened because an `Annotated` type alias using the `type` statement was used, or if the `Field()` function was attached to a single member of a union type.  
    warnings.warn()
```

```
In [11]:  
import torch  
import torch.nn as nn  
import torch.optim as optim  
from torch.utils.data import DataLoader  
from torchvision import datasets, transforms  
from timm import create_model  
from torch.optim.lr_scheduler import ReduceLROnPlateau  
  
from sklearn.metrics import (  
    classification_report, confusion_matrix, matthews_corrcoef,  
    roc_auc_score, average_precision_score, roc_curve,  
    precision_recall_curve, cohen_kappa_score  
)  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
import time  
  
# ======  
# Config for Cotton Dataset  
# ======  
train_dir = "/kaggle/working/cotton_train_aug" # Augmented train  
val_dir = "/kaggle/working/cotton_split/val"  
test_dir = "/kaggle/working/cotton_split/test"  
  
batch_size = 32  
num_epochs = 35  
patience = 5  
num_classes = 4  
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
class_names = ["diseased cotton leaf", "diseased cotton plant", "fresh cotton leaf", "fresh cotton plant"]  
  
# ======  
# Data Transforms & Dataloaders  
# ======  
common_tfms = transforms.Compose([  
    transforms.Resize((224, 224)),  
    transforms.ToTensor(),  
    transforms.Normalize([0.5]*3, [0.5]*3)  
)  
  
train_ds = datasets.ImageFolder(root=train_dir, transform=common_tfms)  
val_ds = datasets.ImageFolder(root=val_dir, transform=common_tfms)  
test_ds = datasets.ImageFolder(root=test_dir, transform=common_tfms)  
  
train_loader = DataLoader(train_ds, batch_size=batch_size, shuffle=True, num_workers=2)  
val_loader = DataLoader(val_ds, batch_size=batch_size, shuffle=False, num_workers=2)  
test_loader = DataLoader(test_ds, batch_size=batch_size, shuffle=False, num_workers=2)  
  
# ======  
# Model  
# ======  
model = create_model("mobilevit_s", pretrained=True, num_classes=num_classes)  
model = model.to(device)  
  
criterion = nn.CrossEntropyLoss()  
optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)  
scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=True)  
  
# ======  
# Training Loop with Early Stopping  
# ======  
best_val_loss = float("inf")  
patience_counter = 0  
  
for epoch in range(num_epochs):  
    start_time = time.time()  
  
    # ---- Train ----  
    model.train()
```

```

train_loss, train_correct = 0, 0
for imgs, labels in train_loader:
    imgs, labels = imgs.to(device), labels.to(device)

    optimizer.zero_grad()
    outputs = model(imgs)
    loss = criterion(outputs, labels)
    loss.backward()
    optimizer.step()

    train_loss += loss.item() * imgs.size(0)
    train_correct += (outputs.argmax(1) == labels).sum().item()

train_loss /= len(train_loader.dataset)
train_acc = train_correct / len(train_loader.dataset)

# ---- Validation ----
model.eval()
val_loss, val_correct = 0, 0
with torch.no_grad():
    for imgs, labels in val_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        outputs = model(imgs)
        loss = criterion(outputs, labels)

        val_loss += loss.item() * imgs.size(0)
        val_correct += (outputs.argmax(1) == labels).sum().item()

val_loss /= len(val_loader.dataset)
val_acc = val_correct / len(val_loader.dataset)

scheduler.step(val_loss)

elapsed = time.time() - start_time
print(f"Epoch [{epoch+1}/{num_epochs}] "
      f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f} "
      f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f} "
      f"Time: {elapsed:.2f}s")

# Early Stopping
if val_loss < best_val_loss:
    best_val_loss = val_loss
    patience_counter = 0
    torch.save(model.state_dict(), "mobilevit_cotton_best.pth")
else:
    patience_counter += 1
    if patience_counter >= patience:
        print("Early stopping triggered!")
        break

print("Training finished Best model saved as mobilevit_cotton_best.pth")

# =====
# Evaluation Functions
# =====
def plot_confusion_matrix(cm, classes, title="Confusion Matrix"):
    plt.figure(figsize=(6,5))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=classes, yticklabels=classes)
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title(title)
    plt.show()

def print_report(y_true, y_pred, y_prob, title="", plot_curves=False):
    print(f"\n--- {title} ---")
    print(classification_report(y_true, y_pred, digits=4))

    cm = confusion_matrix(y_true, y_pred)
    plot_confusion_matrix(cm, class_names, title=f"{title} Confusion Matrix")

    mcc = matthews_corrcoef(y_true, y_pred)
    print(f"MCC: {mcc:.4f}")

    kappa = cohen_kappa_score(y_true, y_pred)
    print(f"Cohen's Kappa: {kappa:.4f}")

# Class-wise NPV + PPV
npv_list, ppv_list = [], []
for i in range(len(cm)):
    TN = np.sum(np.delete(np.delete(cm, i, axis=0), i, axis=1))
    FN = np.sum(cm[i, :]) - cm[i, i]
    FP = np.sum(cm[:, i]) - cm[i, i]

```

```

TP = cm[i, i]

NPV = TN / (TN + FN) if (TN + FN) > 0 else 0
PPV = TP / (TP + FP) if (TP + FP) > 0 else 0
npv_list.append(NPV)
ppv_list.append(PPV)

print(f"Mean NPV: {np.mean(npv_list):.4f}")
print(f"Mean PPV (Precision): {np.mean(ppv_list):.4f}")

if plot_curves:
    y_onehot = np.eye(num_classes)[y_true]
    roc_auc = roc_auc_score(y_onehot, y_prob, average='macro', multi_class='ovr')
    pr_auc = average_precision_score(y_onehot, y_prob, average='macro')
    print(f"ROC AUC: {roc_auc:.4f}, PR AUC: {pr_auc:.4f}")

    # ROC Curve
    plt.figure(figsize=(6,5))
    for i in range(num_classes):
        fpr, tpr, _ = roc_curve(y_onehot[:,i], y_prob[:,i])
        plt.plot(fpr, tpr, label=f"{class_names[i]} (AUC={roc_auc_score(y_onehot[:,i], y_prob[:,i]):.4f})")
    plt.plot([0,1],[0,1],'k--')
    plt.title(f"{title} ROC Curve")
    plt.xlabel("FPR")
    plt.ylabel("TPR")
    plt.legend()
    plt.show()

    # PR Curve
    plt.figure(figsize=(6,5))
    for i in range(num_classes):
        precision, recall, _ = precision_recall_curve(y_onehot[:,i], y_prob[:,i])
        plt.plot(recall, precision, label=f"{class_names[i]} (AP={average_precision_score(y_onehot[:,i], y_prob[:,i]):.4f})")
    plt.title(f"{title} Precision-Recall Curve")
    plt.xlabel("Recall")
    plt.ylabel("Precision")
    plt.legend()
    plt.show()

def evaluate_model(model, loader, title="", plot_curves=False):
    model.eval()
    y_true, y_pred, y_prob = [], [], []
    start_time = time.time()

    with torch.no_grad():
        for imgs, labels in loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            probs = torch.softmax(outputs, dim=1)
            preds = outputs.argmax(1)

            y_true.extend(labels.cpu().numpy())
            y_pred.extend(preds.cpu().numpy())
            y_prob.extend(probs.cpu().numpy())

    infer_time = time.time() - start_time
    y_true, y_pred, y_prob = np.array(y_true), np.array(y_pred), np.array(y_prob)

    print_report(y_true, y_pred, y_prob, title, plot_curves)
    print(f"{title} inference time: {infer_time:.2f} sec")
    return y_true, y_pred, y_prob, infer_time

# =====
# Final Evaluation
# =====
model.load_state_dict(torch.load("mobilevit_cotton_best.pth"))

y_true_train, y_pred_train, y_prob_train, train_infer_time = evaluate_model(model, train_loader, "Train Set", plot_curves)
y_true_val, y_pred_val, y_prob_val, val_infer_time = evaluate_model(model, val_loader, "Validation Set", plot_curves)
y_true_test, y_pred_test, y_prob_test, test_infer_time = evaluate_model(model, test_loader, "Test Set", plot_curves)

print("\n===== Summary =====")
print(f"Training inference time: {train_infer_time:.2f} sec")
print(f"Validation inference time: {val_infer_time:.2f} sec")
print(f"Test inference time: {test_infer_time:.2f} sec")

```

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
`warnings.warn(`

```


Epoch [1/35] Train Loss: 0.4105, Train Acc: 0.9267 Val Loss: 0.0459, Val Acc: 0.9908 Time: 50.33s
Epoch [2/35] Train Loss: 0.0619, Train Acc: 0.9869 Val Loss: 0.0157, Val Acc: 0.9931 Time: 49.89s
Epoch [3/35] Train Loss: 0.0336, Train Acc: 0.9929 Val Loss: 0.0193, Val Acc: 0.9931 Time: 50.00s
Epoch [4/35] Train Loss: 0.0278, Train Acc: 0.9933 Val Loss: 0.0182, Val Acc: 0.9954 Time: 50.18s
Epoch [5/35] Train Loss: 0.0126, Train Acc: 0.9970 Val Loss: 0.0169, Val Acc: 0.9977 Time: 50.06s
Epoch [6/35] Train Loss: 0.0096, Train Acc: 0.9980 Val Loss: 0.0204, Val Acc: 0.9954 Time: 50.04s
Epoch [7/35] Train Loss: 0.0070, Train Acc: 0.9990 Val Loss: 0.0185, Val Acc: 0.9954 Time: 50.21s
Early stopping triggered!
Training finished Best model saved as mobilevit_cotton_best.pth

```

--- Train Set ---

	precision	recall	f1-score	support
0	1.0000	1.0000	1.0000	1264
1	0.9965	1.0000	0.9983	3460
2	1.0000	0.9979	0.9990	1948
3	1.0000	0.9958	0.9979	1924
accuracy			0.9986	8596
macro avg	0.9991	0.9984	0.9988	8596
weighted avg	0.9986	0.9986	0.9986	8596


Train Set Confusion Matrix



MCC: 0.9980
Cohen's Kappa: 0.9980
Mean NPV: 0.9996
Mean PPV (Precision): 0.9991
Train Set inference time: 15.54 sec

--- Validation Set ---

	precision	recall	f1-score	support
0	1.0000	1.0000	1.0000	64
1	0.9942	0.9942	0.9942	171
2	0.9897	0.9796	0.9846	98
3	0.9901	1.0000	0.9950	100
accuracy			0.9931	433
macro avg	0.9935	0.9934	0.9934	433
weighted avg	0.9931	0.9931	0.9931	433



MCC: 0.9904

Cohen's Kappa: 0.9903


Mean NPV: 0.9976

Mean PPV (Precision): 0.9935



Validation Set inference time: 0.97 sec

--- Test Set ---

	precision	recall	f1-score	support
0	1.0000	1.0000	1.0000	92
1	1.0000	1.0000	1.0000	263
2	1.0000	0.9931	0.9965	145
3	0.9933	1.0000	0.9967	149
accuracy			0.9985	649
macro avg	0.9983	0.9983	0.9983	649
weighted avg	0.9985	0.9985	0.9985	649



MCC: 0.9978
 Cohen's Kappa: 0.9978
 Mean NPV: 0.9995
 Mean PPV (Precision): 0.9983
 ROC AUC: 1.0000, PR AUC: 1.0000



Test Set inference time: 1.37 sec

===== Summary =====
Training inference time: 15.54 sec
Validation inference time: 0.97 sec
Test inference time: 1.37 sec

In [12]: `import torch`

```

import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
from timm import create_model
from torch.optim.lr_scheduler import ReduceLROnPlateau

from sklearn.metrics import (
    classification_report, confusion_matrix, matthews_corrcoef,
    roc_auc_score, average_precision_score, roc_curve,
    precision_recall_curve, cohen_kappa_score
)
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import time

# =====
# Config for Cotton Dataset
# =====
train_dir = "/kaggle/working/cotton_train_aug" # Augmented train
val_dir = "/kaggle/working/cotton_split/val"
test_dir = "/kaggle/working/cotton_split/test"

batch_size = 32
num_epochs = 35
patience = 5
num_classes = 4
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
class_names = ["diseased cotton leaf", "diseased cotton plant", "fresh cotton leaf", "fresh cotton plant"]

# =====
# Data Transforms & Dataloaders
# =====
common_tfms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])

train_ds = datasets.ImageFolder(root=train_dir, transform=common_tfms)
val_ds = datasets.ImageFolder(root=val_dir, transform=common_tfms)
test_ds = datasets.ImageFolder(root=test_dir, transform=common_tfms)

train_loader = DataLoader(train_ds, batch_size=batch_size, shuffle=True, num_workers=2)
val_loader = DataLoader(val_ds, batch_size=batch_size, shuffle=False, num_workers=2)
test_loader = DataLoader(test_ds, batch_size=batch_size, shuffle=False, num_workers=2)

# =====
# Model
# =====
model = create_model("tiny_vit_5m_224", pretrained=True, num_classes=num_classes)
model = model.to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)
scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=True)

# =====
# Training Loop with Early Stopping
# =====
best_val_loss = float("inf")
patience_counter = 0

for epoch in range(num_epochs):
    start_time = time.time()

    # ---- Train ----
    model.train()
    train_loss, train_correct = 0, 0
    for imgs, labels in train_loader:
        imgs, labels = imgs.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(imgs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        train_loss += loss.item() * imgs.size(0)
        train_correct += (outputs.argmax(1) == labels).sum().item()

    train_loss /= len(train_loader.dataset)

```

```

train_acc = train_correct / len(train_loader.dataset)

# ---- Validation ----
model.eval()
val_loss, val_correct = 0, 0
with torch.no_grad():
    for imgs, labels in val_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        outputs = model(imgs)
        loss = criterion(outputs, labels)

        val_loss += loss.item() * imgs.size(0)
        val_correct += (outputs.argmax(1) == labels).sum().item()

val_loss /= len(val_loader.dataset)
val_acc = val_correct / len(val_loader.dataset)

scheduler.step(val_loss)

elapsed = time.time() - start_time
print(f"Epoch [{epoch+1}/{num_epochs}] "
      f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f} "
      f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f} "
      f"Time: {elapsed:.2f}s")

# Early Stopping
if val_loss < best_val_loss:
    best_val_loss = val_loss
    patience_counter = 0
    torch.save(model.state_dict(), "tiny_vit_best.pth")
else:
    patience_counter += 1
    if patience_counter >= patience:
        print("Early stopping triggered!")
        break

print("Training finished Best model saved as tiny_vit_best.pth")

# =====
# Evaluation Functions
# =====
def plot_confusion_matrix(cm, classes, title="Confusion Matrix"):
    plt.figure(figsize=(6,5))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=classes, yticklabels=classes)
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title(title)
    plt.show()

def print_report(y_true, y_pred, y_prob, title="", plot_curves=False):
    print(f"\n--- {title} ---")
    print(classification_report(y_true, y_pred, digits=4))

    cm = confusion_matrix(y_true, y_pred)
    plot_confusion_matrix(cm, class_names, title=f"{title} Confusion Matrix")

    mcc = matthews_corrcoef(y_true, y_pred)
    print(f"MCC: {mcc:.4f}")

    kappa = cohen_kappa_score(y_true, y_pred)
    print(f"Cohen's Kappa: {kappa:.4f}")

    # Class-wise NPV + PPV
    npv_list, ppv_list = [], []
    for i in range(len(cm)):
        TN = np.sum(np.delete(np.delete(cm, i, axis=0), i, axis=1))
        FN = np.sum(cm[i, :]) - cm[i, i]
        FP = np.sum(cm[:, i]) - cm[i, i]
        TP = cm[i, i]

        NPV = TN / (TN + FN) if (TN + FN) > 0 else 0
        PPV = TP / (TP + FP) if (TP + FP) > 0 else 0
        npv_list.append(NPV)
        ppv_list.append(PPV)

    print(f"Mean NPV: {np.mean(npv_list):.4f}")
    print(f"Mean PPV (Precision): {np.mean(ppv_list):.4f}")

    if plot_curves:
        y_onehot = np.eye(num_classes)[y_true]
        roc_auc = roc_auc_score(y_onehot, y_prob, average='macro', multi_class='ovr')
        pr_auc = average_precision_score(y_onehot, y_prob, average='macro')

```

```

print(f"ROC AUC: {roc_auc:.4f}, PR AUC: {pr_auc:.4f}")

# ROC Curve
plt.figure(figsize=(6,5))
for i in range(num_classes):
    fpr, tpr, _ = roc_curve(y_onehot[:,i], y_prob[:,i])
    plt.plot(fpr, tpr, label=f"{class_names[i]} (AUC={roc_auc_score(y_onehot[:,i], y_prob[:,i]):.4f})")
plt.plot([0,1],[0,1], 'k--')
plt.title(f"{title} ROC Curve")
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.legend()
plt.show()

# PR Curve
plt.figure(figsize=(6,5))
for i in range(num_classes):
    precision, recall, _ = precision_recall_curve(y_onehot[:,i], y_prob[:,i])
    plt.plot(recall, precision, label=f"{class_names[i]} (AP={average_precision_score(y_onehot[:,i], y_prob[:,i]):.4f})")
plt.title(f"{title} Precision-Recall Curve")
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.legend()
plt.show()

def evaluate_model(model, loader, title="", plot_curves=False):
    model.eval()
    y_true, y_pred, y_prob = [], [], []
    start_time = time.time()

    with torch.no_grad():
        for imgs, labels in loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            probs = torch.softmax(outputs, dim=1)
            preds = outputs.argmax(1)

            y_true.extend(labels.cpu().numpy())
            y_pred.extend(preds.cpu().numpy())
            y_prob.extend(probs.cpu().numpy())

    infer_time = time.time() - start_time
    y_true, y_pred, y_prob = np.array(y_true), np.array(y_pred), np.array(y_prob)

    print_report(y_true, y_pred, y_prob, title, plot_curves)
    print(f"{title} inference time: {infer_time:.2f} sec")
    return y_true, y_pred, y_prob, infer_time

# =====
# Final Evaluation
# =====
model.load_state_dict(torch.load("tiny_vit_best.pth"))

y_true_train, y_pred_train, y_prob_train, train_infer_time = evaluate_model(model, train_loader, "Train Set", plot_curves)
y_true_val, y_pred_val, y_prob_val, val_infer_time = evaluate_model(model, val_loader, "Validation Set", plot_curves)
y_true_test, y_pred_test, y_prob_test, test_infer_time = evaluate_model(model, test_loader, "Test Set", plot_curves)

print("\n===== Summary =====")
print(f"Training inference time: {train_infer_time:.2f} sec")
print(f"Validation inference time: {val_infer_time:.2f} sec")
print(f"Test inference time: {test_infer_time:.2f} sec")

model.safetensors: 0% | 0.00/48.4M [00:00<?, ?B/s]
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
  warnings.warn(

```

```


Epoch [1/35] Train Loss: 0.2403, Train Acc: 0.9359 Val Loss: 0.0451, Val Acc: 0.9931 Time: 39.18s
Epoch [2/35] Train Loss: 0.0320, Train Acc: 0.9929 Val Loss: 0.0309, Val Acc: 0.9931 Time: 39.02s
Epoch [3/35] Train Loss: 0.0201, Train Acc: 0.9948 Val Loss: 0.0203, Val Acc: 0.9954 Time: 39.10s
Epoch [4/35] Train Loss: 0.0083, Train Acc: 0.9984 Val Loss: 0.0248, Val Acc: 0.9908 Time: 38.87s
Epoch [5/35] Train Loss: 0.0094, Train Acc: 0.9973 Val Loss: 0.0267, Val Acc: 0.9931 Time: 38.88s
Epoch [6/35] Train Loss: 0.0059, Train Acc: 0.9987 Val Loss: 0.0175, Val Acc: 0.9954 Time: 38.98s
Epoch [7/35] Train Loss: 0.0036, Train Acc: 0.9991 Val Loss: 0.0165, Val Acc: 0.9977 Time: 39.24s
Epoch [8/35] Train Loss: 0.0040, Train Acc: 0.9986 Val Loss: 0.0170, Val Acc: 0.9977 Time: 39.03s
Epoch [9/35] Train Loss: 0.0073, Train Acc: 0.9978 Val Loss: 0.0245, Val Acc: 0.9954 Time: 38.90s
Epoch [10/35] Train Loss: 0.0043, Train Acc: 0.9988 Val Loss: 0.0433, Val Acc: 0.9908 Time: 39.07s
Epoch [11/35] Train Loss: 0.0024, Train Acc: 0.9994 Val Loss: 0.0188, Val Acc: 0.9977 Time: 38.99s
Epoch [12/35] Train Loss: 0.0007, Train Acc: 1.0000 Val Loss: 0.0187, Val Acc: 0.9977 Time: 39.02s
Early stopping triggered!
Training finished Best model saved as tiny_vit_best.pth

```

--- Train Set ---

	precision	recall	f1-score	support
0	1.0000	1.0000	1.0000	1264
1	1.0000	0.9997	0.9999	3460
2	1.0000	1.0000	1.0000	1948
3	0.9995	1.0000	0.9997	1924
accuracy			0.9999	8596
macro avg	0.9999	0.9999	0.9999	8596
weighted avg	0.9999	0.9999	0.9999	8596

Train Set Confusion Matrix



MCC: 0.9998

Cohen's Kappa: 0.9998

Mean NPV: 1.0000


Mean PPV (Precision): 0.9999

Train Set inference time: 12.94 sec

--- Validation Set ---

	precision	recall	f1-score	support
0	1.0000	1.0000	1.0000	64
1	1.0000	1.0000	1.0000	171
2	1.0000	0.9898	0.9949	98
3	0.9901	1.0000	0.9950	100
accuracy			0.9977	433
macro avg	0.9975	0.9974	0.9975	433
weighted avg	0.9977	0.9977	0.9977	433

Validation Set Confusion Matrix



MCC: 0.9968

Cohen's Kappa: 0.9968


Mean NPV: 0.9993

Mean PPV (Precision): 0.9975



Validation Set inference time: 0.84 sec

--- Test Set ---

	precision	recall	f1-score	support
0	1.0000	1.0000	1.0000	92
1	1.0000	0.9962	0.9981	263
2	0.9931	0.9931	0.9931	145
3	0.9933	1.0000	0.9967	149
accuracy			0.9969	649
macro avg	0.9966	0.9973	0.9970	649
weighted avg	0.9969	0.9969	0.9969	649



MCC: 0.9957
Cohen's Kappa: 0.9957
Mean NPV: 0.9989
Mean PPV (Precision): 0.9966
ROC AUC: 1.0000, PR AUC: 1.0000



Test Set inference time: 1.24 sec

```
===== Summary =====
Training inference time: 12.94 sec
Validation inference time: 0.84 sec
Test inference time: 1.24 sec
```

```
In [13]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
from timm import create_model
```

```

from torch.optim.lr_scheduler import ReduceLROnPlateau

from sklearn.metrics import (
    classification_report, confusion_matrix, matthews_corrcoef,
    roc_auc_score, average_precision_score, roc_curve,
    precision_recall_curve, cohen_kappa_score
)
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import time

# =====
# Config for Cotton Dataset
# =====
train_dir = "/kaggle/working/cotton_train_aug" # Augmented train
val_dir = "/kaggle/working/cotton_split/val"
test_dir = "/kaggle/working/cotton_split/test"

batch_size = 32
num_epochs = 35
patience = 5
num_classes = 4
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
class_names = ["diseased cotton leaf", "diseased cotton plant", "fresh cotton leaf", "fresh cotton plant"]

# =====
# Data Transforms & Dataloaders
# =====
common_tfms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])

train_ds = datasets.ImageFolder(root=train_dir, transform=common_tfms)
val_ds = datasets.ImageFolder(root=val_dir, transform=common_tfms)
test_ds = datasets.ImageFolder(root=test_dir, transform=common_tfms)

train_loader = DataLoader(train_ds, batch_size=batch_size, shuffle=True, num_workers=2)
val_loader = DataLoader(val_ds, batch_size=batch_size, shuffle=False, num_workers=2)
test_loader = DataLoader(test_ds, batch_size=batch_size, shuffle=False, num_workers=2)

# =====
# Model
# =====
model = create_model("levit_128", pretrained=True, num_classes=num_classes) # smaller LeViT
model = model.to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)
scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=True)

# =====
# Training Loop with Early Stopping
# =====
best_val_loss = float("inf")
patience_counter = 0

for epoch in range(num_epochs):
    start_time = time.time()

    # ---- Train ----
    model.train()
    train_loss, train_correct = 0, 0
    for imgs, labels in train_loader:
        imgs, labels = imgs.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(imgs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        train_loss += loss.item() * imgs.size(0)
        train_correct += (outputs.argmax(1) == labels).sum().item()

    train_loss /= len(train_loader.dataset)
    train_acc = train_correct / len(train_loader.dataset)

    # ---- Validation ----
    model.eval()
    val_loss, val_correct = 0, 0

```

```

with torch.no_grad():
    for imgs, labels in val_loader:
        imgs, labels = imgs.to(device), labels.to(device)
        outputs = model(imgs)
        loss = criterion(outputs, labels)

        val_loss += loss.item() * imgs.size(0)
        val_correct += (outputs.argmax(1) == labels).sum().item()

val_loss /= len(val_loader.dataset)
val_acc = val_correct / len(val_loader.dataset)

scheduler.step(val_loss)

elapsed = time.time() - start_time
print(f"Epoch [{epoch+1}/{num_epochs}] "
      f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f} "
      f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f} "
      f"Time: {elapsed:.2f}s")

# Early Stopping
if val_loss < best_val_loss:
    best_val_loss = val_loss
    patience_counter = 0
    torch.save(model.state_dict(), "tiny_vit_best.pth")
else:
    patience_counter += 1
    if patience_counter >= patience:
        print("Early stopping triggered!")
        break

print("Training finished Best model saved as tiny_vit_best.pth")

# =====
# Evaluation Functions
# =====

def plot_confusion_matrix(cm, classes, title="Confusion Matrix"):
    plt.figure(figsize=(6,5))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=classes, yticklabels=classes)
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title(title)
    plt.show()

def print_report(y_true, y_pred, y_prob, title="", plot_curves=False):
    print(f"\n--- {title} ---")
    print(classification_report(y_true, y_pred, digits=4))

    cm = confusion_matrix(y_true, y_pred)
    plot_confusion_matrix(cm, class_names, title=f"{title} Confusion Matrix")

    mcc = matthews_corrcoef(y_true, y_pred)
    print(f"MCC: {mcc:.4f}")

    kappa = cohen_kappa_score(y_true, y_pred)
    print(f"Cohen's Kappa: {kappa:.4f}")

    # Class-wise NPV + PPV
    npv_list, ppv_list = [], []
    for i in range(len(cm)):
        TN = np.sum(np.delete(cm, i, axis=0), i, axis=1)
        FN = np.sum(cm[i, :]) - cm[i, i]
        FP = np.sum(cm[:, i]) - cm[i, i]
        TP = cm[i, i]

        NPV = TN / (TN + FN) if (TN + FN) > 0 else 0
        PPV = TP / (TP + FP) if (TP + FP) > 0 else 0
        npv_list.append(NPV)
        ppv_list.append(PPV)

    print(f"Mean NPV: {np.mean(npv_list):.4f}")
    print(f"Mean PPV (Precision): {np.mean(ppv_list):.4f}")

    if plot_curves:
        y_onehot = np.eye(num_classes)[y_true]
        roc_auc = roc_auc_score(y_onehot, y_prob, average='macro', multi_class='ovr')
        pr_auc = average_precision_score(y_onehot, y_prob, average='macro')
        print(f"ROC AUC: {roc_auc:.4f}, PR AUC: {pr_auc:.4f}")

    # ROC Curve
    plt.figure(figsize=(6,5))
    for i in range(num_classes):

```

```

        fpr, tpr, _ = roc_curve(y_onehot[:,i], y_prob[:,i])
        plt.plot(fpr, tpr, label=f"class_names[i] (AUC={roc_auc_score(y_onehot[:,i], y_prob[:,i]):.4f})")
plt.plot([0,1],[0,1],'k--')
plt.title(f"{title} ROC Curve")
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.legend()
plt.show()

# PR Curve
plt.figure(figsize=(6,5))
for i in range(num_classes):
    precision, recall, _ = precision_recall_curve(y_onehot[:,i], y_prob[:,i])
    plt.plot(recall, precision, label=f"class_names[i] (AP={average_precision_score(y_onehot[:,i], y_prob[:,i]):.4f})")
plt.title(f"{title} Precision-Recall Curve")
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.legend()
plt.show()

def evaluate_model(model, loader, title="", plot_curves=False):
    model.eval()
    y_true, y_pred, y_prob = [], [], []
    start_time = time.time()

    with torch.no_grad():
        for imgs, labels in loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            probs = torch.softmax(outputs, dim=1)
            preds = outputs.argmax(1)

            y_true.extend(labels.cpu().numpy())
            y_pred.extend(preds.cpu().numpy())
            y_prob.extend(probs.cpu().numpy())

    infer_time = time.time() - start_time
    y_true, y_pred, y_prob = np.array(y_true), np.array(y_pred), np.array(y_prob)

    print_report(y_true, y_pred, y_prob, title, plot_curves)
    print(f"{title} inference time: {infer_time:.2f} sec")
    return y_true, y_pred, y_prob, infer_time

# =====
# Final Evaluation
# =====
model.load_state_dict(torch.load("tiny_vit_best.pth"))

y_true_train, y_pred_train, y_prob_train, train_infer_time = evaluate_model(model, train_loader, "Train Set", plot_curves)
y_true_val, y_pred_val, y_prob_val, val_infer_time = evaluate_model(model, val_loader, "Validation Set", plot_curves)
y_true_test, y_pred_test, y_prob_test, test_infer_time = evaluate_model(model, test_loader, "Test Set", plot_curves)

print("\n===== Summary =====")
print(f"Training inference time: {train_infer_time:.2f} sec")
print(f"Validation inference time: {val_infer_time:.2f} sec")
print(f"Test inference time: {test_infer_time:.2f} sec")

model.safetensors: 0% | 0.00/37.1M [00:00<?, ?B/s]
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
  warnings.warn(

```

```

Epoch [1/35] Train Loss: 0.6636, Train Acc: 0.8213 Val Loss: 5.8121, Val Acc: 0.9492 Time: 18.83s
Epoch [2/35] Train Loss: 0.1902, Train Acc: 0.9532 Val Loss: 0.2565, Val Acc: 0.9815 Time: 18.88s
Epoch [3/35] Train Loss: 0.1058, Train Acc: 0.9716 Val Loss: 2.9911, Val Acc: 0.9861 Time: 18.77s
Epoch [4/35] Train Loss: 0.0666, Train Acc: 0.9824 Val Loss: 0.0516, Val Acc: 0.9908 Time: 18.82s
Epoch [5/35] Train Loss: 0.0978, Train Acc: 0.9699 Val Loss: 0.0448, Val Acc: 0.9908 Time: 19.02s
Epoch [6/35] Train Loss: 0.0855, Train Acc: 0.9739 Val Loss: 0.7255, Val Acc: 0.9746 Time: 18.82s
Epoch [7/35] Train Loss: 0.0697, Train Acc: 0.9789 Val Loss: 0.0359, Val Acc: 0.9908 Time: 19.11s
Epoch [8/35] Train Loss: 0.0614, Train Acc: 0.9786 Val Loss: 0.0321, Val Acc: 0.9931 Time: 19.05s
Epoch [9/35] Train Loss: 0.0420, Train Acc: 0.9866 Val Loss: 0.3267, Val Acc: 0.9908 Time: 18.97s
Epoch [10/35] Train Loss: 0.0361, Train Acc: 0.9899 Val Loss: 0.0174, Val Acc: 0.9931 Time: 19.16s
Epoch [11/35] Train Loss: 0.0295, Train Acc: 0.9912 Val Loss: 0.0878, Val Acc: 0.9954 Time: 19.18s
Epoch [12/35] Train Loss: 0.0263, Train Acc: 0.9924 Val Loss: 0.0206, Val Acc: 0.9931 Time: 19.30s
Epoch [13/35] Train Loss: 0.0179, Train Acc: 0.9944 Val Loss: 0.0184, Val Acc: 0.9931 Time: 19.09s
Epoch [14/35] Train Loss: 0.0150, Train Acc: 0.9965 Val Loss: 0.0135, Val Acc: 0.9954 Time: 19.38s
Epoch [15/35] Train Loss: 0.0133, Train Acc: 0.9973 Val Loss: 0.0151, Val Acc: 0.9931 Time: 18.97s
Epoch [16/35] Train Loss: 0.0098, Train Acc: 0.9973 Val Loss: 0.0125, Val Acc: 0.9954 Time: 19.18s
Epoch [17/35] Train Loss: 0.0100, Train Acc: 0.9974 Val Loss: 0.0164, Val Acc: 0.9954 Time: 18.76s
Epoch [18/35] Train Loss: 0.0090, Train Acc: 0.9979 Val Loss: 0.0220, Val Acc: 0.9931 Time: 18.66s
Epoch [19/35] Train Loss: 0.0075, Train Acc: 0.9985 Val Loss: 0.0190, Val Acc: 0.9931 Time: 19.26s
Epoch [20/35] Train Loss: 0.0086, Train Acc: 0.9980 Val Loss: 0.0242, Val Acc: 0.9931 Time: 18.91s
Epoch [21/35] Train Loss: 0.0079, Train Acc: 0.9983 Val Loss: 0.0472, Val Acc: 0.9931 Time: 18.98s
Early stopping triggered!


```

Training finished Best model saved as tiny_vit_best.pth

--- Train Set ---

	precision	recall	f1-score	support
0	0.9781	0.9905	0.9843	1264
1	0.9917	0.9991	0.9954	3460
2	0.9937	0.9743	0.9839	1948
3	0.9990	0.9969	0.9979	1924
accuracy			0.9917	8596
macro avg	0.9906	0.9902	0.9904	8596
weighted avg	0.9918	0.9917	0.9917	8596

Train Set Confusion Matrix




MCC: 0.9885
 Cohen's Kappa: 0.9884
 Mean NPV: 0.9973
 Mean PPV (Precision): 0.9906
 Train Set inference time: 11.03 sec

--- Validation Set ---

	precision	recall	f1-score	support
0	1.0000	1.0000	1.0000	64
1	1.0000	0.9942	0.9971	171
2	1.0000	0.9898	0.9949	98
3	0.9804	1.0000	0.9901	100
accuracy			0.9954	433
macro avg	0.9951	0.9960	0.9955	433
weighted avg	0.9955	0.9954	0.9954	433


Validation Set Confusion Matrix




MCC: 0.9936
 Cohen's Kappa: 0.9936
 Mean NPV: 0.9983
 Mean PPV (Precision): 0.9951
 Validation Set inference time: 0.73 sec


--- Test Set ---

	precision	recall	f1-score	support
0	1.0000	1.0000	1.0000	92
1	0.9925	1.0000	0.9962	263
2	1.0000	0.9862	0.9931	145
3	0.9933	0.9933	0.9933	149
accuracy			0.9954	649
macro avg	0.9964	0.9949	0.9956	649
weighted avg	0.9954	0.9954	0.9954	649



MCC: 0.9935
 Cohen's Kappa: 0.9935
 Mean NPV: 0.9985
 Mean PPV (Precision): 0.9964
 ROC AUC: 0.9979, PR AUC: 0.9979





Test Set inference time: 1.02 sec

===== Summary =====
 Training inference time: 11.03 sec
 Validation inference time: 0.73 sec
 Test inference time: 1.02 sec

```
In [15]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
from timm import create_model
from torch.optim.lr_scheduler import ReduceLROnPlateau

from sklearn.metrics import (
    classification_report, confusion_matrix, matthews_corrcoef,
    roc_auc_score, average_precision_score, roc_curve,
    precision_recall_curve, cohen_kappa_score
)
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import time

# =====
# Config for Cotton Dataset
# =====
train_dir = "/kaggle/working/cotton_train_aug" # Augmented train
val_dir = "/kaggle/working/cotton_split/val"
test_dir = "/kaggle/working/cotton_split/test"

batch_size = 32
num_epochs = 35
patience = 5
num_classes = 4
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
class_names = ["diseased cotton leaf", "diseased cotton plant", "fresh cotton leaf", "fresh cotton plant"]

# =====
# Data Transforms & Dataloaders
# =====
common_tfms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])

train_ds = datasets.ImageFolder(root=train_dir, transform=common_tfms)
val_ds = datasets.ImageFolder(root=val_dir, transform=common_tfms)
test_ds = datasets.ImageFolder(root=test_dir, transform=common_tfms)
```

```

train_loader = DataLoader(train_ds, batch_size=batch_size, shuffle=True, num_workers=2)
val_loader = DataLoader(val_ds, batch_size=batch_size, shuffle=False, num_workers=2)
test_loader = DataLoader(test_ds, batch_size=batch_size, shuffle=False, num_workers=2)

# =====
# Model
# =====
model = create_model("deit_tiny_patch16_224", pretrained=True, num_classes=num_classes)
model = model.to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)
scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=True)

# =====
# Training Loop with Early Stopping
# =====
best_val_loss = float("inf")
patience_counter = 0

for epoch in range(num_epochs):
    start_time = time.time()

    # ---- Train ----
    model.train()
    train_loss, train_correct = 0, 0
    for imgs, labels in train_loader:
        imgs, labels = imgs.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(imgs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        train_loss += loss.item() * imgs.size(0)
        train_correct += (outputs.argmax(1) == labels).sum().item()

    train_loss /= len(train_loader.dataset)
    train_acc = train_correct / len(train_loader.dataset)

    # ---- Validation ----
    model.eval()
    val_loss, val_correct = 0, 0
    with torch.no_grad():
        for imgs, labels in val_loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            loss = criterion(outputs, labels)

            val_loss += loss.item() * imgs.size(0)
            val_correct += (outputs.argmax(1) == labels).sum().item()

    val_loss /= len(val_loader.dataset)
    val_acc = val_correct / len(val_loader.dataset)

    scheduler.step(val_loss)

    elapsed = time.time() - start_time
    print(f"Epoch [{epoch+1}/{num_epochs}] "
          f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f} "
          f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f} "
          f"Time: {elapsed:.2f}s")

    # Early Stopping
    if val_loss < best_val_loss:
        best_val_loss = val_loss
        patience_counter = 0
        torch.save(model.state_dict(), "tiny_vit_best.pth")
    else:
        patience_counter += 1
        if patience_counter >= patience:
            print("Early stopping triggered!")
            break

print("Training finished Best model saved as tiny_vit_best.pth")

# =====
# Evaluation Functions
# =====
def plot_confusion_matrix(cm, classes, title="Confusion Matrix"):
    plt.figure(figsize=(6,5))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=classes, yticklabels=classes)

```

```

plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title(title)
plt.show()

def print_report(y_true, y_pred, y_prob, title="", plot_curves=False):
    print(f"\n--- {title} ---")
    print(classification_report(y_true, y_pred, digits=4))

    cm = confusion_matrix(y_true, y_pred)
    plot_confusion_matrix(cm, class_names, title=f"{title} Confusion Matrix")

    mcc = matthews_corrcoef(y_true, y_pred)
    print(f"MCC: {mcc:.4f}")

    kappa = cohen_kappa_score(y_true, y_pred)
    print(f"Cohen's Kappa: {kappa:.4f}")

    # Class-wise NPV + PPV
    npv_list, ppv_list = [], []
    for i in range(len(cm)):
        TN = np.sum(np.delete(np.delete(cm, i, axis=0), i, axis=1))
        FN = np.sum(cm[i, :]) - cm[i, i]
        FP = np.sum(cm[:, i]) - cm[i, i]
        TP = cm[i, i]

        NPV = TN / (TN + FN) if (TN + FN) > 0 else 0
        PPV = TP / (TP + FP) if (TP + FP) > 0 else 0
        npv_list.append(NPV)
        ppv_list.append(PPV)

    print(f"Mean NPV: {np.mean(npv_list):.4f}")
    print(f"Mean PPV (Precision): {np.mean(ppv_list):.4f}")

    if plot_curves:
        y_onehot = np.eye(num_classes)[y_true]
        roc_auc = roc_auc_score(y_onehot, y_prob, average='macro', multi_class='ovr')
        pr_auc = average_precision_score(y_onehot, y_prob, average='macro')
        print(f"ROC AUC: {roc_auc:.4f}, PR AUC: {pr_auc:.4f}")

        # ROC Curve
        plt.figure(figsize=(6,5))
        for i in range(num_classes):
            fpr, tpr, _ = roc_curve(y_onehot[:,i], y_prob[:,i])
            plt.plot(fpr, tpr, label=f'{class_names[i]} (AUC={roc_auc_score(y_onehot[:,i], y_prob[:,i]):.4f})')
        plt.plot([0,1],[0,1], 'k--')
        plt.title(f'{title} ROC Curve')
        plt.xlabel("FPR")
        plt.ylabel("TPR")
        plt.legend()
        plt.show()

        # PR Curve
        plt.figure(figsize=(6,5))
        for i in range(num_classes):
            precision, recall, _ = precision_recall_curve(y_onehot[:,i], y_prob[:,i])
            plt.plot(recall, precision, label=f'{class_names[i]} (AP={average_precision_score(y_onehot[:,i], y_prob[:,i]):.4f})')
        plt.title(f'{title} Precision-Recall Curve')
        plt.xlabel("Recall")
        plt.ylabel("Precision")
        plt.legend()
        plt.show()

def evaluate_model(model, loader, title="", plot_curves=False):
    model.eval()
    y_true, y_pred, y_prob = [], [], []
    start_time = time.time()

    with torch.no_grad():
        for imgs, labels in loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            probs = torch.softmax(outputs, dim=1)
            preds = outputs.argmax(1)

            y_true.extend(labels.cpu().numpy())
            y_pred.extend(preds.cpu().numpy())
            y_prob.extend(probs.cpu().numpy())

    infer_time = time.time() - start_time
    y_true, y_pred, y_prob = np.array(y_true), np.array(y_pred), np.array(y_prob)

```

```

print_report(y_true, y_pred, y_prob, title, plot_curves)
print(f"{title} inference time: {infer_time:.2f} sec")
return y_true, y_pred, y_prob, infer_time

# =====
# Final Evaluation
# =====
model.load_state_dict(torch.load("tiny_vit_best.pth"))

y_true_train, y_pred_train, y_prob_train, train_infer_time = evaluate_model(model, train_loader, "Train Set", plot_curves)
y_true_val, y_pred_val, y_prob_val, val_infer_time = evaluate_model(model, val_loader, "Validation Set", plot_curves)
y_true_test, y_pred_test, y_prob_test, test_infer_time = evaluate_model(model, test_loader, "Test Set", plot_curves)

print("\n===== Summary =====")
print(f"Training inference time: {train_infer_time:.2f} sec")
print(f"Validation inference time: {val_infer_time:.2f} sec")
print(f"Test inference time: {test_infer_time:.2f} sec")

```

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.


```

warnings.warn(
Epoch [1/35] Train Loss: 0.1376, Train Acc: 0.9522 Val Loss: 0.0358, Val Acc: 0.9931 Time: 24.20s
Epoch [2/35] Train Loss: 0.0242, Train Acc: 0.9928 Val Loss: 0.0451, Val Acc: 0.9885 Time: 24.11s
Epoch [3/35] Train Loss: 0.0210, Train Acc: 0.9935 Val Loss: 0.0170, Val Acc: 0.9977 Time: 24.32s
Epoch [4/35] Train Loss: 0.0174, Train Acc: 0.9938 Val Loss: 0.0421, Val Acc: 0.9861 Time: 24.21s
Epoch [5/35] Train Loss: 0.0178, Train Acc: 0.9943 Val Loss: 0.0171, Val Acc: 0.9954 Time: 24.27s
Epoch [6/35] Train Loss: 0.0090, Train Acc: 0.9974 Val Loss: 0.0497, Val Acc: 0.9838 Time: 24.30s
Epoch [7/35] Train Loss: 0.0050, Train Acc: 0.9985 Val Loss: 0.0147, Val Acc: 0.9954 Time: 24.28s
Epoch [8/35] Train Loss: 0.0016, Train Acc: 0.9995 Val Loss: 0.0172, Val Acc: 0.9977 Time: 24.23s
Epoch [9/35] Train Loss: 0.0003, Train Acc: 1.0000 Val Loss: 0.0178, Val Acc: 0.9977 Time: 24.15s
Epoch [10/35] Train Loss: 0.0003, Train Acc: 1.0000 Val Loss: 0.0179, Val Acc: 0.9977 Time: 24.49s
Epoch [11/35] Train Loss: 0.0002, Train Acc: 1.0000 Val Loss: 0.0180, Val Acc: 0.9977 Time: 24.19s
Epoch [12/35] Train Loss: 0.0002, Train Acc: 1.0000 Val Loss: 0.0181, Val Acc: 0.9977 Time: 24.19s
Early stopping triggered!

```

Training finished Best model saved as tiny_vit_best.pth

--- Train Set ---					
	precision	recall	f1-score	support	
0	1.0000	1.0000	1.0000	1264	
1	1.0000	1.0000	1.0000	3460	
2	1.0000	1.0000	1.0000	1948	
3	1.0000	1.0000	1.0000	1924	
accuracy			1.0000	8596	
macro avg	1.0000	1.0000	1.0000	8596	
weighted avg	1.0000	1.0000	1.0000	8596	




MCC: 1.0000
 Cohen's Kappa: 1.0000
 Mean NPV: 1.0000
 Mean PPV (Precision): 1.0000
 Train Set inference time: 11.02 sec

--- Validation Set ---

	precision	recall	f1-score	support
0	1.0000	1.0000	1.0000	64
1	1.0000	0.9942	0.9971	171
2	0.9898	0.9898	0.9898	98
3	0.9901	1.0000	0.9950	100
accuracy			0.9954	433
macro avg	0.9950	0.9960	0.9955	433
weighted avg	0.9954	0.9954	0.9954	433


Validation Set Confusion Matrix




MCC: 0.9936
 Cohen's Kappa: 0.9936
 Mean NPV: 0.9983
 Mean PPV (Precision): 0.9950
 Validation Set inference time: 0.78 sec


--- Test Set ---

	precision	recall	f1-score	support
0	1.0000	1.0000	1.0000	92
1	1.0000	1.0000	1.0000	263
2	1.0000	0.9931	0.9965	145
3	0.9933	1.0000	0.9967	149
accuracy			0.9985	649
macro avg	0.9983	0.9983	0.9983	649
weighted avg	0.9985	0.9985	0.9985	649



MCC: 0.9978
 Cohen's Kappa: 0.9978
 Mean NPV: 0.9995
 Mean PPV (Precision): 0.9983
 ROC AUC: 1.0000, PR AUC: 0.9999





Test Set inference time: 1.02 sec

```
===== Summary =====
Training inference time: 11.02 sec
Validation inference time: 0.78 sec
Test inference time: 1.02 sec
```

```
In [16]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
from timm import create_model
from torch.optim.lr_scheduler import ReduceLROnPlateau

from sklearn.metrics import (
    classification_report, confusion_matrix, matthews_corrcoef,
    roc_auc_score, average_precision_score, roc_curve,
    precision_recall_curve, cohen_kappa_score
)
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import time

# =====
# Config for Cotton Dataset
# =====
train_dir = "/kaggle/working/cotton_train_aug" # Augmented train
val_dir = "/kaggle/working/cotton_split/val"
test_dir = "/kaggle/working/cotton_split/test"

batch_size = 32
num_epochs = 35
patience = 5
num_classes = 4
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
class_names = ["diseased cotton leaf", "diseased cotton plant", "fresh cotton leaf", "fresh cotton plant"]

# =====
# Data Transforms & Dataloaders
# =====
common_tfms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*3, [0.5]*3)
])

train_ds = datasets.ImageFolder(root=train_dir, transform=common_tfms)
val_ds = datasets.ImageFolder(root=val_dir, transform=common_tfms)
test_ds = datasets.ImageFolder(root=test_dir, transform=common_tfms)
```

```

train_loader = DataLoader(train_ds, batch_size=batch_size, shuffle=True, num_workers=2)
val_loader = DataLoader(val_ds, batch_size=batch_size, shuffle=False, num_workers=2)
test_loader = DataLoader(test_ds, batch_size=batch_size, shuffle=False, num_workers=2)

# =====
# Model
# =====
model = create_model("swin_tiny_patch4_window7_224", pretrained=True, num_classes=num_classes)
model = model.to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model.parameters(), lr=1e-4, weight_decay=1e-4)
scheduler = ReduceLROnPlateau(optimizer, mode="min", patience=2, factor=0.5, verbose=True)

# =====
# Training Loop with Early Stopping
# =====
best_val_loss = float("inf")
patience_counter = 0

for epoch in range(num_epochs):
    start_time = time.time()

    # ---- Train ----
    model.train()
    train_loss, train_correct = 0, 0
    for imgs, labels in train_loader:
        imgs, labels = imgs.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(imgs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        train_loss += loss.item() * imgs.size(0)
        train_correct += (outputs.argmax(1) == labels).sum().item()

    train_loss /= len(train_loader.dataset)
    train_acc = train_correct / len(train_loader.dataset)

    # ---- Validation ----
    model.eval()
    val_loss, val_correct = 0, 0
    with torch.no_grad():
        for imgs, labels in val_loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            loss = criterion(outputs, labels)

            val_loss += loss.item() * imgs.size(0)
            val_correct += (outputs.argmax(1) == labels).sum().item()

    val_loss /= len(val_loader.dataset)
    val_acc = val_correct / len(val_loader.dataset)

    scheduler.step(val_loss)

    elapsed = time.time() - start_time
    print(f"Epoch [{epoch+1}/{num_epochs}] "
          f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f} "
          f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f} "
          f"Time: {elapsed:.2f}s")

    # Early Stopping
    if val_loss < best_val_loss:
        best_val_loss = val_loss
        patience_counter = 0
        torch.save(model.state_dict(), "tiny_vit_best.pth")
    else:
        patience_counter += 1
        if patience_counter >= patience:
            print("Early stopping triggered!")
            break

print("Training finished Best model saved as tiny_vit_best.pth")

# =====
# Evaluation Functions
# =====
def plot_confusion_matrix(cm, classes, title="Confusion Matrix"):
    plt.figure(figsize=(6,5))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=classes, yticklabels=classes)

```

```

plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title(title)
plt.show()

def print_report(y_true, y_pred, y_prob, title="", plot_curves=False):
    print(f"\n--- {title} ---")
    print(classification_report(y_true, y_pred, digits=4))

    cm = confusion_matrix(y_true, y_pred)
    plot_confusion_matrix(cm, class_names, title=f"{title} Confusion Matrix")

    mcc = matthews_corrcoef(y_true, y_pred)
    print(f"MCC: {mcc:.4f}")

    kappa = cohen_kappa_score(y_true, y_pred)
    print(f"Cohen's Kappa: {kappa:.4f}")

    # Class-wise NPV + PPV
    npv_list, ppv_list = [], []
    for i in range(len(cm)):
        TN = np.sum(np.delete(np.delete(cm, i, axis=0), i, axis=1))
        FN = np.sum(cm[i, :]) - cm[i, i]
        FP = np.sum(cm[:, i]) - cm[i, i]
        TP = cm[i, i]

        NPV = TN / (TN + FN) if (TN + FN) > 0 else 0
        PPV = TP / (TP + FP) if (TP + FP) > 0 else 0
        npv_list.append(NPV)
        ppv_list.append(PPV)

    print(f"Mean NPV: {np.mean(npv_list):.4f}")
    print(f"Mean PPV (Precision): {np.mean(ppv_list):.4f}")

    if plot_curves:
        y_onehot = np.eye(num_classes)[y_true]
        roc_auc = roc_auc_score(y_onehot, y_prob, average='macro', multi_class='ovr')
        pr_auc = average_precision_score(y_onehot, y_prob, average='macro')
        print(f"ROC AUC: {roc_auc:.4f}, PR AUC: {pr_auc:.4f}")

        # ROC Curve
        plt.figure(figsize=(6,5))
        for i in range(num_classes):
            fpr, tpr, _ = roc_curve(y_onehot[:,i], y_prob[:,i])
            plt.plot(fpr, tpr, label=f'{class_names[i]} (AUC={roc_auc_score(y_onehot[:,i], y_prob[:,i]):.4f})')
        plt.plot([0,1],[0,1], 'k--')
        plt.title(f'{title} ROC Curve')
        plt.xlabel("FPR")
        plt.ylabel("TPR")
        plt.legend()
        plt.show()

        # PR Curve
        plt.figure(figsize=(6,5))
        for i in range(num_classes):
            precision, recall, _ = precision_recall_curve(y_onehot[:,i], y_prob[:,i])
            plt.plot(recall, precision, label=f'{class_names[i]} (AP={average_precision_score(y_onehot[:,i], y_prob[:,i]):.4f})')
        plt.title(f'{title} Precision-Recall Curve')
        plt.xlabel("Recall")
        plt.ylabel("Precision")
        plt.legend()
        plt.show()

def evaluate_model(model, loader, title="", plot_curves=False):
    model.eval()
    y_true, y_pred, y_prob = [], [], []
    start_time = time.time()

    with torch.no_grad():
        for imgs, labels in loader:
            imgs, labels = imgs.to(device), labels.to(device)
            outputs = model(imgs)
            probs = torch.softmax(outputs, dim=1)
            preds = outputs.argmax(1)

            y_true.extend(labels.cpu().numpy())
            y_pred.extend(preds.cpu().numpy())
            y_prob.extend(probs.cpu().numpy())

    infer_time = time.time() - start_time
    y_true, y_pred, y_prob = np.array(y_true), np.array(y_pred), np.array(y_prob)

```

```

print_report(y_true, y_pred, y_prob, title, plot_curves)
print(f"\{title\} inference time: {infer_time:.2f} sec")
return y_true, y_pred, y_prob, infer_time

# =====
# Final Evaluation
# =====
model.load_state_dict(torch.load("tiny_vit_best.pth"))

y_true_train, y_pred_train, y_prob_train, train_infer_time = evaluate_model(model, train_loader, "Train Set", plot_curves)
y_true_val, y_pred_val, y_prob_val, val_infer_time = evaluate_model(model, val_loader, "Validation Set", plot_curves)
y_true_test, y_pred_test, y_prob_test, test_infer_time = evaluate_model(model, test_loader, "Test Set", plot_curves)

print("\n===== Summary =====")
print(f"Training inference time: {train_infer_time:.2f} sec")
print(f"Validation inference time: {val_infer_time:.2f} sec")
print(f"Test inference time: {test_infer_time:.2f} sec")

```

model.safetensors: 0% | 0.00/114M [00:00<?, ?B/s]

/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.

warnings.warn(


Epoch [1/35] Train Loss: 0.1291, Train Acc: 0.9549 Val Loss: 0.0643, Val Acc: 0.9700 Time: 69.90s
 Epoch [2/35] Train Loss: 0.0262, Train Acc: 0.9912 Val Loss: 0.0094, Val Acc: 0.9954 Time: 69.79s
 Epoch [3/35] Train Loss: 0.0235, Train Acc: 0.9922 Val Loss: 0.0113, Val Acc: 0.9954 Time: 69.84s
 Epoch [4/35] Train Loss: 0.0163, Train Acc: 0.9960 Val Loss: 0.0247, Val Acc: 0.9931 Time: 69.93s
 Epoch [5/35] Train Loss: 0.0135, Train Acc: 0.9958 Val Loss: 0.0447, Val Acc: 0.9931 Time: 69.73s
 Epoch [6/35] Train Loss: 0.0041, Train Acc: 0.9988 Val Loss: 0.0172, Val Acc: 0.9931 Time: 69.85s
 Epoch [7/35] Train Loss: 0.0013, Train Acc: 0.9998 Val Loss: 0.0245, Val Acc: 0.9931 Time: 69.82s
 Early stopping triggered!

Training finished Best model saved as tiny_vit_best.pth

--- Train Set ---

	precision	recall	f1-score	support
0	0.9992	1.0000	0.9996	1264
1	0.9869	1.0000	0.9934	3460
2	0.9979	0.9949	0.9964	1948
3	0.9989	0.9777	0.9882	1924
accuracy			0.9938	8596
macro avg	0.9957	0.9931	0.9944	8596
weighted avg	0.9939	0.9938	0.9938	8596

Train Set Confusion Matrix




MCC: 0.9914
 Cohen's Kappa: 0.9914
 Mean NPV: 0.9980
 Mean PPV (Precision): 0.9957
 Train Set inference time: 24.34 sec

--- Validation Set ---

	precision	recall	f1-score	support
0	1.0000	1.0000	1.0000	64
1	0.9942	1.0000	0.9971	171
2	1.0000	0.9898	0.9949	98
3	0.9900	0.9900	0.9900	100
accuracy			0.9954	433
macro avg	0.9960	0.9949	0.9955	433
weighted avg	0.9954	0.9954	0.9954	433


Validation Set Confusion Matrix




MCC: 0.9936
 Cohen's Kappa: 0.9936
 Mean NPV: 0.9985
 Mean PPV (Precision): 0.9960
 Validation Set inference time: 1.44 sec


--- Test Set ---

	precision	recall	f1-score	support
0	1.0000	1.0000	1.0000	92
1	0.9962	1.0000	0.9981	263
2	1.0000	1.0000	1.0000	145
3	1.0000	0.9933	0.9966	149
accuracy			0.9985	649
macro avg	0.9991	0.9983	0.9987	649
weighted avg	0.9985	0.9985	0.9985	649



MCC: 0.9978
 Cohen's Kappa: 0.9978
 Mean NPV: 0.9995
 Mean PPV (Precision): 0.9991
 ROC AUC: 1.0000, PR AUC: 1.0000





Test Set inference time: 2.08 sec

===== Summary =====
Training inference time: 24.34 sec
Validation inference time: 1.44 sec
Test inference time: 2.08 sec